



An Introduction to Artificial Intelligence

— Lecture Notes in Progress —

Prof. Dr. Karl Stroetmann

March 22, 2018

These lecture notes, their \LaTeX sources, and the programs discussed in these lecture notes are all available at

<https://github.com/karlstroetmann/Artificial-Intelligence>.

In particular, the lecture notes are found in the directory `Lecture-Notes` in the file `artificial-intelligence.pdf`. The lecture notes are subject to continuous change. Provided the program `git` is installed on your computer, the repository containing the lecture notes can be cloned using the command

```
git clone https://github.com/karlstroetmann/Artificial-Intelligence.git.
```

Once you have cloned the repository, the command

```
git pull
```

can be used to load the current version of these lecture notes from `github`. As this is the second time that I give this lecture, these lecture notes are still incomplete and will be changed frequently during the semester. If you find any typos or inconsistencies, please email me at:

karl.stroetmann@dhbw-mannheim.de.

Contents

1	Introduction	3
1.1	What is Artificial Intelligence?	3
1.2	Literature	5
2	Search	6
2.1	The Sliding Puzzle	9
2.2	Breadth First Search	11
2.2.1	A Queue Based Implementation of Breadth First Search	13
2.3	Depth First Search	16
2.3.1	Getting Rid of the Parent Dictionary	17
2.3.2	A Recursive Implementation of Depth First Search	18
2.4	Iterative Deepening	19
2.4.1	A Recursive Implementation of Iterative Deepening	21
2.5	Bidirectional Breadth First Search	21
2.6	Best First Search	23
2.7	The A* Search Algorithm	27
2.7.1	Completeness and Optimality of A* Search	29
2.8	Bidirectional A* Search	33
2.9	Iterative Deepening A* Search	33
2.10	The A*-IDA* Search Algorithm	37
3	Constraint Satisfaction	42
3.1	Formal Definition of Constraint Satisfaction Problems	42
3.1.1	Example: Map Colouring	44
3.1.2	Example: The Eight Queens Puzzle	44
3.1.3	Applications	47
3.2	Brute Force Search	47
3.3	Backtracking Search	49
3.4	Constraint Propagation	52
3.5	Consistency Checking	57
3.6	Local Search	61
4	Playing Games	65
4.1	Tic-Tac-Toe	66
4.2	The Minimax Algorithm	68
4.3	α - β -Pruning	70
4.4	Depth Limited Search	72
5	Linear Regression	74
5.1	Simple Linear Regression	74
5.1.1	Assessing the Quality of Linear Regression	75
5.1.2	Putting the Theory to the Test	76

5.1.3	Testing the Statistical Significance	80
5.2	General Linear Regression	81
5.2.1	Some Useful Gradients	83
5.2.2	Deriving the Normal Equation	83
5.2.3	Testing the Statistical Significance	85
5.2.4	Implementation	85
6	Classification	88
6.1	Introduction	88
6.1.1	Notation	90
6.1.2	Applications of Classification	90
6.2	Digression: The Method of Gradient Ascent	91
6.3	Logistic Regression	93
6.3.1	The Sigmoid Function	94
6.3.2	The Model of Logistic Regression	96
6.3.3	Implementing Logistic Regression	98
7	Neural Networks	103
7.1	Feedforward Neural Networks	103
7.2	Backpropagation	106
7.2.1	Definition of some Auxiliary Variables	106
7.2.2	The Hadamard Product	107
7.2.3	Backpropagation: The Equations	107
7.2.4	Proof of the Backpropagation Equations	108
7.3	Stochastic Gradient Descent	110
7.4	Implementation	111

Chapter 1

Introduction

1.1 What is Artificial Intelligence?

Before we start to dive into the subject of **Artificial Intelligence** we have to answer the following question:

What is **Artificial Intelligence**?

Historically, there have been a number of different answers to this question [RN09]. We will look at these different answers and discuss them.

1. **Artificial Intelligence** is the study of creating machines that **think** like humans.

As we have a working prototype of intelligence, namely humans, it is quite natural to try to build machines that work in a way similar to humans, thereby creating artificial intelligence. As a first step in this endeavor we would have to study how humans actually think and thus we would have to study the brain. Unfortunately, as of today, no one really knows how the brain works. Although there are branches of science devoted to studying the human thought processes and the human brain, namely **cognitive science** and **computational neuroscience**, this approach has not proven to be fruitful for creating thinking machines, the reason being that the current understanding of the human thought processes is just not sufficient.

2. **Artificial Intelligence** is the science of machines that **act** like humans.

Since we do not know how humans think, we cannot build machines that think like people. Therefore, the next best thing might be to build machines that act and behave like humans. Actually, the **Turing Test** is based on this idea: Turing suggested that if we want to know whether we have succeeded in building an intelligent machine, we should place it at the other end of a chat line. If we cannot distinguish the computer from a human, then we have succeeded at creating intelligence.

However, with respect to the kind of Artificial Intelligence that is needed in industry, this approach isn't very useful. To illustrate the point, consider an analogy with aerodynamics: In aerodynamics we try to build planes that fly fast and efficiently, not planes that flap their wings like birds do, as the later approach has failed historically, e.g. **Daedalus and Icarus**.

3. **Artificial Intelligence** is the science of creating machines that **think logically**.

The idea with this approach is to create machines that are based on **mathematical logic**. If a goal is given to these machines, then these machines use logical reasoning in order to deduce those actions that need to be performed in order to best achieve the given goals. Unfortunately, this approach had only limited success: In playing games the approach was quite successful for dealing with games like checkers or chess. However, the approach was mostly unsuccessful for dealing with many real world problems. There were two main reasons for its failure:

- (a) In order for the logical approach to be successful, the environment has to be **completely** described by mathematical axioms. It has turned out that our knowledge of the real world is often not sufficient to completely describe the environment via axioms.

- (b) Even if we have complete knowledge, it often turns out that describing every possible case via logic formulae is just unwieldy. Consider the following formula:

$$\forall x : (\text{bird}(x) \rightarrow \text{flies}(x))$$

The problem with this formula is that although it appears to be common sense, there are a number of counter examples:

- i. Penguins, emus, and ostriches don't fly.
However, if we put a penguin into a plane, it turns out the penguin will fly.
- ii. Birds that have just hatched do not fly.
- iii. Birds with clipped wings do not fly.
- iv. ...

Trying to model all eventualities with logic has turned out to be too unwieldy to be practical.

- (c) In real life situations we often deal with **uncertainty**. Classical logic does not perform well when it has to deal with uncertainties.

4. **Artificial Intelligence** is the science of creating machines that **act rationally**.

All we really want is to build machines that, given the knowledge we have, try to **optimize the expected results**: In our world, there is lots of uncertainty. We cannot hope to create machines that always make the decisions that turn out to be optimal. What we can hope is to create machines that will make decisions that turn out to be good on average. For example, suppose we try to create a program for asset management: We cannot hope to build a machine that always buys the best company share in the stock market. Rather, our goal should be to build a program that maximizes our expected profits in the long term.

It has turned out that the main tool needed for this approach is not mathematical logic but rather **numerical analysis** and **mathematical statistics**. The shift from logic to numerical analysis and statistics has been the most important reason for the success of Artificial Intelligence in the recent years. Another important factor is the **enhanced performance of modern hardware**.

Now that we have clarified the notion of artificial intelligence, we should set its goals. As we can never achieve more than what we aim for, we have every reason to be ambitious here. For example, my personal vision of Artificial Intelligence goes like this: Imagine 70 years from now you (not feeling too well) have a conversation with **Siri**. Instead of asking Siri for the best graveyard in the vicinity, you think about all the sins you have committed. As Siri has accompanied you for your whole life, she knows about these sins better than you. Hence, the conversation with Siri works out as follows:

You (with trembling voice):

Hey Siri, does God exist?

Siri (with the voice of Darth Vader):

Your voice seems troubled, let me think ...

After a small pause which almost drains the battery of your phone completely, Siri gets back with a soothing announcement:

You don't have to worry any more, I have fixed the problem. He is dead now.

May **The Force** be with us on achieving our goals!

1.2 Literature

The main sources of these lecture notes are the following:

1. A course on artificial intelligence that was offered on the EDX platform. The course materials are available at

<http://ai.berkeley.edu/home.html>.

2. The book

[Introduction to Artificial Intelligence](#)

written by Stuart Russel and Peter Norvig [RN09].

3. A course on artificial intelligence that is offered on [Udacity](#). The title of the course is

[Intro to Artificial Intelligence](#)

and the course is given by [Peter Norvig](#), who is director of research at Google and [Sebastian Thrun](#), who is the chairman of [Udacity](#).

The programs presented in these lecture notes are expected to run with version 2.7 of [SETLX](#).

Chapter 2

Search

In this chapter we discuss various [search algorithms](#). First, we define the notion of a [search problem](#). As one of the examples, we will discuss the [sliding puzzle](#). Then we introduce various algorithms for solving search problems. In particular, we present

1. [breadth first search](#),
2. [depth first search](#),
3. [iterative deepening](#),
4. [bidirectional breadth first search](#),
5. [A* search](#),
6. [bidirectional A* search](#),
7. [iterative deepening A* search](#), and
8. [A*-IDA* search](#).

Definition 1 (Search Problem) A [search problem](#) is a tuple of the form

$$\mathcal{P} = \langle Q, \text{nextStates}, \text{start}, \text{goal} \rangle$$

where

1. Q is the set of states, also known as the [state space](#).
2. nextStates is a function taking a state as input and returning the set of those states that can be reached from the given state in one step, i.e. we have

$$\text{nextStates} : Q \rightarrow 2^Q.$$

The function nextStates gives rise to the [transition relation](#) R , which is a relation on Q , i.e. $R \subseteq Q \times Q$. This relation is defined as follows:

$$R := \{ \langle s_1, s_2 \rangle \in Q \times Q \mid s_2 \in \text{nextStates}(s_1) \}.$$

If either $\langle s_1, s_2 \rangle \in R$ or $\langle s_2, s_1 \rangle \in R$, then s_1 and s_2 are called [neighboring states](#).

3. start is the [start state](#), hence $\text{start} \in Q$.
4. goal is the [goal state](#), hence $\text{goal} \in Q$.

Sometimes, instead of a single goal there is a set of goal states G .

A **path** is a list $[s_1, \dots, s_n]$ such that $\langle s_i, s_{i+1} \rangle \in R$ for all $i \in \{1, \dots, n-1\}$. The **length** of this path is defined as the length of the list. A path $[s_1, \dots, s_n]$ is a **solution** to the search problem P iff the following conditions are satisfied:

1. $s_1 = \text{start}$, i.e. the first element of the path is the start state.
2. $s_n = \text{goal}$, i.e. the last element of the path is the goal state.

A path $p = [s_1, \dots, s_n]$ is a **minimal solution** to the search problem \mathcal{P} iff it is a solution and, furthermore, the length of p is minimal among all other solutions. \diamond

Remark: In the literature, a **state** is often called a **node**. In these lecture notes, I will also refer to states as nodes. \diamond

Example: We illustrate the notion of a search problem with the following example, which is also known as the **missionaries and cannibals problem**: Three missionaries and three infidels have to cross a river that runs from the west to the east. Initially, they are on the northern shore. There is just one small boat and that boat has only room for at most two passengers. Both the missionaries and the infidels can steer the boat. However, if at any time the missionaries are confronted with a majority of infidels on either shore of the river, then the missionaries have a problem.

```

1  problem := [m, i] |-> m > 0 && m < i;
2
3  noProblemAtAll := [m, i] |-> !problem(m, i) && !problem(3 - m, 3 - i);
4
5  nextStates := procedure(s) {
6      [m, i, b] := s;
7      if (b == 1) { // The boat is on the northern shore.
8          return { [m - mb, i - ib, 0]
9                  : mb in {0 .. m}, ib in {0 .. i}
10                 | mb + ib in {1, 2} && noProblemAtAll(m - mb, i - ib)
11                };
12      } else {
13          return { [m + mb, i + ib, 1]
14                  : mb in {0 .. 3 - m}, ib in {0 .. 3 - i}
15                 | mb + ib in {1, 2} && noProblemAtAll(m + mb, i + ib)
16                };
17      }
18  };
19  start := [3, 3, 1];
20  goal  := [0, 0, 0];

```

Figure 2.1: The missionary and cannibals problem codes as a search problem.

Figure 2.1 shows a formalization of the missionaries and cannibals problem as a search problem. We discuss this formalization line by line.

1. Line 1 defines the auxiliary function **problem**.

If m is the number of missionaries on a given shore, while i is the number of infidels on that same shore, then $\text{problem}(m, i)$ is **true** iff the missionaries have a problem on that shore.

2. Line 3 defines the auxiliary function **noProblemAtAll**.

If m is the number of missionaries on the northern shore and i is the number of infidels on that shore, then the expression $\text{noProblemAtAll}(m, i)$ is true, if there is no problem for the missionaries on either shore.

The implementation of this function uses the fact that if m is the number of missionaries on the northern shore, then $3 - m$ is the number of missionaries on the southern shore. Similarly, if i is the number of infidels on the northern shore, then the number of infidels on the southern shore is $3 - i$.

3. Lines 5 to 18 define the function `nextStates`. A state s is represented as a triple of the form

$$s = [m, i, b] \quad \text{where } m \in \{0, 1, 2, 3\}, i \in \{0, 1, 2, 3\}, b \in \{0, 1\}.$$

Here m is the number of missionaries on the northern shore, i is the number of infidels on the northern shore, and b is the number of boats on the northern shore.

- (a) Line 6 extracts the components m , i , and b from the state s .
- (b) Line 7 checks whether the boat is on the northern shore.
- (c) If this is the case, then the states reachable from the given state s are those states where `mb` missionaries and `ib` infidels cross the river. After `mb` missionaries and `ib` infidels have crossed the river and reached the southern shore, $m - mb$ missionaries and $i - ib$ infidels remain on the northern shore. Of course, after the crossing the boat is no longer on the northern shore. Therefore, the new state has the form

$$[m - mb, i - ib, 0].$$

This explains line 8.

- (d) Since the number `mb` of missionaries leaving the northern shore can not be greater than the number m of all missionaries on the northern shore, we have the condition

$$mb \in \{0, \dots, m\}.$$

There is a similar condition for the number of infidels crossing:

$$ib \in \{0, \dots, i\}.$$

This explains line 9.

- (e) Furthermore, we have to check that the number of persons crossing the river is at least 1 and at most 2. This explains the condition

$$mb + ib \in \{1, 2\}.$$

Finally, there should be no problem in the new state on either shore. This is checked using the expression

$$\text{noProblemAtAll}(m - mb, i - ib).$$

These two checks are performed in line 10.

4. If the boat is on the southern shore instead, then the missionaries and the infidels will be crossing the river from the southern shore to the northern shore. Therefore, the number of missionaries and infidels on the northern shore is now increased. Hence, in this case the new state has the form

$$[m + mb, i + ib, 1].$$

As the number of missionaries on the southern shore is $3 - m$ and the number of infidels on the southern shore is $3 - i$, `mb` is now a member of the set $\{0, \dots, 3 - m\}$, while `ib` is a member of the set $\{0, \dots, 3 - i\}$.

5. Finally the start state and the goal state are defined in line 19 and line 20.

The code in Figure 2.1 does not define the set of states Q of the search problem. The reason is that, in order to solve the problem, we do not need to define this set. If we wanted to, we could define the set of states as follows:

`States := { [m,i,b] : m in {0..3}, i in {0..3}, b in {0,1} | noProblemAtAll(m, i) };`

Figure 2.2 shows a graphical representation of the transition relation of the missionaries and cannibals puzzle. In that figure, for every state both the northern and the southern shore are shown. The start state is covered with a blue ellipse, while the goal state is covered with a green ellipse. The figure clearly shows that the problem is solvable and that there is a solution involving just 11 crossings of the river. \diamond



Figure 2.2: A graphical representation of the missionaries and cannibals problem.

2.1 The Sliding Puzzle

The 3×3 sliding puzzle uses a square board of length 3. This board is subdivided into $3 \times 3 = 9$ squares of length 1. Of these 9 squares, 8 are occupied with square tiles that are numbered from 1 to 8. One square remains empty. Figure 2.3 on page 10 shows two possible states of this sliding puzzle. The 4×4 sliding puzzle is similar to the 3×3 sliding puzzle but it is played on a square board of length 4 instead. The 4×4 sliding puzzle is also known as the 15 puzzle, while the 3×3 puzzle is called the 8 puzzle.

In order to solve the 3×3 sliding puzzle shown in Figure 2.3 we have to transform the state shown on the left of Figure 2.3 into the state shown on the right of this figure. The following operations are permitted when transforming a state of the sliding puzzle:

1. If a tile is to the left of the free square, this tile can be moved to the right.
2. If a tile is to the right of the free square, this tile can be moved to the left.
3. If a tile is above the free square, this tile can be moved down.
4. If a tile is below the free square, this tile can be moved up.

Figure 2.3: The 3×3 sliding puzzle.

In order to get a feeling for the complexity of the sliding puzzle, you can check the page

<http://mypuzzle.org/sliding>.

The sliding puzzle is much more complex than the missionaries and cannibals problem because the state space is much larger. For the case of the 3×3 sliding puzzle, there are 9 squares that can be positioned in $9!$ different ways. It turns out that only half of these positions are reachable from a given start state. Therefore, the effective number of states for the 3×3 sliding puzzle is

$$9!/2 = 181,440.$$

This is already a big number, but 181,440 states can still be stored in a modern computer. However, the 4×4 sliding puzzle has

$$16!/2 = 10,461,394,944,000$$

different states reachable from a given start state. If a state is represented as matrix containing 16 numbers and we store every number using just 4 bits, we still need $16 \cdot 4 = 64$ bits or 8 bytes for every state. Hence we would need a total of

$$(16!/2) \cdot 8 = 83,691,159,552,000$$

bytes to store every state. We would thus need about 84 Terabytes to store the set of all states. As few computers are equipped with this kind of memory, it is obvious that we won't be able to store the entire state space in memory.

Figure 2.4 shows how the 3×3 sliding puzzle can be formulated as a search problem. We discuss this program line by line.

1. `findTile` is an auxiliary procedure that takes a `number` and a `state` and returns the row and column where the tile labelled with `number` can be found.

Here, a state is represented as a list of lists. For example, the states shown in Figure 2.3 are represented as shown in line 26 and line 30. The empty tile is coded as 0.

2. `moveDir` takes a `state`, the `row` and the `column` where to find the empty square and a direction in which the empty square should be moved. This direction is specified via the two variables `dx` and `dy`. The tile at the position $\langle \text{row} + \text{dx}, \text{col} + \text{dy} \rangle$ is moved into the position $\langle \text{row}, \text{col} \rangle$, while the tile at position $\langle \text{row} + \text{dx}, \text{col} + \text{dy} \rangle$ becomes empty.
3. Given a `state`, the procedure `nextStates` computes the set of all states that can be reached in one step from `state`. The basic idea is to find the position of the empty tile and then try to move the empty tile in all possible directions. If the empty tile is found at position $[\text{row}, \text{col}]$ and the direction of the movement is given as $[\text{dx}, \text{dy}]$, then in order to ensure that the empty tile can be moved to the position $[\text{row} + \text{dx}, \text{col} + \text{dy}]$, we have to ensure that both

$$\text{row} + \text{dx} \in \{1, \dots, n\} \quad \text{and} \quad \text{col} + \text{dy} \in \{1, \dots, n\}$$

```

1  findTile := procedure(number, state) {
2      n := #state;
3      L := [1 .. n];
4      for (row in L, col in L | state[row][col] == number) {
5          return [row, col];
6      }
7  };
8  moveDir := procedure(state, row, col, dx, dy) {
9      state[row][col] := state[row + dx][col + dy];
10     state[row + dx][col + dy] := 0;
11     return state;
12 };
13 nextStates := procedure(state) {
14     n := #state;
15     [row, col] := findTile(0, state);
16     newStates := [];
17     directions := [ [1, 0], [-1, 0], [0, 1], [0, -1] ];
18     L := [1 .. n];
19     for ([dx, dy] in directions) {
20         if (row + dx in L && col + dy in L) {
21             newStates += [ moveDir(state, row, col, dx, dy) ];
22         }
23     }
24     return newStates;
25 };
26 start := [ [8, 0, 6],
27            [5, 4, 7],
28            [2, 3, 1]
29 ];
30 goal := [ [0, 1, 2],
31           [3, 4, 5],
32           [6, 7, 8]
33 ];

```

Figure 2.4: The 3×3 sliding puzzle.

hold, where n is the size of the board.

Next, we want to develop an algorithm that can solve puzzles of the kind described so far. The most basic algorithm to solve search problems is **breadth first search**. We discuss this algorithm next.

2.2 Breadth First Search

Informally, breadth first search, abbreviated as BFS, works as follows:

1. Given a search problems $\langle Q, \text{nextStates}, \text{start}, \text{goal} \rangle$, we initialize a set **Frontier** to contain the state **start**.

In general, **Frontier** contains those states that have just been discovered and whose successors have not yet been seen.

2. As long as the set **Frontier** does not contain the state **goal**, we recompute this set by adding all states to it that can be reached in one step from a state in **Frontier**. Then, the states that had been previously present in **Frontier** are removed. These old states are then saved into a set **Visited**.

In order to avoid loops, an implementation of breadth first search keeps track of those states that have been visited. These states are collected in a set **Visited**. Once a state has been added to the set **Visited**, it will never be revisited again. Furthermore, in order to keep track of the path leading to the goal, we have a dictionary **Parent**. For every state s that is in **Frontier**, **Parent**[s] is the state that caused s to be added to the set **Frontier**, i.e. we have

$$s \in \text{nextStates}(\text{Parent}[s]).$$

```

1  search := procedure(start, goal, nextStates) {
2      Frontier := { start };
3      Visited  := {}; // set of nodes that have been expanded
4      Parent   := {};
5      while (Frontier != {}) {
6          NewFrontier := {};
7          for ( s in Frontier, ns in nextStates(s)
8              | !(ns in Visited) && !(ns in Frontier)
9              )
10             {
11                 NewFrontier += { ns };
12                 Parent[ns]  := s;
13                 if (ns == goal) {
14                     return pathTo(goal, Parent);
15                 }
16             }
17             Visited += Frontier;
18             Frontier := NewFrontier;
19         }
20     };

```

Figure 2.5: Breadth first search.

Figure 2.5 on page 12 shows an implementation of breadth first search in SETLX. We discuss this implementation line by line:

1. **Frontier** is the set of all those states that have been encountered but whose neighbours have not yet been explored. Initially, it contains the state **start**.
2. **Visited** is the set of all those states, all whose neighbours have already been added to the set **Frontier**. In order to avoid infinite loops, these states must not be visited again.
3. **Parent** is a dictionary keeping track of the state leading to a given state.
4. As long as the set **Frontier** is not empty, we add all neighbours of states in **Frontier** that have not yet been visited to the set **NewFrontier**. When doing this, we keep track of the path leading to a new state **ns** by storing its parent in the dictionary **Parent**.
5. If the new state happens to be the state **goal**, we return a path leading from **start** to **goal**. The procedure **pathTo()** is shown in Figure 2.6 on page 13.
6. After we have collected all successors of states in **Frontier**, the states in the set **Frontier** have been visited and are therefore added to the set **Visited**, while the **Frontier** is updated to **NewFrontier**.

The procedure call **pathTo(state, Parent)** constructs a path reaching from **start** to **state** in reverse by looking up the parent states.

If we try breadth first search to solve the missionaries and cannibals problem, we immediately get the solution shown in Figure 2.7. 15 nodes had to be expanded to find this solution. To keep this in perspective,

```

1  pathTo := procedure(state, Parent) {
2      Path := [];
3      while (state != om) {
4          Path += [state];
5          state := Parent[state];
6      }
7      return reverse(Path);
8  };

```

Figure 2.6: The procedure `pathTo()`.

we note that Figure 2.2 shows that the entire state space contains 16 states. Therefore, with the exception of one state, we have inspected all the states. This is a typical behaviour for breadth first search.

1	MMM	KKK	B	~~~~~			
2				> KK >			
3	MMM	K		~~~~~		KK	B
4				< K <			
5	MMM	KK	B	~~~~~		K	
6				> KK >			
7	MMM			~~~~~		KKK	B
8				< K <			
9	MMM	K	B	~~~~~		KK	
10				> MM >			
11	M	K		~~~~~	MM	KK	B
12				< M K <			
13	MM	KK	B	~~~~~	M	K	
14				> MM >			
15		KK		~~~~~	MMM	K	B
16				< K <			
17		KKK	B	~~~~~	MMM		
18				> KK >			
19		K		~~~~~	MMM	KK	B
20				< K <			
21		KK	B	~~~~~	MMM	K	
22				> KK >			
23				~~~~~	MMM	KKK	B

Figure 2.7: A solution of the missionaries and cannibals problem.

Next, let us try to solve the 3×3 sliding puzzle. It takes about 6 seconds to solve this problem on my computer¹, while 181439 states are touched. Again, we see that breadth first search touches nearly all the states reachable from the start state.

2.2.1 A Queue Based Implementation of Breadth First Search

In the literature, for example in Figure 3.11 of Russell & Norvig [RN09], breadth first search is often implemented using a **queue** data structure. Figure 2.8 on page 14 shows an implementation of breadth first search that uses

¹ I happen to own an iMac from 2017. This iMac is equipped with 32 Gigabytes of main memory and a quad core 3.4 GHz “Intel Core i5” processor. I suspect this to be the I5-7500 (Kaby Lake) processor.

a queue to store the set **Frontier**. However, when we run this version, it turns out that the solution of the 3×3 sliding puzzle needs about 58 seconds, which is a lot slower than our set based implementation that has been presented in Figure 2.5.

```

1  search := procedure(start, goal, nextStates) {
2      Queue := [ start ];
3      Visited := {};
4      Parent := {};
5      while (Queue != []) {
6          state := Queue[1];
7          Queue := Queue[2..];
8          if (state == goal) {
9              return pathTo(state, Parent);
10         }
11         Visited += { state };
12         newStates := nextStates(state);
13         for (ns in newStates | !(ns in Visited) && Parent[ns] == om) {
14             Parent[ns] := state;
15             Queue += [ ns ];
16         }
17     }
18 };

```

Figure 2.8: A queue based implementation of breadth first search.

The solution of the 3×3 sliding puzzle that is found by breadth first search is shown in Figure 2.9 and Figure 2.10.

We conclude our discussion of breadth first search by noting the two most important properties of breadth first search.

1. Breadth first search is **complete**: If there is a solution to the given search problem, then breadth first search is going to find it.
2. The solution found by breadth first search is **optimal**, i.e. it is one of the shortest possible solutions.

Proof: Both of these claims can be shown simultaneously. Consider the implementation of breadth first search shown in Figure 2.5. An easy induction on the number of iterations of the **while** loop shows that after n iterations of the **while** loop, the set **Frontier** contains exactly those states that have a distance of n to the state **start**. This claim is obviously true before the first iteration of the while loop as in this case, **Frontier** only contains the state **start**. In the induction step we assume the claim is true after n iterations. Then, in the next iteration all states that can be reached in one step from a state in **Frontier** are added to the new **Frontier**, provided there is no shorter path to these states. There is a shorter path to these states if these states are already a member of the set **Visited**. Hence, the claim is true after $n + 1$ iterations also.

Now, if there is a path from **start** to **goal**, there must also be a shortest path. Assume this path has a length of k . Then, **goal** is reached in the iteration number k and the shortest path is returned. \square

The fact that breadth first search is both complete and the path returned is optimal is rather satisfying. However, breadth first search still has a big downside that makes it unusable for many problems: If the **goal** is far from the **start**, breadth first search will use a lot of memory because it will store a large part of the state space in the set **Visited**. In many cases, the state space is so big that this is not possible. For example, it is impossible to solve the more interesting cases of the 4×4 sliding puzzle.

1	+---+---+---+		+---+---+---+		+---+---+---+		+---+---+---+
2	8 6		8 6		5 8 6		5 8 6
3	+---+---+---+		+---+---+---+		+---+---+---+		+---+---+---+
4	5 4 7	==>	5 4 7	==>	4 7	==>	2 4 7
5	+---+---+---+		+---+---+---+		+---+---+---+		+---+---+---+
6	2 3 1		2 3 1		2 3 1		3 1
7	+---+---+---+		+---+---+---+		+---+---+---+		+---+---+---+
8							
9	+---+---+---+		+---+---+---+		+---+---+---+		+---+---+---+
10	5 8 6		5 8 6		5 8 6		5 8
11	+---+---+---+		+---+---+---+		+---+---+---+		+---+---+---+
12	2 4 7	==>	2 4 7	==>	2 4	==>	2 4 6
13	+---+---+---+		+---+---+---+		+---+---+---+		+---+---+---+
14	3 1		3 1		3 1 7		3 1 7
15	+---+---+---+		+---+---+---+		+---+---+---+		+---+---+---+
16							
17	+---+---+---+		+---+---+---+		+---+---+---+		+---+---+---+
18	5 8		5 8		2 5 8		2 5 8
19	+---+---+---+		+---+---+---+		+---+---+---+		+---+---+---+
20	2 4 6	==>	2 4 6	==>	4 6	==>	4 6
21	+---+---+---+		+---+---+---+		+---+---+---+		+---+---+---+
22	3 1 7		3 1 7		3 1 7		3 1 7
23	+---+---+---+		+---+---+---+		+---+---+---+		+---+---+---+
24							
25	+---+---+---+		+---+---+---+		+---+---+---+		+---+---+---+
26	2 5 8		2 5 8		2 5 8		2 5
27	+---+---+---+		+---+---+---+		+---+---+---+		+---+---+---+
28	4 1 6	==>	4 1 6	==>	4 1	==>	4 1 8
29	+---+---+---+		+---+---+---+		+---+---+---+		+---+---+---+
30	3 7		3 7		3 7 6		3 7 6
31	+---+---+---+		+---+---+---+		+---+---+---+		+---+---+---+
32							
33	+---+---+---+		+---+---+---+		+---+---+---+		+---+---+---+
34	2 5		2 5		4 2 5		4 2 5
35	+---+---+---+		+---+---+---+		+---+---+---+		+---+---+---+
36	4 1 8	==>	4 1 8	==>	1 8	==>	1 8
37	+---+---+---+		+---+---+---+		+---+---+---+		+---+---+---+
38	3 7 6		3 7 6		3 7 6		3 7 6
39	+---+---+---+		+---+---+---+		+---+---+---+		+---+---+---+
40							
41	+---+---+---+		+---+---+---+		+---+---+---+		+---+---+---+
42	4 2 5		4 2 5		4 2 5		4 2
43	+---+---+---+		+---+---+---+		+---+---+---+		+---+---+---+
44	1 7 8	==>	1 7 8	==>	1 7	==>	1 7 5
45	+---+---+---+		+---+---+---+		+---+---+---+		+---+---+---+
46	3 6		3 6		3 6 8		3 6 8
47	+---+---+---+		+---+---+---+		+---+---+---+		+---+---+---+

Figure 2.9: The first 24 steps in the solution of the 3×3 sliding puzzle.

Figure 2.10: The last 7 steps in the solution of the 3×3 sliding puzzle.

2.3 Depth First Search

To overcome the memory limitations of breadth first search, the **depth first search** algorithm has been developed. The basic idea is to replace the queue of Figure 2.8 by a stack. The resulting algorithm is shown in Figure 2.11 on page 16. Basically, in this implementation, a path is searched to its end before trying an alternative. This way, we might be able to find a **goal** that is far away from **start** without exploring the whole state space.

```

1  search := procedure(start, goal, nextStates) {
2      Stack := [ start ];
3      Parent := {};
4      while (Stack != []) {
5          state := Stack[-1];
6          Stack := Stack[..-2];
7          if (state == goal) {
8              return pathTo(state, Parent);
9          }
10         newStates := nextStates(state);
11         for (ns in newStates | ns != start && Parent[ns] == om) {
12             Parent[ns] := state;
13             Stack      += [ns];
14         }
15     }
16 };

```

Figure 2.11: The depth first search algorithm.

Actually, it is not necessary to understand the details of the implementation shown in Figure 2.11 on page 16. The reason is that the recursive implementation of depth first search that is presented in the following subsection is superior to the implementation shown in Figure 2.11 on page 16. When we test the implementation shown above with the 3×3 sliding puzzle, it takes about 68 seconds to find a solution. The solution that is found has a length of 41,553 steps. As the shortest path from **start** to **goal** has 31 steps, the solution found by depth first search is very far from optimal. All this is rather disappointing news. The only good news is that there is no longer a need to keep the set **Visited** around. However, we still have to maintain the set **Parent**. If we

were more ambitious, we could eliminate the use of this dictionary also, but the resulting implementation would be rather unwieldy. Fortunately, we will be able to get rid of the set `Parent` with next to no effort when we develop a recursive implementation of depth first search in the following subsection.

2.3.1 Getting Rid of the Parent Dictionary

It can be argued that the implementation of depth first search discussed previously is not really depth first search because it uses the dictionary `Parent`. As states are only added to `Parent` and never removed, at the end of the search this dictionary will contain all states that have been visited. This defeats the most important advantage of depth first search which is the fact that it should only store the current path that is investigated. Therefore, it has been suggested (for example compare Russel and Norvig [RN09]) that instead of storing single states, the stack should store the full paths leading to these states. This leads to the implementation shown in Figure 2.12 on page 17.

```

1  search := procedure(start, goal, nextStates) {
2      Stack := [ [start] ];
3      while (Stack != []) {
4          Path := Stack[-1];
5          Stack := Stack[..-2];
6          state := Path[-1];
7          if (state == goal) {
8              return Path;
9          }
10         newStates := nextStates(state);
11         for (ns in newStates | !(ns in Path)) {
12             Stack += [ Path + [ns] ];
13         }
14     }
15 };

```

Figure 2.12: An path-based implementation of depth first search.

Unfortunately, it turns out that the paths get very long and hence need a lot of memory to be stored and this fact defeats the main idea of this implementation. As a result, the procedure `search` that is given in Figure 2.12 on page 17 is not able to solve the instance of the 3×3 sliding puzzle that was shown in Figure 2.3 on page 10.

Exercise 1: Assume that n is a positive natural number and that the set of states Q_n is defined as

$$Q_n := \{ \langle a, b \rangle \mid a \in \{0, \dots, n\} \wedge b \in \{0, \dots, n\} \}.$$

Furthermore, the states `start` and `goal` are defined as

$$\text{start} := \langle 0, 0 \rangle \quad \text{and} \quad \text{goal} := \langle n, n \rangle.$$

Next, the function `nextStates` is defined as

$$\text{nextStates}(\langle a, b \rangle) := \begin{cases} \{ \langle a+1, b \rangle, \langle a, b+1 \rangle \} & \text{if } a < n \text{ and } b < n \\ \{ \langle a+1, b \rangle, \langle a+1, 0 \rangle \} & \text{if } a < n \text{ and } b = n \\ \{ \langle 0, b+1 \rangle, \langle a, b+1 \rangle \} & \text{if } a = n \text{ and } b < n \\ \{ \} & \text{if } a = n \text{ and } b = n \end{cases}$$

Finally, the search problem \mathcal{P} is defined as

$$\mathcal{P} := \langle Q, \text{nextStates}, \text{start}, \text{goal} \rangle.$$

Assume that states can be stored using 8 bytes. Furthermore, assume that we use the algorithm given in Figure

2.12 on page 17 to solve the search problem \mathcal{P} . How many bytes do we need to store all the states on the stack in the moment that the `goal` is reached? How big is this number if $n = 10,000$?

- (a) **Note** that the question only asks for the memory needed to store the states. The memory needed to store the stack itself and the various lists on the stack comes on top of the memory needed to store the states. However, to answer this exercise correctly, you should ignore this type of memory.
- (b) In order to understand how the stack evolves, we need to know that in SETLX sets are ordered ascendingly. Furthermore, pairs are ordered lexicographically in SETLX, i.e. we have

$$\langle x_1, y_1 \rangle < \langle x_2, y_2 \rangle \iff x_1 < x_2 \vee (x_1 = x_2 \wedge y_1 < y_2).$$

Hence, when we have a state $\langle a, b \rangle$ such that both $a < n$ and $b < n$, the set `nextStates`($\langle a, b \rangle$) is ordered as follows:

$$\{\langle a, b+1 \rangle, \langle a+1, b \rangle\}.$$

◇

2.3.2 A Recursive Implementation of Depth First Search

Sometimes, the depth first search algorithm is presented as a recursive algorithm, since this leads to an implementation that is slightly shorter and is easier to understand. What is more, we no longer need the dictionary `Parent` to record the parent of each node. The resulting implementation is shown in Figure 2.13 on page 18.

```

1  search := procedure(start, goal, nextStates) {
2      return dfs(start, goal, nextStates, [start]);
3  };
4  dfs := procedure(state, goal, nextStates, Path) {
5      if (state == goal) {
6          return Path;
7      }
8      newStates := nextStates(state);
9      for (ns in newStates | !(ns in Path)) {
10         result := dfs(ns, goal, nextStates, Path + [ns]);
11         if (result != om) {
12             return result;
13         }
14     }
15 };

```

Figure 2.13: A recursive implementation of depth first search.

The only purpose of the procedure `search` is to call the procedure `dfs`, which needs one additional argument. This argument is called `Path`. The idea is that `Path` is a path leading from the state `start` to the current `state` that is the first argument of the procedure `dfs`. On the first invocation of `dfs`, the parameter `state` is equal to `start` and therefore `Path` is initialized as the list containing only `start`.

The implementation of `dfs` works as follows:

1. If `state` is equal to `goal`, our search is successful. Since by assumption the list `Path` is a path connecting `start` and `state` and we have checked that `state` is equal to `goal`, we can return `Path` as our solution.
2. Otherwise, `newStates` is the set of states that are reachable from `state` in one step. Any of the states `ns` in this set could be the next state on a path that leads to `goal`. Therefore, we try recursively to reach `goal` from every state `ns`. Note that we have to change `Path` to the list

`Path + [ns]`

when we call the procedure `dfs` recursively. This way, we retain the invariant of `dfs` that the list `Path` is a path connecting `start` with `state`.

3. We still have to avoid running in circles. In the recursive version of depth first search, this is achieved by checking that the state `ns` is not already a member of the list `Path`. In the non-recursive version of depth first search, we had used the set `Parent` instead. The current implementation no longer has a need for the dictionary `Parent`. This is very fortunate since it reduces the memory requirements of depth first search considerably.
4. If one of the recursive calls of `dfs` returns a list, this list is a solution to our search problem and hence it is returned. However, if instead the undefined value `om` is returned, the `for` loop needs to carry on and test the other successors of `state`.
5. Note that the recursive invocation of `dfs` returns `om` if the end of the `for` loop is reached and no solution has been returned so far. The reason is that there is no `return` statement at the end of the procedure `dfs`. Hence, if the last line of the procedure `dfs` is reached, `om` is returned by default.

For the 3×3 puzzle, it takes about one second to compute the solution. In this case, the length of the solution is still 3653 steps, which is unsatisfying. The good news is that this program does not need much memory. The only variable that uses considerable memory is the variable `Path`. If we can somehow keep the list `Path` short, then the recursive version of depth first search uses only a tiny fraction of the memory needed by breadth first search.

2.4 Iterative Deepening

The fact that the recursive version of depth first search took just one second to find a solution is very impressive. The question is whether it might be possible to force depth first search to find the shortest solution. The answer to this question leads to an algorithm that is known as **iterative deepening**. The main idea behind iterative deepening is to run depth first with a **depth limit** d . This limit enforces that a solution has at most a length of d . If no solution is found at a depth of d , the new depth $d + 1$ can be tried next and the process can be continued until a solution is found. The program shown in Figure 2.14 on page 20 implements this strategy. We proceed to discuss the details of this program.

1. The procedure `search` initializes the variable `limit` to 1 and tries to find a solution to the search problem that has a length that is less than or equal to `limit`. If a solution is found, it is returned. Otherwise, the variable `limit` is incremented by one and a new instance of depth first search is started. This process continues until either
 - a solution is found or
 - the sun rises in the west.
2. The procedure `depthLimitedSearch` implements depth first search but takes care to compute only those paths that have a length of at most `limit`. The implementation shown in Figure 2.14 is stack based. In this implementation, the stack contains paths leading from `start` to the state at the end of a given path. Hence it is similar to the implementation of depth first search shown in Figure 2.11 on page 16.
3. The stack is initialized to contain the path [`start`].
4. In the `while`-loop, the first thing that happens is that the `Path` on top of the stack is removed from the stack. The state at the end of this `Path` is called `state`. If this `state` happens to be the `goal`, a solution to the search problem has been found and this solution is returned.
5. Otherwise, we check the length of `Path`. If this length is greater than or equal to the `limit`, the `Path` can be discarded as we have already checked that it does not end in the `goal`.
6. Otherwise, the neighbours of `state` are computed. For every neighbour `ns` of `state` that has not yet been encountered in `Path`, we extend `Path` to a new list that ends in `ns`.

```

1  search := procedure(start, goal, nextStates) {
2      limit := 1;
3      while (true) {
4          Path := depthLimitedSearch(start, goal, nextStates, limit);
5          if (Path != om) { return Path; }
6          limit += 1;
7      }
8  };
9  depthLimitedSearch := procedure(start, goal, nextStates, limit) {
10     Stack := [ [start] ];
11     while (Stack != []) {
12         Path := Stack[-1];
13         Stack := Stack[..-2];
14         state := Path[-1];
15         if (state == goal) { return Path; }
16         if (#Path >= limit) { continue; }
17         for (ns in nextStates(state) | !(ns in Path)) {
18             Stack += [ Path + [ns] ];
19         }
20     }
21 };

```

Figure 2.14: Iterative deepening implemented in SETLX.

7. This process is iterated until the **Stack** is exhausted.

The nice thing about the program presented in this section is the fact that it does not use much memory. The reason is that the stack can never have a size that is longer than **limit** and therefore the overall memory that is needed can be bounded by $\mathcal{O}(\text{limit}^2)$. However, when we run this program to solve the 3×3 sliding puzzle, the algorithm takes about 25 minutes. There are two reasons for this:

1. First, it is quite wasteful to run the search for a depth limit of 1, 2, 3, \dots all the way up to 31. Essentially, all the computations done with a limit less than 31 are essentially wasted.
2. Given a state s that is reachable from the **start**, there often is a huge number of different paths that lead from **start** to s . The version of iterative deepening presented in this section tries all of these paths and hence needs a large amount of time.

Exercise 2: Assume the set of states Q is defined as

$$Q := \{ \langle a, b \rangle \mid a \in \mathbb{N} \wedge b \in \mathbb{N} \}.$$

Furthermore, the states **start** and **goal** are defined as

$$\text{start} := \langle 0, 0 \rangle \quad \text{and} \quad \text{goal} := \langle n, n \rangle \text{ where } n \in \mathbb{N}.$$

Next, the function **nextStates** is defined as

$$\text{nextStates}(\langle a, b \rangle) := \{ \langle a + 1, b \rangle, \langle a, b + 1 \rangle \}.$$

Finally, the search problem \mathcal{P} is defined as

$$\mathcal{P} := \langle Q, \text{nextStates}, \text{start}, \text{goal} \rangle.$$

Given a natural number n , compute the number of different solutions of this search problem and prove your claim. \diamond

Exercise 3: If there is no solution, the implementation of iterative deepening that is shown in Figure 2.14 does not terminate. The reason is that the function `depthLimitedSearch` does not distinguish between the following two reasons for failure:

1. It can fail to find a solution because the depth limit is reached.
2. It can also fail because it has tried all paths without hitting the depth limit but the **Stack** is exhausted.

Improve the implementation of iterative deepening so that it will always terminate eventually, provided the state space is finite. \diamond

2.4.1 A Recursive Implementation of Iterative Deepening

If we implement iterative deepening recursively, then we know that the call stack is bounded by the length of the shortest solution. Figure 2.15 on page 21 shows a recursive implementation of iterative deepening. Unfortunately, the running time of the recursive implementation of iterative deepening is still quite big: On my computer, the recursive implementation takes about 23 minutes.

```

1  search := procedure(start, goal, nextStates) {
2      limit := 1;
3      while (true) {
4          result := dfsLimited(start, goal, nextStates, [start], limit);
5          if (result != om) {
6              return result;
7          }
8          limit += 1;
9      }
10 };
11 dfsLimited := procedure(state, goal, nextStates, Path, limit) {
12     if (state == goal) {
13         return Path;
14     }
15     if (limit == 0) {
16         return; // limit exceeded
17     }
18     for (ns in nextStates(state) | !(ns in Path)) {
19         result := dfsLimited(ns, goal, nextStates, Path + [ns], limit - 1);
20         if (result != om) {
21             return result;
22         }
23     }
24 };

```

Figure 2.15: A recursive implementation of iterative deepening.

2.5 Bidirectional Breadth First Search

Breadth first search first visits all states that have a distance of 1 from start, then all states that have a distance of 2, then of 3 and so on until finally the goal is found. If the length of the shortest path from **start** to **goal** is d , then all states that have a distance of at most d will be visited. In many search problems, the number of states grows exponentially with the distance, i.e. there is a **branching factor** b such that the set of all states that

have a distance of at most d from **start** is roughly

$$1 + b + b^2 + b^3 + \dots + b^d = \frac{b^{d+1} - 1}{b - 1} = \mathcal{O}(b^d).$$

At least this is true in the beginning of the search. As the size of the memory that is needed is the most constraining factor when searching, it is important to cut down this size. One simple idea is to start searching both from the node **start** and the node **goal** simultaneously. The justification is that we can hope that the path starting from **start** and the path starting from **goal** will meet in the middle and hence they will both have a size of approximately $d/2$. If this is the case, only

$$2 \cdot \frac{b^{d/2} - 1}{b - 1}$$

nodes need to be explored and even for modest values of b this number is much smaller than

$$\frac{b^{d+1} - 1}{b - 1}$$

which is the number of nodes expanded in breadth first search. For example, assume that the branching factor $b = 2$ and that the length of the shortest path leading from **start** to **goal** is 40. Then we need to explore

$$2^{40} - 1 = 1,099,511,627,775$$

in breadth first search, while we only have to explore

$$2 \cdot (2^{40/2} - 1) = 2,097,150$$

with bidirectional breadth first search. While it is certainly feasible to keep two million states in memory, keeping a trillion states in memory is impossible on most devices.

Figure 2.16 on page 23 shows the implementation of bidirectional breadth first search. Essentially, we have to copy the breadth first program shown in Figure 2.5. Let us discuss the details of the implementation.

1. The variable **FrontierA** is the frontier that starts from the state **start**, while **FrontierB** is the frontier that starts from the state **goal**.
2. **VisitedA** is the set of states that have been visited starting from **start**, while **VisitedB** is the set of states that have been visited starting from **goal**.
3. For every state s that is in **FrontierA**, **ParentA**[s] is the state that caused s to be added to the set **FrontierA**. Similarly, for every state s that is in **FrontierB**, **ParentB**[s] is the state that caused s to be added to the set **FrontierB**.
4. The bidirectional search keeps running for as long as both sets **FrontierA** and **FrontierB** are non-empty and a path has not yet been found.
5. Initially, the **while** loop adds the frontier sets to the visited sets as all the neighbours of the frontier sets will now be explored.
6. Then the **while** loop computes those states that can be reached from **FrontierA** and have not been visited from **start**. If a state **ns** is a neighbour of a state **s** from the set **FrontierA** and the state **ns** has already been encountered during the search that started from **goal**, then a path leading from **start** to **goal** has been found and this path is returned. The function **combinePaths** that computes this path by combining the path that leads from **start** to **ns** and then from **ns** to **goal** to is shown in Figure 2.17 on page 23.
7. Next, the same computation is done with the role of the states **start** and **goal** exchanged.

On my computer, bidirectional breadth first search solves the 3×3 sliding puzzle in less than a second! However, bidirectional breadth first search is still not able to solve the 4×4 sliding puzzle since the portion of the search space that needs to be computed is just too big to fit into memory.

```

1  search := procedure(start, goal, nextStates) {
2      FrontierA := { start };
3      VisitedA  := {}; // set of nodes expanded starting from start
4      ParentA   := {};
5      FrontierB := { goal };
6      VisitedB  := {}; // set of nodes expanded starting from goal
7      ParentB   := {};
8      while (FrontierA != {} && FrontierB != {}) {
9          VisitedA += FrontierA;
10         VisitedB += FrontierB;
11         NewFrontier := {};
12         for (s in FrontierA, ns in nextStates(s) | !(ns in VisitedA)) {
13             NewFrontier += { ns };
14             ParentA[ns] := s;
15             if (ns in VisitedB) {
16                 return combinePaths(ns, ParentA, ParentB);
17             }
18         }
19         FrontierA := NewFrontier;
20         NewFrontier := {};
21         for (s in FrontierB, ns in nextStates(s) | !(ns in VisitedB)) {
22             NewFrontier += { ns };
23             ParentB[ns] := s;
24             if (ns in VisitedA) {
25                 return combinePaths(ns, ParentA, ParentB);
26             }
27         }
28         FrontierB := NewFrontier;
29     }
30 };

```

Figure 2.16: Bidirectional breadth first search.

```

1  combinePaths := procedure(node, ParentA, ParentB) {
2      Path1 := pathTo(node, ParentA);
3      Path2 := pathTo(node, ParentB);
4      return Path1[..-2] + reverse(Path2);
5  };

```

Figure 2.17: Combining two paths.

2.6 Best First Search

Up to now, all the search algorithms we have discussed were essentially blind. Given a state s and all of its neighbours, they had no idea which of the neighbours they should pick because they had no conception which of these neighbours might be more promising than the other neighbours. If a human tries to solve a problem, she usually will develop a feeling that certain states are more favourable than other states because they seem to be closer to the solution. In order to formalise this procedure, we next define the notion of a [heuristic](#).

Definition 2 (Heuristic) Given a search problem

$$\mathcal{P} = \langle Q, \text{nextStates}, \text{start}, \text{goal} \rangle,$$

a **heuristic** is a function

$$h : Q \rightarrow \mathbb{R}$$

that computes an approximation of the distance of a given state s to the goal state goal . The heuristic is **admissible** if it always underestimates the true distance, i.e. if the function

$$d : Q \rightarrow \mathbb{R}$$

computes the true distance of a state s to the goal, then we must have

$$h(s) \leq d(s) \quad \text{for all } s \in Q.$$

Hence, the heuristic is admissible iff it is **optimistic**: An admissible heuristic must never overestimate the distance to the goal, but it is free to underestimate this distance.

Finally, the heuristic h is called **consistent** iff we have

$$h(\text{goal}) = 0 \quad \text{and} \quad h(s_1) \leq 1 + h(s_2) \quad \text{for all } s_2 \in \text{nextStates}(s_1). \quad \diamond$$

Let us explain the idea behind the notion of consistency. First, if we are already at the goal, the heuristic should notice this and hence return $h(\text{goal}) = 0$. Secondly, assume we are at the state s_1 and s_2 is a neighbour of s_1 , i.e. we have that

$$s_2 \in \text{nextStates}(s_1).$$

Now if our heuristic h assumes that the distance of s_2 from the **goal** is $h(s_2)$, then the distance of s_1 from the **goal** can be at most $1 + h(s_2)$ because starting from s_1 we can first go to s_2 in one step and then from s_2 to **goal** in $h(s_2)$ steps for a total of $1 + h(s_2)$ steps. Of course, it is possible that there exists a cheaper path from s_1 leading to the **goal** than the one that visits s_2 first. Hence we have the inequality

$$h(s_1) \leq 1 + h(s_2).$$

Theorem 3 Every consistent heuristic is also admissible.

Proof: Assume that the heuristic h is consistent. Assume further that $s \in Q$ is some state such that there is a path p from s to the **goal**. Assume this path has the form

$$p = [s_n, s_{n-1}, \dots, s_1, s_0], \quad \text{where } s_n = s \text{ and } s_0 = \text{goal}.$$

Then the length of p is n and we have to show that $h(s) \leq n$. In order to prove this claim, we show that we have

$$h(s_k) \leq k \quad \text{for all } k \in \{0, 1, \dots, n\}.$$

This claim is shown by induction on k .

B.C.: $k = 0$.

We have $h(s_0) = h(\text{goal}) = 0 \leq 0$ because the fact that h is consistent implies $h(\text{goal}) = 0$.

I.S.: $k \mapsto k + 1$.

We have to show that $h(s_{k+1}) \leq k + 1$ holds. This is shown as follows:

$$\begin{aligned} h(s_{k+1}) &\leq 1 + h(s_k) && \text{because } s_k \in \text{nextStates}(s_{k+1}) \text{ and } h \text{ is consistent} \\ &\leq 1 + k && \text{because } h(s_k) \leq k \text{ by induction hypotheses} \end{aligned}$$

This concludes the proof. \square

It is natural to ask whether the last theorem can be reversed, i.e. whether every admissible heuristic is also consistent. The answer to this question is negative since there are **some contorted** heuristics that are

admissible but that fail to be consistent. However, in practice it turns out that most admissible heuristics are also consistent. Therefore, when we construct consistent heuristics later, we will start with admissible heuristics, since these are easy to find. We will then have to check that these heuristics are also consistent.

Examples: In the following, we will discuss several heuristics for the sliding puzzle.

1. The simplest heuristic that is admissible is the function $h(s) := 0$. Since we have

$$0 \leq 1 + 0,$$

this heuristic is obviously consistent, but this heuristic is too trivial to be of any use.

2. The next heuristic is the **number of misplaced tiles** heuristic. For a state s , this heuristic counts the number of tiles in s that are not in their final position, i.e. that are not in the same position as the corresponding tile in **goal**. For example, in Figure 2.3 on page 10 in the state depicted to the left, only the tile with the label 4 is in the same position as in the state depicted to the right. Hence, there are 7 misplaced tiles.

As every misplaced tile must be moved at least once and every step in the sliding puzzle moves at most one tile, it is obvious that this heuristic is admissible. It is also consistent. First, the **goal** has no misplaced tiles, hence its heuristic is 0. Second, in every step of the sliding puzzle only one tile is moved. Therefore the number of misplaced tiles in two neighbouring state can differ by at most one.

Unfortunately, the number of misplaced tiles heuristic is very crude and therefore not particularly useful.

3. The **Manhattan heuristic** improves on the previous heuristic. For two points $\langle x_1, y_1 \rangle, \langle x_2, y_2 \rangle \in \mathbb{R}^2$ the **Manhattan distance** of these points is defined as

$$d_1(\langle x_1, y_1 \rangle, \langle x_2, y_2 \rangle) := |x_1 - x_2| + |y_1 - y_2|.$$

If we associate **Cartesian coordinates** with the tiles of the sliding puzzle such that the tile in the upper left corner has coordinates $\langle 1, 1 \rangle$ and the coordinates of the tile in the lower right corner is $\langle 3, 3 \rangle$, then the Manhattan distance of two positions measures how many steps it takes to move a tile from the first position to the second position if we are allowed to move the tile horizontally or vertically regardless of the fact that the intermediate positions might be blocked by other tiles. To compute the Manhattan heuristic for a state s with respect to the **goal**, we first define the position $\text{pos}(t, s)$ for all tiles $t \in \{1, \dots, 8\}$ in a given state s as follows:

$$\text{pos}(t, s) = \langle \text{row}, \text{col} \rangle \stackrel{\text{def}}{\iff} s[\text{row}][\text{col}] = t,$$

i.e. given a state s , the expression $\text{pos}(t, s)$ computes the Cartesian coordinates of the tile t with respect to s . Then we can define the Manhattan heuristic h for the 3×3 puzzle as follows:

$$h(s) := \sum_{t=1}^8 d_1(\text{pos}(t, s), \text{pos}(t, \text{goal})).$$

The Manhattan heuristic measure the number of moves that would be needed if we wanted to put every tile of s into its final positions and if we were allowed to slide tiles over each other. Figure 2.18 on page 26 shows how the Manhattan distance can be computed. The code given in that figure works for a general $n \times n$ sliding puzzle. It takes two states **stateA** and **stateB** and computes the Manhattan distance between these states.

- (a) First, the size **n** of the puzzle is computed by checking the number of rows of **stateA**.
- (b) Next, the **for** loop iterates over all rows and columns of **stateA** that do not contain a blank tile. Remember that the blank tile is coded using the number 0. The tile at position $\langle \text{rowA}, \text{colA} \rangle$ in **stateA** is computed using the expression **stateA**[**rowA**][**colA**] and the corresponding position $\langle \text{rowB}, \text{colB} \rangle$ of this tile in state **stateB** is computed using the function **findTile**.
- (c) Finally, the Manhattan distance between the two positions $\langle \text{rowA}, \text{colA} \rangle$ and $\langle \text{rowB}, \text{colB} \rangle$ is added to the **result**.

```

1  manhattan := procedure(stateA, stateB) {
2      n := #stateA;
3      L := [1 .. n];
4      result := 0;
5      for (rowA in L, colA in L | stateA[rowA][colA] != 0) {
6          [rowB, colB] := findTile(stateA[rowA][colA], stateB);
7          result += abs(rowA - rowB) + abs(colA - colB);
8      }
9      return result;
10 };

```

Figure 2.18: The Manhattan distance between two states.

The Manhattan distance is admissible. The reason is that if $s_2 \in \text{nextStates}(s_1)$, then there can be only one tile t such that the position of t in s_1 is different from the position of t in s_2 . Furthermore, this position differs by either one row or one column. Therefore,

$$|h(s_1) - h(s_2)| = 1$$

and hence $h(s_1) \leq 1 + h(s_2)$. □

Now we are ready to present **best first search**. This algorithm is derived from the stack based version of depth first search. However, instead of using a stack, the algorithm uses a **priority queue**. In this priority queue, the paths are ordered with respect to the estimated distance of the state at the end of the path from the **goal**. We always expand the path next that seems to be closest to the goal.

```

1  bestFirstSearch := procedure(start, goal, nextStates, heuristic) {
2      PrioQueue := { [heuristic(start, goal), [start]] };
3      while (PrioQueue != {}) {
4          [_, Path] := fromB(PrioQueue);
5          state := Path[-1];
6          if (state == goal) { return Path; }
7          newStates := nextStates(state);
8          for (ns in newStates | !(ns in Path)) {
9              PrioQueue += { [heuristic(ns, goal), Path + [ns]] };
10         }
11     }
12 };

```

Figure 2.19: The best first search algorithm.

The procedure **bestFirstSearch** shown in Figure 2.19 on page 26 takes four parameters. The first three of these parameters are the same as in the previous search algorithm. The last parameter **heuristic** is a function that takes to states and then estimates the distance between these states. Later, we will use the Manhattan distance to serve as this **heuristic**. The details of the implementation are as follows:

1. The variable **PrioQueue** serves as a priority queue. We take advantage of the fact that SETLX stores sets as ordered binary trees that store their elements in increasing order. Hence, the smallest element of a set is the first element.

Furthermore, SETLX orders pairs lexicographically. Hence, we store the paths in **PrioQueue** as pairs of the form

$\langle \text{estimate}, \text{Path} \rangle$.

Here **Path** is a list of states starting from the node **start**. If the last node on this list is called **state**, then we have

estimate = **heuristic**(**state**, **goal**),

i.e. **estimate** is the estimated distance between **state** and **goal** and hence an estimate of the number of steps needed to complete **Path** into a solution. This ensures, that the best **Path**, i.e. the path whose end state is nearest to the **goal** is at the beginning of the set **PrioQueue**.

2. As long as **PrioQueue** is not empty, we take the **Path** from the beginning of this priority queue and remove it from the queue. The state at the end of **Path** is named **state**.
3. If this **state** is the **goal**, a solution has been found and is returned.
4. Otherwise, the states reachable from **state** are inserted into the priority queue. When these states are inserted, we have to compute their estimated distance to **goal** since this distance is used as the priority in **PrioQueue**.

Best first search solves the instance of the 3×3 puzzle shown in Figure 2.3 on page 10 in less than a quarter of a second. However, the solution that is found takes 75 steps. While this is not as ridiculous as the solution found by depth first search, it is still far from an optimal solution. Furthermore, best first search is still not strong enough to solve the 4×4 puzzle shown in Figure 2.23 on page 35.

It should be noted that the fact that the Manhattan distance is a **consistent** heuristic is of no consequence for best first search. Only the A* algorithm, which is presented next, makes use of this fact.

2.7 The A* Search Algorithm

We have seen that best first search can be very fast. However, the solution returned by best first search is not optimal. In contrast, the A* algorithm described next guarantees that a shortest path is found provided the heuristic used is consistent. The basic idea is that the A* search algorithm works similar to the queue based version of breadth first search, but instead of using a simple queue, a priority queue is used instead. The priority $f(s)$ of every state s is given as

$$f(s) := g(s) + h(s),$$

where $g(s)$ computes the length of the path leading from **start** to s and $h(s)$ is the heuristical estimate of the distance from s to **goal**. The details of the A* algorithm are given in Figure 2.20 on page 28 and discussed below. The function **aStarSearch** takes 4 parameters:

1. **start** is a state. This state represents the start state of the search problem.
2. **goal** is the goal state.
3. **nextStates** is a function that takes a state as a parameter. For a state s ,

nextStates(s)

computes the set of all those states that can be reached from s in a single step.

4. **heuristic** is a function that takes two parameters. For two states s_1 and s_2 , the expression

heuristic(s_1, s_2)

computes an estimate of the distance between s_1 and s_2 .

The function **aStarSearch** maintains 5 variables that are crucial for the understanding of the algorithm.

1. **Parent** is a dictionary associating a parent state with those states that have already been encountered during the search, i.e. we have

Parent[s_2] = $s_1 \Rightarrow s_2 \in \text{nextStates}(s_1)$.

Once the goal has been found, this dictionary is used to compute the path from **start** to **goal**.

```

1  aStarSearch := procedure(start, goal, nextStates, heuristic) {
2      Parent    := {};                               // back pointers, represented as dictionary
3      Distance  := { [start, 0] };
4      estGoal   := heuristic(start, goal);
5      Estimate  := { [start, estGoal] }; // estimated distances
6      Frontier  := { [estGoal, start] }; // priority queue
7      while (Frontier != {}) {
8          [_, state] := fromB(Frontier);
9          if (state == goal) {
10             return pathTo(state, Parent);
11         }
12         stateDist := Distance[state];
13         for (neighbour in nextStates(state)) {
14             oldEstimate := Estimate[neighbour];
15             newEstimate := stateDist + 1 + heuristic(neighbour, goal);
16             if (oldEstimate == om || newEstimate < oldEstimate) {
17                 Parent[neighbour] := state;
18                 Distance[neighbour] := stateDist + 1;
19                 Estimate[neighbour] := newEstimate;
20                 Frontier += { [newEstimate, neighbour] };
21                 if (oldEstimate != om) {
22                     Frontier -= { [oldEstimate, neighbour] };
23                 }
24             }
25         }
26     }
27 };

```

Figure 2.20: The A* search algorithm.

2. **Distance** is a dictionary. For every state s that is encountered during the search, this dictionary records the length of the shortest path from **start** to s .
3. **Estimate** is a dictionary. For every state s encountered in the search, **Estimate**[s] is an estimate of the length that a path from **start** to **goal** would have if it would pass through the state s . This estimate is calculated using the equation

$$\text{Estimate}[s] = \text{Distance}[s] + \text{heuristic}(s, \text{goal}).$$

Instead of recalculating this sum every time we need it, we store it in the dictionary **Estimate**. This increases the efficiency of the algorithm.

4. **Frontier** is a **priority queue**. The elements of **Frontier** are pairs of the form

$$[d, s] \quad \text{such that} \quad d = \text{Estimate}[s],$$

i.e. if $[d, s] \in \text{Frontier}$, then the state s has been encountered in the search and it is estimated that a path leading from **start** to **goal** and passing through s would have a length of d .

Now that we have established the key variables, the A* algorithm runs in a **while** loop that does only terminate if either a solution is found or the priority queue **Frontier** is exhausted.

1. First, the **state** with the smallest estimated distance for a path running from **start** to **goal** and passing through **state** is chosen from the priority queue **Frontier**. Note that the call to **fromB** does not only return the pair

[stateEstimate, state]

from **Frontier** that has the lowest value of **stateEstimate**, but also removes this pair from the priority queue.

2. Now if this **state** is the **goal** a solution has been found. Hence, in this case the solution is returned and the function **aStarSearch** terminates.
3. Otherwise, we check the length of the path leading from **start** to **state**. This length is stored in **stateDist**. Effectively, this is the distance between **start** and **state**.
4. Next, we have a loop that iterates over all neighbours of **state**.
 - (a) For every **neighbour** we check the estimated length of a solution passing through **neighbour** and store this length in **oldEstimate**. Note that **oldEstimate** is undefined, i.e. it has the value **om**, if we haven't yet encountered the node **neighbour** in our search.
 - (b) If a solution would go from **start** to **state** and from there proceed to **neighbour**, the estimated length of this solution would be

stateDist + 1 + heuristic(neighbour, goal).

Therefore this value is stored in **newEstimate**.

- (c) Next, we need to check whether this new solution that first passes through **state** and then proceeds to **neighbour** is better than the previous solution that passes through **neighbour**. This check is done by comparing **newEstimate** and **oldEstimate**. Note that we have to take care of the fact that there might be no valid **oldEstimate**.

In case the new solution seems to be better than the old solution, we have to update the **Parent** dictionary, the **Distance** dictionary, and the **Estimate** dictionary. Furthermore, we have to update the priority queue **Frontier**. Here, we have to take care to remove the previous entry for the state **neighbour** if it exists, which is the case if **oldEstimate** is not **om**.

The A* algorithm has been discovered by Hart, Nilsson, and Raphael and was first published in 1968 [HNR68]. However, there was a subtle bug in the first publication which was corrected in 1972 [HNR72].

When we run A* on the 3×3 sliding puzzle, it takes about 17 seconds to solve the instance shown in Figure 2.3 on page 10. If we just look at the time, this seems to be disappointing. However, the good news is that now only 10,061 states are touched in the search for a solution. This is more than a tenfold reduction when compared with breadth first search. The fact that the running time is, nevertheless, quite high results from the complexity of computing the Manhattan distance.

2.7.1 Completeness and Optimality of A* Search

In order to prove the completeness and the optimality of the A* search algorithm it is convenient to reformulate the algorithm: In particular, we replace the dictionary **Parent** by a dictionary **PathDict**. For every state s that is reached during A* search, **PathDict**[s] returns a path that connects the node **start** with the node s . Furthermore, we have added a set **Explored** to the implementation. This set is only needed for the proof of the optimality of the path found by A* search. Figure 2.21 on page 30 shows this version of the algorithm.

Theorem 4 (Completeness and Optimality of A* Search) Assume

$$\mathcal{P} = \langle Q, \text{nextStates}, \text{start}, \text{goal} \rangle$$

is a search problem and

$$\text{heuristic} : Q \rightarrow \mathbb{N}$$

is a consistent heuristic for the search problem \mathcal{P} . Then the A* search algorithm is complete, i.e. if there is a path from **start** to **goal**, then the search is successful and, furthermore, the solution that is computed is a shortest path leading from **start** to **goal**.

```

1  aStarSearch := procedure(start, goal, nextStates, heuristic) {
2      estGoal := heuristic(start, goal);
3      Estimate := { [start, estGoal] }; // estimated total distance
4      Frontier := { [estGoal, start] }; // priority queue
5      PathDict := { [start, [start]] };
6      Distance := { [start, 0] };
7      Explored := {};
8      while (Frontier != {}) {
9          [_, state] := fromB(Frontier);
10         Explored += { state };
11         if (state == goal) {
12             return PathDict[state];
13         }
14         stateDist := Distance[state];
15         for (neighbour in nextStates(state)) {
16             oldEstimate := Estimate[neighbour];
17             newEstimate := stateDist + 1 + heuristic(neighbour, goal);
18             if (oldEstimate == om || newEstimate < oldEstimate) {
19                 Distance[neighbour] := stateDist + 1;
20                 Estimate[neighbour] := newEstimate;
21                 Frontier += { [newEstimate, neighbour] };
22                 PathDict[neighbour] := PathDict[state] + [neighbour];
23                 if (oldEstimate != om) {
24                     Frontier -= { [oldEstimate, neighbour] };
25                 }
26             }
27         }
28     }
29 };

```

Figure 2.21: A path based implementation of A* search.

Proof: In order to prove this theorem, we first need to establish a number of notions that are needed in order to improve our understanding of the A* algorithm. The first of these notions is the notion of **explored** states. A state is said to have been **explored** iff it has been removed from the priority queue **Frontier**. To emphasize this notion I have added the set **Explored** in the implementation of A* search that is shown in Figure 2.21 on page 30. A state is explored once it has been added to the set **Explored**. Note that the variable **Explored** serves no purpose in the implementation of A* search: This variable is only written but is never read. I have added the variable **Explored** to aid this proof. Furthermore, we define a state as **visited** iff it has been entered into the dictionary **Distance**. Before we prove the theorem, let us establish the following claim:

Claim One: If a state s is visited and $P := \text{PathDict}[s]$, then P is a path from **start** to s and the length of P is equal to **Distance** $[s]$.

We establish this claim by a straightforward computational induction.

BC: The first path that is entered into the dictionary **PathDict** is the path **[start]**. This path connects the node **start** with the node **start**. Obviously, this path has the length 0 and that is exactly the value that is entered in the dictionary **Distance** in line 6.

IS: If we add in line 22 the path

$\text{PathDict}[\text{state}] + [\text{neighbour}]$

as a path for the state `neighbour`, then by induction hypotheses we know that $P := \text{PathDict}[\text{state}]$ is a path leading from `start` to `state` and that, furthermore, the length of the path P is equal to $\text{Distance}[\text{state}]$. When we append `neighbour` to the path P , the length of the resulting path is 1 more than the length of P and it is obvious that this new path

$$P + [\text{neighbour}]$$

connects `start` with `neighbour`. The induction hypotheses tell us that $\text{Distance}[\text{state}]$ is equal to the length of P . Therefore, the length of the new path is $\text{Distance}[\text{state}] + 1$ and since stateDist is defined as $\text{Distance}[\text{state}]$, the correct length is stored in line 23.

This concludes the proof of Claim One. ✓

Another notion we need to establish is a function g that is defined for all states s and returns the distance of s from the state `start`, i.e. we have

$$g(s) := \text{dist}(\text{start}, s) \quad \text{for all } s \in Q.$$

Here, the expression $\text{dist}(\text{start}, s)$ returns the length of the shortest path from `start` to s . From the definition of the dictionary Distance it is obvious that

$$g(s) \leq \text{Distance}[s].$$

The reason is that every time an entry for a state s is added to the dictionary Distance , we have found a path from `start` to s that has the length $\text{Distance}[s]$. This might not be the shortest path, hence $g(s)$ might be less than the length of this path.

Next, we introduce the notion of the **estimated total distance** of a state s that has been visited during the search. The **estimated total distance** of a state s is written as $f(s)$ and denotes the estimated length of a path starting in `start` and ending in `goal` that visits the state s in between. Formally, $f(s)$ is defined as follows:

$$f(s) := \text{Distance}[s] + \text{heuristic}[s].$$

Here, $\text{Distance}[s]$ is the number of steps that it takes to reach the state s from `start`, while $\text{heuristic}[s]$ is the estimated distance of a path from s to `goal`. Note that the priority queue `Frontier` is ordered according to the estimated total distance. The shorter the estimated total distance of a node s is, the higher is the priority of s . In order to prove the theorem we need to establish the following claim.

Claim Two: If a state $s \in Q$ has been explored, then we have

$$\text{Distance}[s] = g(s),$$

i.e. the path leading from `start` to s is guaranteed to be a shortest path.

Proof of Claim Two: The proof of Claim Two is a proof by contradiction. We assume that s is a state that has been explored such that

$$g(s) < \text{Distance}[s],$$

i.e. we assume that $P_1 := \text{PathDict}[s]$ is not a shortest path from `start` to s . Then, there must be another path P_2 from `start` to s that is shorter than P_1 . Both P_1 and P_2 start at the same state `start`. Hence, these two paths must have the form

$$P_1 = [\text{start}, x_1, \dots, x_k, x_{k+1}, \dots, x_{n-1}, s], \quad P_2 = [\text{start}, x_1, \dots, x_k, y_{k+1}, \dots, y_{m-1}, s]$$

where $x_{k+1} \neq y_{k+1}$. Here, x_k is the state at which the paths P_1 and P_2 **diverge**. The number k could be 0. In that case the two paths would already diverge at the state `start`, but in general k will be some non-negative integer. Since we assume that P_1 is longer than P_2 we will have $m < n$. Now it is time to make use of the fact that the heuristic h is consistent. First, since $s \in \text{nextStates}(y_{m-1})$ we have that

$$\text{heuristic}(y_{m-1}) \leq 1 + \text{heuristic}(s).$$

Next, as we have $y_{m-1} \in \text{nextStates}(y_{m-2})$ the consistency of heuristic implies

$$\text{heuristic}(y_{m-2}) \leq 1 + \text{heuristic}(y_{m-1}) \leq 2 + \text{heuristic}(s).$$

Continuing this way we conclude that

$$\text{heuristic}(y_{m-i}) \leq i + \text{heuristic}(s) \quad \text{for all } i \in \{1, \dots, m - k - 1\}.$$

Define $i := m - k - 1$. Since $m - i = m - (m - k - 1) = k + 1$ we have shown that

$$\text{heuristic}(y_{k+1}) \leq m - k - 1 + \text{heuristic}(s).$$

This implies

$$\begin{aligned} f(y_{k+1}) &= \text{Distance}[y_{k+1}] + \text{heuristic}(y_{k+1}) \\ &\leq \text{Distance}[y_{k+1}] + m - k - 1 + \text{heuristic}(s) \\ &= (k + 1) + m - k - 1 + \text{heuristic}(s) && \text{since } \text{Distance}[y_{k+1}] = k + 1 \\ &= m + \text{heuristic}(s). \end{aligned}$$

Similarly, we can see that

$$f(y_{k+2}) \leq k + 2 + (m - k - 2) + \text{heuristic}(s) = m + \text{heuristic}(s)$$

and in general we have

$$f(y_j) \leq m + \text{heuristic}(s) \quad \text{for all } j \in \{k + 1, \dots, m - 1\}.$$

Next, we compute $f(s)$. We have

$$f(s) = \text{Distance}[s] + \text{heuristic}(s) = n + \text{heuristic}(s).$$

However, since we have $m < n$ this implies

$$f(y_{k+1}) = m + \text{heuristic}(s) < n + \text{heuristic}(s) = f(s).$$

Now $f(s)$ is the priority of the state s in the priority queue, while $f(y_{k+1})$ is the priority attached to the state y_{k+1} . At the latest, the state y_{k+1} is put onto the priority queue **Frontier** immediately after the state x_k is explored. But then the priority of y_{k+1} is higher than the priority of the state s and hence it is explored prior to s . Since all states of the form

$$y_j \quad \text{for } j \in \{k + 1, \dots, m - 1\}$$

have a priority that is higher than the state s , these states would all have been explored prior to the state s . In particular, once the state y_{m-1} is explored, the estimated total distance of s would be $m + \text{heuristic}(s)$ contradicting the fact that it is $n + \text{heuristic}(s)$. This contradiction proves that our assumption

$$\text{Distance}[s] > g(s)$$

must be wrong and therefore we know that the equation $\text{Distance}[s] = g(s)$ must hold for all states s that have been explored. Hence the validity of Claim Two has been established. ✓

Claim Two implies that if the A* algorithm finds a path from **start** to **goal**, then this path must be optimal. This follows from the fact that before we check whether the search has reached the **goal**, the **state** that is compared to **goal** has been explored and, according to Claim Two, must therefore be optimal.

In order to finish the proof we still need to show that we will find a path provided a path exists. Let us therefore assume that there is a path P from **start** to **goal**. Let us assume that this path has the form

$$P = [s_0, s_1, \dots, s_{n-1}, s_n] \quad \text{where } s_0 = \text{start} \text{ and } s_n = \text{goal}.$$

The proof proceeds by contradiction. We assume that the search never discovers the state **goal**. Let k be the smallest index such that the state s_k is explored but the state s_{k+1} is not explored. Since we have

- (a) $s_0 = \text{start}$,
- (b) the state **start** is explored at the beginning of the search,
- (c) $s_n = \text{goal}$,
- (d) the state **goal** is never explored by our assumption,

an index k such that s_k is explored but the state s_{k+1} is never explored must exist. Now when s_k is explored, s_{k+1} is a neighbour of s_k and is put on the **Frontier**. Define d to be the estimated total distance of s_k . Since **heuristic** is assumed to be consistent, the estimated total distance of s_{k+1} is at most $d + 2$. The reason is that

$$\text{Distance}[s_{k+1}] = \text{Distance}[s_k] + 1 \quad \text{and} \quad \text{heuristic}[s_{k+1}] \leq 1 + \text{heuristic}[s_k],$$

provided we assume that $s_{k+1} \in \text{neighbour}(s_k)$ also implies $s_k \in \text{neighbour}(s_{k+1})$.

There are only finite number of states that have an estimated total distance less or equal than $d + 2$. Only these states have a priority that is as least as high as the priority of s_{k+1} . Once all these states have been explored and removed from the priority queue **Frontier**, s_{k+1} will be explored contrary to our assumption that it is never explored. This contradiction proves that **goal** is eventually explored and therefore A* search is complete. \square

2.8 Bidirectional A* Search

So far, the best search algorithm we have encountered is bidirectional breadth first search. However, in terms of memory consumption, the A* algorithm also looks very promising. Hence, it might be a good idea to combine these two algorithms. Figure 2.22 on page 34 shows the resulting program. This program relates to the A* algorithm shown in Figure 2.20 on page 28 as the algorithm for bidirectional search shown in Figure 2.16 on page 23 relates to breadth first search shown in Figure 2.5 on page 12. The only new idea is that we alternate between the A* search starting from **start** and the A* search starting from **goal** depending on the estimated total distance:

- (a) As long as the search starting from **start** is more promising, we remove states from **FrontierA**.
- (b) Once the total estimated distance of a path starting from **goal** is less than the best total estimated distance of a path starting from **start**, we switch and remove states from **FrontierB**.

When we run bidirectional A* search for the 3×3 sliding puzzle shown in Figure 2.3 on page 10, the program takes one second but only uses 2,963 states. Therefore, I have tried to solve the 4×4 sliding puzzle shown in Figure 2.23 on page 35 using bidirectional A* search. A solution of 44 steps was found in 50 seconds. Only 20,624 states had to be processed to compute this solution! None of the other algorithms presented so far was able to compute the solution.

2.9 Iterative Deepening A* Search

So far, we have combined A* search with bidirectional search and achieved good results. When memory space is too limited for bidirectional A* search to be possible, we can instead combine A* search with **iterative deepening**. The resulting search technique is known as **iterative deepening A* search** and is commonly abbreviated as IDA*. It has been invented by Richard Korf [Kor85]. Figure 2.24 on page 35 shows an implementation of IDA* in SETLX. We proceed to discuss this program.

1. As in the A* search algorithm, the function **idaStarSearch** takes four parameters.

- (a) **start** is a state. This state represents the start state of the search problem.
- (b) **goal** is the goal state.
- (c) **nextStates** is a function that takes a state s as a parameter and computes the set of all those states that can be reached from s in a single step.
- (d) **heuristic** is a function that takes two parameters s_1 and s_2 , where s_1 and s_2 are states. The expression

$$\text{heuristic}(s_1, s_2)$$

computes an estimate of the distance between s_1 and s_2 . It is assumed that this estimate is optimistic, i.e. **heuristic** has to be admissible.

```

1  aStarSearch := procedure(start, goal, nextStates, heuristic) {
2      ParentA   := {};                               ParentB   := {};
3      DistanceA := { [start, 0] };                     DistanceB := { [goal, 0] };
4      estimate  := heuristic(start, goal);
5      EstimateA := { [start, estimate] }; EstimateB := { [goal, estimate] };
6      FrontierA := { [estimate, start] }; FrontierB := { [estimate, goal] };
7      while (FrontierA != {} && FrontierB != {}) {
8          [guessA, stateA] := first(FrontierA); stateADist := DistanceA[stateA];
9          [guessB, stateB] := first(FrontierB); stateBDist := DistanceB[stateB];
10         if (guessA <= guessB) {
11             FrontierA -= { [guessA, stateA] };
12             for (neighbour in nextStates(stateA)) {
13                 oldEstimate := EstimateA[neighbour];
14                 newEstimate := stateADist + 1 + heuristic(neighbour, goal);
15                 if (oldEstimate == om || newEstimate < oldEstimate) {
16                     ParentA[neighbour] := stateA;
17                     DistanceA[neighbour] := stateADist + 1;
18                     EstimateA[neighbour] := newEstimate;
19                     FrontierA += { [newEstimate, neighbour] };
20                     if (oldEstimate != om) {
21                         FrontierA -= { [oldEstimate, neighbour] };
22                     }
23                 }
24                 if (DistanceB[neighbour] != om) {
25                     return combinePaths(neighbour, ParentA, ParentB);
26                 }
27             }
28         } else {
29             FrontierB -= { [guessB, stateB] };
30             for (neighbour in nextStates(stateB)) {
31                 oldEstimate := EstimateB[neighbour];
32                 newEstimate := stateBDist + 1 + heuristic(start, neighbour);
33                 if (oldEstimate == om || newEstimate < oldEstimate) {
34                     ParentB[neighbour] := stateB;
35                     DistanceB[neighbour] := stateBDist + 1;
36                     EstimateB[neighbour] := newEstimate;
37                     FrontierB += { [newEstimate, neighbour] };
38                     if (oldEstimate != om) {
39                         FrontierB -= { [oldEstimate, neighbour] };
40                     }
41                 }
42                 if (DistanceA[neighbour] != om) {
43                     return combinePaths(neighbour, ParentA, ParentB);
44                 }
45             }
46         }
47     }
48 };

```

Figure 2.22: Bidirectional A* search.

```

1  start := [ [ 1, 2, 0, 4 ],
2             [ 14, 7, 12, 10 ],
3             [ 3, 5, 6, 13 ],
4             [ 15, 9, 8, 11 ]
5             ];
6  goal  := [ [ 1, 2, 3, 4 ],
7             [ 5, 6, 7, 8 ],
8             [ 9, 10, 11, 12 ],
9             [ 13, 14, 15, 0 ]
10            ];

```

Figure 2.23: A start state and a goal state for the 4×4 sliding puzzle.

```

1  idaStarSearch := procedure(start, goal, nextStates, heuristic) {
2      limit := heuristic(start, goal);
3      while (true) {
4          Path := search([start], goal, nextStates, limit, heuristic);
5          if (isList(Path)) {
6              return Path;
7          }
8          limit := Path;
9      }
10 };
11 search := procedure(Path, goal, nextStates, limit, heuristic) {
12     state := Path[-1];
13     distance := #Path - 1;
14     total := distance + heuristic(state, goal);
15     if (total > limit) {
16         return total;
17     }
18     if (state == goal) {
19         return Path;
20     }
21     smallest := mathConst("Infinity");
22     for (ns in nextStates(state) | !(ns in Path)) {
23         result := search(Path+[ns], goal, nextStates, limit, heuristic);
24         if (isList(result)) {
25             return result;
26         }
27         smallest := min([result, smallest]);
28     }
29     return smallest;
30 };

```

Figure 2.24: Iterative deepening A* search.

2. The function `idaStarSearch` initializes `limit` to be an estimate of the distance between `start` and `goal`. As we assume that the function `heuristic` is optimistic, we know that there is no path from `start` to `goal` that is shorter than `limit`. Hence, we start our search by assuming that we might find a path that has a length of `limit`.

3. Next, we start a **while** loop. In this loop, we call the function **search** to compute a path from **start** to **goal** that has a length of at most **limit**. The function **search** uses A* search and is described in detail below. When the function **search** returns, there are two cases:

- (a) **search** does find a path. In this case, this path is returned in the variable **Path** and the value of this variable is a list. This list is returned as the solution to the search problem.
- (b) **search** is not able to find a path within the given **limit**. In this case, **search** will not return a list representing a path but instead it will return a number. This number will specify the minimal length that any path leading from **start** to **goal** needs to have. This number is then used to update the **limit** which is used for the next invocation of **search**.

Note that the fact that **search** is able to compute this new **limit** is a significant enhancement over iterative deepening. While we had to test every single possible length in iterative deepening, now the fact that we can intelligently update the **limit** results in a considerable saving of computation time.

We proceed to discuss the function **search**. This function takes 5 parameters, which we describe next.

- 1. **Path** is a list of states. This list starts with the state **start**. If **state** is the last state on this list, then **Path** represents a path leading from **start** to **state**.
- 2. **goal** is another state. The purpose of the recursive invocations of **search** is to find a path from **state** to **goal**, where **state** is the last element of the list **Path**.
- 3. **nextStates** is a function that takes a state s as input and computes the set of states that are reachable from s in one step.
- 4. **limit** is the upper bound for the path from **start** to **goal**. If the procedure **search** is not able to find a path from **start** to **goal** that has a length of at most **limit**, then the search is unsuccessful. In that case, instead of a path the function **search** returns a new estimate for the distance between **start** and **goal**. Of course, this new estimate will be bigger than **limit**.
- 5. **heuristic** is a function taking to states as arguments. The invocation **heuristic**(s_1, s_2) computes an *estimate* of the distance between the states s_1 and s_2 . It is assumed that this estimate is optimistic, i.e. the value returned by **heuristic**(s_1, s_2) is less or equal than the true distance between s_1 and s_2 .

We proceed to describe the implementation of the function **search**.

- 1. The variable **state** is assigned the last element of **Path**. Hence, **Path** connects **start** and **state**.
- 2. The length of the path connecting **start** and **state** is stored in **distance**.
- 3. Since **heuristic** is assumed to be optimistic, if we want to extend **Path**, then the best we can hope for is to find a path from **start** to **goal** that has a length of

$$\text{distance} + \text{heuristic}(\text{state}, \text{goal}).$$

This length is computed and saved in the variable **total**.

- 4. If **total** is bigger than **limit**, it is not possible to find a path from **start** to **goal** passing through **state** that has a length of at most **limit**. Hence, in this case we return **total** to communicate that the limit needs to be increased to have at least a value of **total**.
- 5. If we are lucky and **state** is equal to **goal**, the search is successful and **Path** is returned.
- 6. Otherwise, we iterate over all nodes **ns** reachable from **state** that have not already been visited by **Path**. If **ns** is a node of this kind, we extend the **Path** so that this node is visited next. The resulting path has the form

$$\text{Path} + [\text{ns}].$$

Next, we recursively start a new search starting from the node **ns**. If this search is successful, the resulting path is returned. Otherwise, the search returns the minimum distance that is needed to reach the state **goal** from the state **start** on a path via the state **ns**. If this distance is smaller than the distance **smallest** that is returned from visiting previous neighbouring nodes, the variable **smallest** variable is updated accordingly. This way, if the **for** loop is not able to return a path leading to **goal**, the variable **smallest** contains a lower bound for the distance that is needed to reach **goal** by a path that extends the given **Path**.

Note: At this point, a natural question is to ask whether the **for** loop should collect all paths leading to **goal** and then only return the path that is shortest. However, this is not necessary: Every time the function **search** is invoked it is already guaranteed that there is no path that is shorter than the parameter **limit**. Therefore, if **search** is able to find a path that has a length of at most **limit**, this path is known to be optimal.

Iterative deepening A* is a complete search algorithm that does find an optimal path, provided that the employed heuristic is optimistic. On the instance of the 3×3 sliding puzzle shown on Figure 2.3 on page 10, this algorithm takes about 1.7 seconds to solve the puzzle. For the 4×4 sliding puzzle, the algorithm takes about 154 seconds. Although this is more than the time needed by bidirectional A* search, the good news is that the IDA* algorithm does not need much memory since basically only the path discovered so far is stored in memory. Hence, IDA* is a viable alternative if the available memory is not sufficient to support the bidirectional A* algorithm.

2.10 The A*-IDA* Search Algorithm

So far, from all of the search algorithms we have tried, the bidirectional A* search has performed best. However, bidirectional A* search is only feasible if sufficient memory is available. While IDA* requires more time, its memory consumption is much lower than the memory consumption of bidirectional A*. Hence, it is natural to try to combine the A* algorithm and the IDA* algorithm. Concretely, the idea is to run an A* search from the **start** node until memory is more or less exhausted. Then, we start IDA* from the **goal** node and search until we find any of the nodes discovered by the A* search that had been started from the **start** node.

An implementation of the A*-IDA* algorithm is shown in Figure 2.25 on page 38 and Figure 2.26 on page 39. We begin with a discussion of the procedure **aStarIdaStarSearch**.

1. The procedure takes 5 arguments.
 - (a) **start** and **goal** are nodes. The procedure tries to find a path connecting **start** and **goal**.
 - (b) **nextStates** is a function that takes a state s as input and computes the set of states that are reachable from s in one step.
 - (c) **heuristic** computes an *estimate* of the distance between s_1 and s_2 . It is assumed that this estimate is optimistic, i.e. the value returned by **heuristic**(s_1, s_2) is less or equal than the true distance between s_1 and s_2 .
 - (d) **size** is the maximal number of states that the A* search is allowed to explore before the algorithm switches over to IDA* search.
2. The basic idea behind the A*-IDA* algorithm is to first use A* search to find a path from **start** to **goal**. If this is successfully done without visiting more than **size** nodes, the algorithm terminates and returns the path that has been found. Otherwise, the algorithm switches over to an IDA* search that starts from **goal** and tries to connect goal to any of the nodes that have been encountered during the A* search. To this end, the procedure **aStarIdaStarSearch** maintains the following variables.
 - (a) **Parent** is a dictionary associating a parent state with those states that have already been encountered during the search, i.e. we have

$$\text{Parent}[s_2] = s_1 \Rightarrow s_2 \in \text{nextStates}(s_1).$$

Once the goal has been found, this dictionary is used to compute the path from **start** to **goal**.

```

1  aStarIdaStarSearch := procedure(start, goal, nextStates, heuristic, size) {
2      Parent      := {};           // dictionary: state |-> predecessor(state)
3      Distance    := { [start, 0] }; // dictionary: state |-> distance(start, state)
4      est         := heuristic(start, goal);
5      Estimate    := { [start, est] }; // dictionary: state |-> distance(start, state, goal)
6      Frontier    := { [est, start] }; // priority queue
7      while (#Distance < size && Frontier != {}) {
8          if (#Distance % 1000 == 0) {
9              print("#Distance = $#Distance$");
10         }
11         [_, state] := fromB(Frontier);
12         if (state == goal) {
13             print("Number of states touched: $#Distance$");
14             return pathTo(state, Parent);
15         }
16         stateDist := Distance[state];
17         for (neighbour in nextStates(state)) {
18             oldEstimate := Estimate[neighbour];
19             newEstimate := stateDist + 1 + heuristic(neighbour, goal);
20             if (oldEstimate == om || newEstimate < oldEstimate) {
21                 Parent[neighbour] := state;
22                 Distance[neighbour] := stateDist + 1;
23                 Estimate[neighbour] := newEstimate;
24                 Frontier += { [newEstimate, neighbour] };
25                 if (oldEstimate != om) {
26                     Frontier -= { [oldEstimate, neighbour] };
27                 }
28             }
29         }
30     }
31     Path := deepeningSearch(goal, start, nextStates, heuristic, Distance);
32     return pathTo(Path[-1], Parent) + reverse(Path)[2..];
33 };

```

Figure 2.25: The A*-IDA* search algorithm, part I.

- (b) **Distance** is a dictionary that remembers for every state s that is encountered during the A* search the length of the shortest path from **start** to s .
- (c) **Estimate** is a dictionary. For every state s encountered in the A* search, **Estimate**[s] is an estimate of the length that a path from **start** to **goal** would have if it would pass through the state s . This estimate is calculated using the equation

$$\text{Estimate}[s] = \text{Distance}[s] + \text{heuristic}(s, \text{goal}).$$

Instead of recalculating this sum every time we need it, we store it in the dictionary **Estimate**.

- (d) **Frontier** is a **priority queue**. The elements of **Frontier** are pairs of the form

$$[d, s] \quad \text{such that} \quad d = \text{Estimate}[s],$$

i.e. if $[d, s] \in \text{Frontier}$, then the state s has been encountered in the A* search and it is estimated that a path leading from **start** to **goal** and passing through s would have a length of d .

3. The A* search runs exactly as discussed previously. The only difference is that the **while** loop is terminated once the dictionary **Distance** has more than **size** entries. If we are lucky, the A* search is

already able to find the **goal** and the algorithm terminates.

4. Otherwise, the procedure **deepeningSearch** is called. This procedure starts an iterative deepening A* search from the node **goal**. This search terminates as soon as a state is found that has already been encountered during the A* search. The set of these nodes is given to the procedure **deepeningSearch** via the parameter **Distance**. The procedure **deepeningSearch** returns a pair. The first component of this pair is the state **s**. This is the state in **Distance** that has been reached by the IDA* search. The second component is the path **P** that leads from the node **s** to the node **goal** but that does not include the node **s**. In order to compute a path from **start** to **goal**, we still have to compute a path from **start** to **s**. This path is then combined with the path **P** and the resulting path is returned.

```

1  deepeningSearch := procedure(goal, start, nextStates, heuristic, Distance) {
2      limit := 0;
3      while (true) {
4          print("limit = $limit$");
5          Path := search([goal], start, nextStates, limit, heuristic, Distance);
6          if (isList(Path)) {
7              return Path;
8          }
9          limit := Path;
10     }
11 };
12 search := procedure(Path, start, nextStates, limit, heuristic, Distance) {
13     state := Path[-1];
14     total := #Path - 1 + heuristic(state, start);
15     if (total > limit) {
16         return total;
17     }
18     if (Distance[state] != om) {
19         return Path;
20     }
21     smallest := mathConst("Infinity");
22     for (ns in nextStates(state) | !(ns in Path) ) {
23         result := search(Path+[ns], start, nextStates, limit, heuristic, Distance);
24         if (isList(result)) {
25             return result;
26         }
27         smallest := min([smallest, result]);
28     }
29     return smallest;
30 };

```

Figure 2.26: The A*-IDA* search algorithm, part II.

Iterative deepening A*-IDA* is a complete search algorithm. On the instance of the 3×3 sliding puzzle shown on Figure 2.3 on page 10, this algorithm takes about 0.8 seconds to solve the puzzle. For the 4×4 sliding puzzle, if the algorithm is allowed to visit at most 3000 states, the algorithm takes less than 6 seconds.

Exercise 4: Assume that you have 3 water buckets: The first bucket can hold 12 litres of water, the second bucket can hold 8 litres, while the last bucket can hold 3 litres. There are three types of action:

1. A bucket can be completely filled.
2. A bucket can be completely emptied.

3. The content of one bucket can be poured into another bucket. Then, there are two cases.
 - (a) If the second bucket has enough free space for all the water in the first bucket, then the first bucket is emptied and all the water from the first bucket is poured into the second bucket.
 - (b) However, if there is not enough space in the second bucket, then the second bucket is filled completely, while the water that does not fit into the second bucket remains in the first bucket.

Your goal is to measure out exactly one litre. Write a program that computes a plan to achieve this goal. \diamond

Exercise 5: The founder of **Taoism**, the Chinese philosopher **Laozi** once said:

“A journey of a thousand miles begins but with a single step”.

This proverb is the foundation of **taoistic search**. The idea is, instead of trying to reach the goal directly, we rather define some intermediate states which are easier to reach than the goal state and that are nearer to the goal than the start state. To make this idea more precise, consider the following instance of the 15-puzzle, where the states **Start** and **Goal** are given as follows:

Start := <pre> +---+---+---+---+ 15 14 8 12 +---+---+---+---+ 10 11 13 9 +---+---+---+---+ 6 2 5 1 +---+---+---+---+ 3 4 7 +---+---+---+---+ </pre>	Goal := <pre> +---+---+---+---+ 1 2 3 4 +---+---+---+---+ 5 6 7 8 +---+---+---+---+ 9 10 11 12 +---+---+---+---+ 13 14 15 +---+---+---+---+ </pre>
---	--

In order to solve the puzzle, we could try to first move the tiles numbered with 1 and 2 into the upper left corner. The resulting state would have the following form:

```

+---+---+---+---+
| 1 | 2 | * | * |
+---+---+---+---+
| * | * | * | * |
+---+---+---+---+
| * | * | * | * |
+---+---+---+---+
| * | * | * | * |
+---+---+---+---+

```

Here, the character “*” is used as a wildcard character, i.e. we do not care about the actual character in the state, for we only want to ensure that the first two tiles are positioned correctly. Once we have reached a state specified by the pattern given above, we could then proceed to reach a state that is described by the following pattern:

```

+---+---+---+---+
| 1 | 2 | * | * |
+---+---+---+---+
| 5 | 6 | * | * |
+---+---+---+---+
| * | * | * | * |
+---+---+---+---+
| * | * | * | * |
+---+---+---+---+

```

This way, slowly but surely we will reach the goal. I have prepared a framework for taoistic search. The file <https://github.com/karlstroetmann/Artificial-Intelligence/blob/master/Set1X/sliding-puzzle-frame.stlx> contains a framework for the sliding puzzle where some functions are left unimplemented. The file

<https://github.com/karlstroetmann/Artificial-Intelligence/blob/master/Set1X/a-star-lao-tzu.stlx> is also required. It contains an adapted form of A* search. More Details will be given in the lecture. Your task is to implement the missing functions in the file `sliding-puzzle-frame.stlx`. ◇

Chapter 3

Constraint Satisfaction

In this chapter we discuss algorithms for solving **constraint satisfaction problems**. These problems can be seen as a refinement of the search problems discussed in the previous chapter: In a search problem, the states are abstract and therefore have no structure that can be exploited to guide the search. In a **constraint satisfaction problem**, the states have a structure that can be exploited to guide the search. This chapter is structured as follows:

1. The first section defines the notion of a constraint satisfaction problem. In order to illustrate this concept, two examples of constraint satisfaction problems are presented. After that, we discuss applications of constraint satisfaction problems.
2. The simplest algorithm to solve a constraint satisfaction problem is via **brute force search**. The idea behind brute force search is to test all possible variable assignments.
3. In most cases, the search space is too big to be enumerated completely. **Backtracking search** improves on brute force search by mixing the generation of assignments with the testing of the constraints. In many cases, this approach can drastically improve the performance of the search algorithm.
4. Backtracking search can be refined by using **constraint propagation** and by using the **most restricted variable** heuristic.
5. Furthermore, checking the **consistency** of the values assigned to different variables can reduce the size of the search space considerably.
6. Finally, **local search** is a completely different approach to solve constraint satisfaction problems.

3.1 Formal Definition of Constraint Satisfaction Problems

Formally, we define a **constraint satisfaction problem** as a triple

$$\mathcal{P} := \langle \text{Vars}, \text{Values}, \text{Constraints} \rangle$$

where

1. **Vars** is a set of strings which serve as **variables**,
2. **Values** is a set of **values** that can be assigned to the variables in **Vars**.
3. **Constraints** is a set of formulæ from **first order logic**. Each of these formulæ is called a **constraint** of \mathcal{P} .

In order to be able to interpret these formulæ, we need a **first order structure** $\mathcal{S} = \langle \mathcal{U}, \mathcal{J} \rangle$. Here, \mathcal{U} is the **universe** of \mathcal{S} and we will assume that this universe is identical to the set **Values**, that is we have

$$\mathcal{U} = \text{Values}.$$

The second component \mathcal{J} defines the [interpretations](#) of the function symbols and predicate symbols that are used in the formulæ defining the constraints. In what follows we assume that these interpretations are understood from the context of the constraint satisfaction problem \mathcal{P} .

In the following, the abbreviation CSP is short for [constraint satisfaction problem](#). Given a CSP

$$\mathcal{P} = \langle \text{Vars}, \text{Values}, \text{Constraints} \rangle,$$

a [variable assignment](#) for \mathcal{P} is a function

$$A : \text{Vars} \rightarrow \text{Values}.$$

A variable assignment A is a [solution](#) of the CSP \mathcal{P} if, given the assignment A , all constraints of \mathcal{P} are satisfied. Finally, a [partial variable assignment](#) B for \mathcal{P} is a function

$$B : \text{Vars} \rightarrow \text{Values} \cup \{\Omega\} \quad \text{where } \Omega \text{ denotes the undefined value.}$$

Hence, a partial variable assignment does not assign values to all variables. Instead, it assigns values only to a subset of the set Vars . The [domain](#) $\text{dom}(B)$ of a partial variable assignment B is the set of those variables that are assigned a value different from Ω , i.e. we define

$$\text{dom}(B) := \{x \in \text{Vars} \mid B(x) \neq \Omega\}.$$

We proceed to illustrate the definitions given so far with two examples.




Figure 3.1: A map of Australia.

3.1.1 Example: Map Colouring

In **map colouring** a map showing different state borders is given and the task is to colour the different states such that no two states that have a common border share the same colour. Figure 3.1 on page 43 shows a map of Australia. There are seven different states in Australia:

1. Western Australia, abbreviated as WA,
2. Northern Territory, abbreviated as NT,
3. South Australia, abbreviated as SA,
4. Queensland, abbreviated as Q,
5. New South Wales, abbreviated as NSW,
6. Victoria, abbreviated as V, and
7. Tasmania, abbreviated as T.

Figure 3.1 would certainly look better if different states had been coloured with different colours. For the purpose of this example let us assume that we have only three colours available. The question then is whether it is possible to colour the different states in a way that no two neighbouring states share the same colour. This problem can be formalized as a constraint satisfaction problem. To this end we define:

1. $\text{Vars} := \{\text{WA}, \text{NT}, \text{SA}, \text{Q}, \text{NSW}, \text{V}, \text{T}\},$
2. $\text{Values} := \{\text{red}, \text{green}, \text{blue}\},$
3. $\text{Constraints} :=$
 $\{\text{WA} \neq \text{NT}, \text{WA} \neq \text{SA}, \text{NT} \neq \text{SA}, \text{NT} \neq \text{Q}, \text{SA} \neq \text{Q}, \text{SA} \neq \text{NSW}, \text{SA} \neq \text{V}, \text{Q} \neq \text{NSW}, \text{NSW} \neq \text{V}, \text{V} \neq \text{T}\}$

Then $\mathcal{P} := \langle \text{Vars}, \text{Values}, \text{Constraints} \rangle$ is a constraint satisfaction problem. If we define the assignment A such that

1. $A(\text{WA}) = \text{blue},$
2. $A(\text{NT}) = \text{red},$
3. $A(\text{SA}) = \text{green},$
4. $A(\text{Q}) = \text{blue},$
5. $A(\text{NSW}) = \text{red},$
6. $A(\text{V}) = \text{blue},$
7. $A(\text{T}) = \text{red},$

then you can check that the assignment A is indeed a solution to the constraint satisfaction problem \mathcal{P} .

3.1.2 Example: The Eight Queens Puzzle

The **eight queens problem** asks to put 8 queens onto a chessboard such that no queen can attack another queen. In **chess**, a queen can attack all pieces that are either in the same row, the same column, or the same diagonal. If we want to put 8 queens on a chessboard such that no two queens can attack each other, we have to put exactly one queen in every row: If we would put more than one queen in a row, the queens in that row can attack each other. If we would leave a row empty, then, given that the other rows contain at most one queen, there would be less than 8 queens on the board. Therefore, in order to model the eight queens problem as a constraint satisfaction problem, we will use the following set of variables:

$$\text{Vars} := \{V_1, V_2, V_3, V_4, V_5, V_6, V_7, V_8\},$$

where for $i \in \{1, \dots, 8\}$ the variable V_i specifies the column of the queen that is placed in row i . As the columns run from one to eight, we define the set **Values** as

$$\mathbf{Values} := \{1, 2, 3, 4, 5, 6, 7, 8\}.$$

Next, let us define the constraints. There are three different types of constraints.

1. We have constraints that express that no two queens positioned in different rows share the same column. To capture these constraints, we define

$$\mathbf{SameRow} := \{V_i \neq V_j \mid i \in \{1, \dots, 8\} \wedge j \in \{1, \dots, 8\} \wedge j < i\}.$$

Here the condition $i < j$ ensures that, for example, we have the constraint $V_2 \neq V_1$ but not the constraint $V_1 \neq V_2$, as the latter constraint would be redundant if the former constraint has already been established.

2. We have constraints that express that no two queens positioned in different rows share the same rising diagonal. To capture these constraints, we define

$$\mathbf{SameRising} := \{i + V_i \neq j + V_j \mid i \in \{1, \dots, 8\} \wedge j \in \{1, \dots, 8\} \wedge j < i\}.$$

A more detailed explanation of these and the following constraints is given in my lecture notes on [logic](#).

3. We have constraints that express that no two queens positioned in different rows share the same falling diagonal. To capture these constraints, we define

$$\mathbf{SameFalling} := \{i - V_i \neq j - V_j \mid i \in \{1, \dots, 8\} \wedge j \in \{1, \dots, 8\} \wedge j < i\}.$$

Then, the set of constraints is defined as

$$\mathbf{Constraints} := \mathbf{SameRow} \cup \mathbf{SameRising} \cup \mathbf{SameFalling}$$

and the eight queens problem can be stated as the constraint satisfaction problem

$$\mathcal{P} := \langle \mathbf{Vars}, \mathbf{Values}, \mathbf{Constraints} \rangle.$$

If we define the assignment A such that

$$A(1) := 4, A(2) := 8, A(3) := 1, A(4) := 2, A(5) := 6, A(6) := 2, A(7) := 7, A(8) := 5,$$

then it is easy to see that this assignment is a solution of the eight queens problem. This solution is shown in Figure 3.2 on page 46.

Later, when we implement procedures to solve CSPs, we will represent variable assignments and partial variable assignments as binary relations. For example, A would then be represented as the relation

$$A = \{ \langle V_1, 4 \rangle, \langle V_2, 8 \rangle, \langle V_3, 1 \rangle, \langle V_4, 2 \rangle, \langle V_5, 6 \rangle, \langle V_6, 2 \rangle, \langle V_7, 7 \rangle, \langle V_8, 5 \rangle \}.$$

If we define

$$B := \{ \langle V_1, 4 \rangle, \langle V_2, 8 \rangle, \langle V_3, 1 \rangle \},$$

then B is a partial assignment and $\text{dom}(B) = \{V_1, V_2, V_3\}$. This partial assignment is shown in Figure 3.3 on page 46.

Figure 3.4 on page 47 shows a SETLX program that can be used to create the eight queens puzzle as a CSP. The code shown in this figure is more general than the eight queens puzzle: Given a natural number n , the function call `queensCSP(n)` creates a constraint satisfaction problem \mathcal{P} that generalizes the eight queens problem to the problem of putting n queens on a board of size n times n .

The beauty of [constraint programming](#) is the fact that we will be able to develop a so called [constraint solver](#) that takes as input a CSP like the one produced by the program shown in Figure 3.4 and that is then capable of computing a solution. In effect, this enables us to use [declarative programming](#): Instead of specifying an algorithm that solves a problem we confine ourselves to specifying the problem precisely and then let a general purpose problem solver do the job of finding the solution. This approach of declarative programming was one of the main ideas incorporated in the programming language [Prolog](#). While [Prolog](#) could not live up to the promises made in declarative programming as a viable general purpose programming method, constraint programming has proven to be very useful in a number of domains.



Figure 3.2: A solution of the eight queens problem.

Figure 3.3: The partial assignment $\{\langle V_1, 4 \rangle, \langle V_2, 8 \rangle, \langle V_3, 1 \rangle\}$.

Exercise 6: Formulate the n queens puzzle as a search problem. Represent states as lists of the form

$$[Q_1, \dots, Q_k] \quad \text{where } k \in \{0, \dots, n\} \text{ and } Q_i \in \{1, \dots, n\}.$$

Here, Q_i specifies the column of the queen positioned in the i -th row. If we assume that every row has at most one queen, the set of all states is then defined as

```

1  createCSP := procedure(n) {
2      Variables := { "V$i$" : i in {1..n} };
3      Values    := { 1 .. n };
4      Constraints := {};
5      for (i in [2..n], j in [1..i-1]) {
6          Constraints += { "V$i$ != V$j$" };
7          Constraints += { "$i$ + V$i$ != $j$ + V$j$" };
8          Constraints += { "$i$ - V$i$ != $j$ - V$j$" };
9      }
10     return [Variables, Values, Constraints];
11 };

```

Figure 3.4: SETLX code to create the CSP representing the eight-queens puzzle.

$$\mathcal{Q} := \{[Q_1, \dots, Q_k] \mid k \in \{0, \dots, n\} \wedge \forall i \in \{1, \dots, n\} : Q_i \in \{1, \dots, n\} \wedge \forall i, j \in \{1, \dots, k\} : (i \neq j \rightarrow Q_i \neq Q_j)\}.$$

- (a) Derive a mathematical expression that returns the number of states for a given value of n . What is the number of states if $n = 8$?
- (b) Which of the search algorithms presented in the previous chapter could be used to solve the 8 queens puzzle? Which of these algorithms will work best?
- (c) Solve the 8 queens puzzle using an appropriate search algorithm.

Note that we do not know the goal of this search problem. Instead, you have to adapt the search algorithm so that it takes a function `isGoal` instead of a state `goal`. For a given state s , the function `isGoal(s)` returns true iff s is a solution of the n queens puzzle. \diamond

3.1.3 Applications

Besides the toy problems discussed so far, there are a number of industrial applications of constraint satisfaction problems. The most important application seem to be variants of **scheduling problems**. A simple example of a scheduling problem is the problem of generating a time table for a school. A school has various teachers, each of which can teach some subjects but not others. Furthermore, there are a number of classes that must be taught in different subjects. The problem is then to assign teachers to classes and to create a time table. In practice, **crew scheduling** is an important problem. For example, airlines have to solve a crew scheduling problem in order to efficiently assign crews to aircrafts.

3.2 Brute Force Search

The most straightforward algorithm to solve a CSP is to test all possible combinations of assigning values to variables. If there are n different values that can be assigned to k variables, this amounts to checking n^k different assignments. For example, for the eight queens problem there are 8 variables and 8 possible values leading to

$$8^8 = 2^{24} = 16,777,216$$

different assignments that need to be tested. Given the clock speed of modern computers, checking a million assignments per second is plausible. Hence, this approach is able to solve the eight queens problem in less than 5 minutes. The approach of testing all possible combinations is known as **brute force search**. An implementation of brute force search is shown in Figure 3.5 on page 48.

The procedure `solve` gets a constraint satisfaction problem CSP as its input. This CSP is given a triple of the form

CSP = [Variables, Values, Constraints].

```

1  solve := procedure(CSP) {
2      return brute_force_search({}, CSP);
3  };
4  brute_force_search := procedure(Assignment, CSP) {
5      [Variables, Values, Constraints] := CSP;
6      if (#Assignment == #Variables) { // all variables have been assigned
7          if (check_all_constraints(Assignment, Constraints)) {
8              return Assignment;
9          }
10         return;
11     }
12     var := rnd(Variables - domain(Assignment));
13     for (value in Values) {
14         NewAss := Assignment + { [var, value] };
15         result := brute_force_search(NewAss, CSP);
16         if (result != om) {
17             return result;
18         }
19     }
20 };

```

Figure 3.5: Solving a CSP via brute force search.

The sole purpose of `search` is to call the function `brute_force_search`. The implementation of this function is recursive and the function takes two arguments.

1. `Assignment` is a partial assignment of values to variables. Initially, this assignment will be empty. Every recursive call of `brute_force_search` adds the assignment of one variable to the given assignment.
2. `CSP` is a triple of the form

$$\text{CSP} = [\text{Variables}, \text{Values}, \text{Constraints}].$$

Here, `Constraints` is a set of first order logic formulæ that are given as strings. These strings have to follow the syntax of SETLX formulæ so that they can be evaluated using the SETLX function `eval`.

The implementation of `brute_force_search` works as follows:

1. If all variables have been assigned a value, the dictionary `Assignment` will have the same number of entries as the set `Variables` has elements. Hence, in that case `Assignment` is a complete assignment of all variables and we now have to test whether all constraints are satisfied. This is done using the auxiliary procedure `check_all_constraints` that is shown in Figure 3.6 on page 49. If the current `Assignment` does indeed satisfy all constraints, it is a solution to the given `CSP` and is therefore returned. If, instead, some constraint is violated, then `brute_force_search` returns the undefined value Ω .
2. If the assignment is not yet complete, we randomly¹ pick a variable `var` from the set of those `Variables` that still have no value assigned. Then, for every possible `value` in the set `Values`, we augment the current partial `Assignment` to the new assignment

$$\text{Assignment} + \{ [\text{var}, \text{value}] \}.$$

Next, the algorithm recursively tries to find a solution for this new partial assignment. If this recursive call succeeds, the solution it has computed is returned. Otherwise, another `value` is tried.

¹ Experimental evidence suggests that picking a random variable performs much better than sorting the variables and always picking the first unassigned variable.

```

21  check_all_constraints := procedure(Assignment, Constraints) {
22      return forall(f in Constraints | eval_constraint(Assignment, f));
23  };
24  eval_constraint := procedure(Assignment, Formula) {
25      for ([var, value] in Assignment) {
26          execute("$var$ := $value$;");
27      }
28      return eval(Formula);
29  };

```

Figure 3.6: Auxiliary procedures for brute force search.

The function `check_all_constraints` takes a complete variable `Assignment` as its first input. The second input is the set of `Constraints`. For all formulæ `f` from the set `Constraints`, it checks whether `f` is satisfied provided all variables are assigned as specified in `Assignment`. This check is done using the auxiliary function `eval_constraint`.

The procedure `eval_constraint` takes an `Assignment` and a `Formula` that is supposed to be a constraint. Its purpose is to evaluate `Formula` using the `Assignment`. In order for this to be possible we have to assume that

$$\text{var}(\text{Formula}) \subseteq \text{dom}(\text{Assignment}),$$

i.e. all variables occurring in `Formula` have a value assigned in `Assignment`.

The `Formula` is evaluated by first assigning the variables assigned in `Assignment` to their corresponding values. Once these assignments have been executed in the `for`-loop, the `Formula` can be evaluated using the procedure `eval`, which is one of the predefined procedures in SETLX.

When I tested this brute force search with the eight queens problem, it took about 300 seconds to compute a solution. In contrast, the seven queens problem took roughly 14 seconds. As we have

$$\frac{8^8}{7^7} \approx 20 \qquad \frac{300}{14} \approx 21$$

this shows that the computation time does indeed grow with the number of possible assignments that have to be checked. However, the correspondence is not exact. The reason is that we stop our search as soon as a solution is found. If we are lucky and the given CSP is easy to solve, this might happen when we have checked only a small portion of the set of all possible assignments.

3.3 Backtracking Search

Due to the combinatorial explosion of the search space, brute force search is only viable when we deal with small problems. One approach to solve a CSP that is both simple and at least more efficient than brute force search is [backtracking](#). Figure 3.7 on page 50 shows a simple CSP solver that employs backtracking. We discuss this program next. The procedure `solve` takes a constraint satisfaction problem `CSP` as input and tries to find a solution.

1. First, the CSP is split into its components.
2. Next, for every constraint `f` of the given CSP, we compute the set of variables that are used in `f`. This is done using the procedure `collectVars` that is shown in Figure 3.8 on page 51. These variables are then stored together with the constraint `f` and the correspondingly modified data structure is stored in CSP and is called an [augmented CSP](#).

The reason to compute and store these variables is efficiency: When we later check whether a constraint `f` is satisfied for a partial variable assignment `Assignment` where `Assignment` is stored as a dictionary,

```

1  solve := procedure(CSP) {
2      [Vars, Vals, Constrs] := CSP;
3      CSP := [Vars, Vals, { [f, collectVars(f)] : f in Constrs }];
4      check {
5          return bt_search({}, CSP);
6      }
7  };
8  bt_search := procedure(Assignment, CSP) {
9      [Variables, Values, Constraints] := CSP;
10     if (#Assignment == #Variables) {
11         return Assignment;
12     }
13     var := rnd(Variables - domain(Assignment));
14     for (value in Values) {
15         check {
16             if (is_consistent(var, value, Assignment, Constraints)) {
17                 return bt_search(Assignment + { [var, value] }, CSP);
18             }
19         }
20     }
21     backtrack;
22 };

```

Figure 3.7: A backtracking CSP solver.

we only need to check the constraint f iff all of the variables occurring in f are elements of the domain of **Assignment**. It would be wasteful to compute these variables every time.

3. Next, we call the procedure **bt_search** to compute a solution of **CSP**. This procedure is enclosed in a **check**-block. Conceptually, this **check**-block is the same as if we had enclosed the call to **bt_search** in a **try-catch**-block as shown below:

```

try {
    return bt_search({}, csp);
} catch(e) {}

```

The point is that the procedure **bt_search** either returns a solution or, if it is not able to find a solution, it throws a special kind of exception, a so called **backtrack exception**. The **check**-block ensures that this exception is silently discarded. It is just a syntactical convenience that is more concise than using **try** and **catch**.

Next, we discuss the implementation of the procedure **bt_search**. This procedure receives a partial assignment **Assignment** as input together with an augmented **CSP**. This partial assignment is **consistent** with **CSP**: If f is a constraint of **CSP** such that all the variables occurring in f are members of $\text{dom}(\text{Assignment})$, then evaluating f using **Assignment** yields **true**. Initially, this partial assignment is empty and hence trivially consistent. The idea is to extend this partial assignment until it is a complete assignment that satisfies all constraints of the given **CSP**.

1. First, the augmented **CSP** is split into its components.
2. Next, if **Assignment** is already a complete variable assignment, i.e. if the dictionary **Assignment** has as many elements as there are variables, then it must be a solution of the **CSP** and, therefore, it is returned.

```

1  collectVars := procedure(f) {
2      return varsTerm(parseTerm(f));
3  };
4  varsTerm := procedure(f) {
5      match (f) {
6          case v | isVariable(v): return { varName(v) };
7          case n | isNumber(n):   return {};
8          case lhs + rhs:         return varsTerm(lhs) + varsTerm(rhs);
9          case lhs - rhs:         return varsTerm(lhs) + varsTerm(rhs);
10         case lhs * rhs:         return varsTerm(lhs) + varsTerm(rhs);
11         case lhs / rhs:         return varsTerm(lhs) + varsTerm(rhs);
12         case lhs \ rhs:         return varsTerm(lhs) + varsTerm(rhs);
13         case lhs % rhs:         return varsTerm(lhs) + varsTerm(rhs);
14         case lhs == rhs:        return varsTerm(lhs) + varsTerm(rhs);
15         case lhs != rhs:        return varsTerm(lhs) + varsTerm(rhs);
16         case !formula:          return varsTerm(formula);
17         case lhs && rhs:         return varsTerm(lhs) + varsTerm(rhs);
18         case lhs || rhs:        return varsTerm(lhs) + varsTerm(rhs);
19         case lhs => rhs:         return varsTerm(lhs) + varsTerm(rhs);
20         case lhs <==> rhs:       return varsTerm(lhs) + varsTerm(rhs);
21         case lhs <!=> rhs:       return varsTerm(lhs) + varsTerm(rhs);
22         default:                return +/ [ varsTerm(t) : t in args(f) ];
23     }
24 };

```

Figure 3.8: The procedure `collectVars`.

3. Otherwise, we have to extend the partial **Assignment**. In order to do so, we first have to select a variable **var** that has not yet been assigned a value in **Assignment** so far.
4. Next, it is tried to assign a **value** to the selected variable **var**. After assigning a **value** to **var**, we immediately check whether this assignment would be consistent with the constraints using the procedure **is_consistent**. If the partial **Assignment** turns out to be consistent, the partial **Assignment** is extended to the new partial assignment

Assignment + { [var, value] }.

Then, the procedure **bt_search** is called recursively to complete this new partial assignment. If this is successful, the resulting assignment is a solution that is returned. Otherwise, the recursive call of **bt_search** will instead raise an exception. This exception is muted by the **check**-block that surrounds the call to **bt_search**. In that case, the **for**-loop generates a new possible **value** that can be assigned to the variable **var**. If all possible values have been tried and none was successful, the **for**-loop ends and the **backtrack**-statement is executed. Effectively, this statement raises an exception that is caught by one of the **check**-blocks that are above the **backtrack**-statement in the call stack.

We still need to discuss the implementation of the auxiliary procedure **is_consistent** shown in 3.9. This procedure takes a variable **var**, a **value**, a partial **Assignment** and a set of **Constraints**. It is assumed that **Assignment** is **partially consistent** with respect to the set **Constraints**, i.e. for every formula **f** occurring in **Constraints** such that

$$\text{vars}(f) \subseteq \text{dom}(\text{Assignment})$$

holds, the formula **f** evaluates to **true** given the **Assignment**. The purpose of **is_consistent** is to check, whether the extended assignment

```

1  is_consistent := procedure(var, value, Assignment, Constraints) {
2      NA := Assignment + { [var, value] };
3      return forall([F, Vars] in Constraints |
4          var in Vars && Vars <= domain(NA) => eval_constraint(NA, F)
5          );
6  };

```

Figure 3.9: Auxiliary procedures for the CSP solver shown in Figure 3.7

$$NA := \text{Assignment} \cup \{(\text{var}, \text{value})\}$$

that assigns **value** to the variable **var** is still partially consistent with **Constraints**. To this end, the **for**-loop iterates over all **Formula** in **Constraints**. However, we only have to check those **Formula** that contain the variable **var** and, furthermore, have the property that

$$\text{Vars}(\text{Formula}) \subseteq \text{dom}(NA),$$

i.e. all variables occurring in **Formula** need to have a value assigned in **NA**. The reasoning is as follows:

1. If **var** does not occur in **Formula**, then adding **var** to **Assignment** cannot change the result of evaluating **Formula** and as **Assignment** is assumed to be partially consistent with respect to **Formula**, **NA** is also partially consistent with respect to **Formula**.
2. If $\text{dom}(NA) \not\subseteq \text{Vars}(\text{Formula})$, then **Formula** can not be evaluated anyway.

If we use backtracking, we can solve the 8 queens problem in less than a second.

3.4 Constraint Propagation

Once we choose a value for a variable, this choice influences the values that are still available for other variables. For example, suppose we place the queen in row 1 in the second column, then no other queen can be placed in that column. Additionally, the queen in row 2 can then not be placed in any of the first three columns. It turns out that elaborating this idea can enhance the performance of backtracking search considerably. Figure 3.10 on page 52 shows an implementation of **constraint propagation**. This implementation is also able to handle **unary constraints**, i.e. constraints that contain only a single variable.

```

1  solve := procedure(CSP) {
2      [Variables, Values, Constrs] := CSP;
3      ValuesPerVar := { [v, Values] : v in Variables };
4      Annotated    := { [f, collectVars(f)] : f in Constrs };
5      UnaryConstrs := { [f, V] : [f, V] in Annotated | #V == 1 };
6      OtherConstrs := { [f, V] : [f, V] in Annotated | #V >= 2 };
7      check {
8          for ([f, V] in UnaryConstrs) {
9              var := arb(V);
10             ValuesPerVar[var] := solve_unary(f, var, ValuesPerVar[var]);
11         }
12         return bt_search({}, [Variables, ValuesPerVar, OtherConstrs]);
13     }
14 };

```

Figure 3.10: Constraint Propagation.

In order to implement constraint propagation, it is necessary to administer the values that can be used to instantiate the different variables separately, i.e. for every variable v we need to know which values are admissible for v . To this end, we need a dictionary that contains the set of possible values for every variable v . Initially, this dictionary assigns the set **Values** to every variable. Next, we take care of the unary constraints and shrink these sets accordingly. Then we solve the binary constraints and shrink these sets even more once we assign values to variables.

1. The procedure **solve** receives a **CSP**. This **CSP** is first split into its three components.
2. The first task of **solve** is to create the dictionary **ValuesPerVar**. Given a variable v , this dictionary assigns the set of values that can be used to instantiate this variable. Initially, this set is the same for all variables and is equal to **Values**.
3. Next, for every constraint f , the dictionary **Annotated** associates the variables occurring in a constraint with this constraint.
4. In order to solve the unary constraints we first have to find them. The set **UnaryConstrs** contains all those pairs $[f, V]$ from the set of annotated constraints **Annotated** such that the set of variables V contains just a single variable.
5. Similarly, the set **OtherConstrs** contains those constraints that involve two or more variables.
6. In order to solve the unary constraints, we iterate over all unary constraints and shrink the set of values associated with the variable occurring in the constraint as dictated by the constraint. This is done using the function **solve_unary**.
7. Then, the dictionaries **ValuesPerVar** and **OtherConstrs** are combined with the set of **Variables** into a triple that is used as the argument to the call of **bt_search**.

```

1  solve_unary := procedure(constraint, variable, Values) {
2      LegalValues := {};
3      for (value in Values) {
4          Assignment := { [variable, value] };
5          if (eval_constraint(Assignment, constraint)) {
6              LegalValues += { value };
7          }
8      }
9      if (LegalValues == {}) {
10         backtrack; // CSP is unsolvable
11     }
12     return LegalValues;
13 };

```

Figure 3.11: Implementation of **solve_unary**.

The function **solve_unary** shown in Figure 3.11 on page 53 takes a unary **constraint**, a **variable** and the set of **Values** that can be assigned to this **variable**. It returns the subset of values that can be substituted for the **variable** without violating the given **constraint**.

1. To achieve its goal, **solve_unary** iterates over all possible **Values**.
2. Next, for every **value** in the set **Values**, an **Assignment** is created that assign the **value** to the **variable**.
3. Then, the **constraint** is evaluated using the **Assignment**.

4. If the **constraint** is satisfied under this **Assignment**, the **value** is added to the set **LegalValues**, which is the set of values which satisfy the unary **constraint**.
5. If the set **LegalValues** is still empty after we have tried all possible **Values**, then the unary **constraint** is unsatisfiable and hence the CSP is unsolvable.
6. Otherwise, the set of **LegalValues** is returned.

```

1  bt_search := procedure(Assignment, CSP) {
2      [Variables, Values, Constraints] := CSP;
3      if (#Assignment == #Variables) {
4          return Assignment;
5      }
6      x := most_constrained_variable(Assignment, Values);
7      for (val in Values[x]) {
8          check {
9              if (is_consistent(x, val, Assignment, Constraints)) {
10                 NewVals := propagate(x, val, Assignment, Constraints, Values);
11                 CSP      := [Variables, NewVals, Constraints];
12                 return bt_search(Assignment + { [x, val] }, CSP);
13             }
14         }
15     }
16     backtrack;
17 };

```

Figure 3.12: Implementation of **bt_search**.

The procedure **bt_search** shown in Figure 3.12 on page 54 is called with a partial **Assignment** that is guaranteed to be consistent and a **CSP**. It tries to complete **Assignment** and thereby compute a solution of the **CSP**.

1. First, it decomposes the **CSP** into its components:
 - (a) **Values** is a dictionary assigning to every variable **v** the set of values that can be used to instantiate **v**. On recursive invocations of **bt_search** these sets will shrink.
 - (b) **Constraints** is a dictionary that assigns to every constraint **f** the set of variables occurring in **f**.
2. If the partial **Assignment** is already complete, i.e. if it assigns a value to every variable, then a solution to the given **CSP** has been found and this solution is returned.
3. Otherwise, we choose a variable **x** such that the number of values that can still be used to instantiate **x** is minimal. This variable is computed using the procedure **most_constrained_variable** that is shown in Figure 3.13 on page 55.
4. Next, all values that are still available for **x** are tried. Note that since **Values[x]** is, in general, smaller than the set of all values of the **CSP**, the **for**-loop in this version of backtracking search is more efficient than the version discussed in the previous section.
5. If it is consistent to assign **val** to the variable **x**, we propagate the consequences of this assignment using the procedure **propagate** shown in Figure 3.13 on page 55. This procedure updates the values that are still allowed for the variables of the **CSP** once the value **val** has been assigned to the variable **x**.
6. Finally, the partial variable **Assignment** is updated to include the assignment of **val** to **x** and the recursive call to **bt_search** tries to complete this new assignment.

```

1  most_constrained_variable := procedure(Assignment, Values) {
2      Unassigned := { [x, U] : [x, U] in Values | Assignment[x] == om };
3      minSize    := min({ #U : [x, U] in Unassigned });
4      return rnd({ x : [x, U] in Unassigned | #U == minSize });
5  };
6  propagate := procedure(x, v, Assignment, Constraints, Values) {
7      NewAssignment := Assignment + { [x, v] };
8      Values[x]     := { v };
9      for ([Formula, Vars] in Constraints | x in Vars && #Vars == 2) {
10         y := arb(Vars - { x });
11         if (!(y in domain(NewAssignment))) {
12             NewValues := Values[y];
13             for (w in NewValues) {
14                 A2 := NewAssignment + { [y, w] };
15                 if (!eval_constraint(A2, Formula)) {
16                     NewValues -= { w };
17                 }
18             }
19             if (NewValues == {}) { backtrack; }
20             Values[y] := NewValues;
21         }
22     }
23     return Values;
24 };

```

Figure 3.13: Constraint Propagation.

Figure 3.13 on page 55 show the implementation of the procedures `most_constrained_variable` and `propagate`. The procedure `most_constrained_variable` takes a partial `Assignment` and a dictionary `Values` returning the set of allowed values for all variables as input.

1. First, this procedure computes the set of `Unassigned` variables. For every variable `x` that has not yet been assigned a value in `Assignment` this set contains the pair `[x, U]`, where `U` is the set of allowable values for the variable `x`.
2. Next, it computes from all unassigned variables `x` the number of values `#U` that can be assigned to `x`. From these numbers, the minimum is computed and stored in `minSize`.
3. Finally, the set of variables that are maximally constrained is computed. These are all those variables `x` such that `Values[x]` has size `minSize`. From these variables, a random variable is returned.

The logic behind choosing a maximally constrained variables is that these variables are the most difficult to get right. If we have a partial assignment that is inconsistent, then we will discover this fact earlier if we try the most difficult variables first. This might save us a lot of unnecessary computation.

The function `propagate` takes the following inputs:

- (a) `x` is a variable.
- (b) `v` is a value that is assigned to the variable `x`.
- (c) `Assignment` is a partial assignment that contains only assignments for variables that are different from `x`.
- (d) `Constraints` is a set of annotated constraints, i.e. this set contains pairs of the form `[Formula, Vars]`, where `Formula` is a constraint and `Vars` is the set of variables occurring in `Formula`.

(e) **Values** is a dictionary assigning sets of values to the different variables.

The function **propagate** updates the dictionary **Values** by taking into account the consequences of assigning the value **v** to the variable **x**.

1. The partial **Assignment** is updated to **NewAssignment**, where **NewAssignment** maps **x** to **v**.
2. As **x** now has the value of **v**, the corresponding entry in the dictionary **Values** is changed accordingly.
3. Next, **propagate** iterates over all binary **Constraints** such that **x** occurs in **Formula**.
4. The remaining variable occurring in **Formula** is called **y**. There can be only one other variable in **Formula** because the test **#Vars == 2** assures that **Formula** is a binary constraint.
5. If the variable **y** has already been assigned a value, then there is nothing to do because we already know that the assignment **NewAssignment** is consistent. This is ensured by the previous call to the function **is_consistent** in the body of the procedure **bt_search**.
However, if **y** is still unassigned, it is necessary to update **Values[y]**.
6. To this end, all values **w** in **Values[y]** are tested. If it turns out that assigning the value **w** to the variable **y** violates the constraint **Formula**, then **Values[y]** must not contain the value **w** and, accordingly, this value is removed from **Values[y]**.
7. If it turns out that **Values[y]** has been reduced to the empty set, then this means that the partial assignment **NewAssignment** can not be completed. Hence, the search has to **backtrack**.

I have tested the program described in this section using the n queens puzzle. I have found that the time needed to solve ten instances of 16 queens problem went down from 72 seconds to just 5 seconds. The procedure is even able to solve the 32 queens problem, taking about 4 seconds on average, while the version of backtracking search that does not use constraint propagation took more than 3 minutes on average.

Exercise 7: There are many different versions of the *zebra puzzle*. The version below is taken from *Wikipedia*. The puzzle reads as follows:

1. There are five houses.
2. The Englishman lives in the red house.
3. The Spaniard owns the dog.
4. Coffee is drunk in the green house.
5. The Ukrainian drinks tea.
6. The green house is immediately to the right of the ivory house.
7. The Old Gold smoker owns snails.
8. Kools are smoked in the yellow house.
9. Milk is drunk in the middle house.
10. The Norwegian lives in the first house.
11. The man who smokes Chesterfields lives in the house next to the man with the fox.
12. Kools are smoked in the house next to the house where the horse is kept.
13. The Lucky Strike smoker drinks orange juice.
14. The Japanese smokes Parliaments.
15. The Norwegian lives next to the blue house.

16. Who drinks water?
17. Who owns the zebra?

In order to solve the puzzle, we also have to know the following facts:

1. Each of the five houses is painted in a **different** colour.
2. The inhabitants of the five houses are of **different** nationalities,
3. they own **different** pets,
4. they drink **different** beverages, and
5. they smoke **different** brands of cigarettes.

Formulate the zebra puzzle as a constraint satisfaction problem and solve the puzzle using the program discussed in this section. You should also try to solve the puzzle using the program given in the previous section. Compare the results. \diamond

3.5 Consistency Checking

So far, the constraints in the constraints satisfaction problems discussed are either **unary constraints** or **binary constraints**: A **unary** constraint is a constraint f such that the formula f contains only one variable, while a **binary** constraint contains two variables. If we have a constraint satisfaction problem that involves also constraints that mention more than two variables, then the constraint propagation shown in the previous section does not apply. For example, consider the **cryptarithmic puzzle** shown in Figure 3.14 on page 57. The idea is that the letters “S”, “E”, “N”, “D”, “M”, “O”, “R”, “Y” are interpreted as variables ranging over the set of decimal digits, i.e. these variables can take values in the set $\{0, 1, 2, 3, 4, 5, 6, 7, 8, 9\}$. Then, the string “SEND” is interpreted as a decimal number, i.e. it is interpreted as the number

$$S \cdot 10^3 + E \cdot 10^2 + N \cdot 10^1 + D \cdot 10^0.$$

The strings “MORE” and “MONEY” are interpreted similarly. To make the problem interesting, the assumption is that different variables have different values. Furthermore, the decimals at the beginning of a number should be different from 0.



$$\begin{array}{r} \text{SEND} \\ + \text{MORE} \\ \hline \text{MONEY} \end{array}$$

Figure 3.14: A cryptarithmic puzzle

A naïve approach to solve this problem would be to code it as a constraint satisfaction problem that has, among others, the following constraint:

$$(S \cdot 10^3 + E \cdot 10^2 + N \cdot 10 + D) + (M \cdot 10^3 + O \cdot 10^2 + R \cdot 10 + E) = M \cdot 10^4 + O \cdot 10^3 + N \cdot 10^2 + E \cdot 10 + Y.$$

The problem with this constraint is that it involves far too many variables. As this constraint can only be checked when all the variables have values assigned to them, the backtracking search would essentially boil down to a mere brute force search. In order to do better, we have to perform the addition in Figure 3.14 column by column, just as it is taught in elementary school. Figure 3.15 on page 58 shows how this can be implemented in SETLX.

Notice that we have introduced three additional variables “C1”, “C2”, “C3”. These variables serve as the **carry digits**. For example, “C1” is the carry digit that we get when we do the addition of the last places of the

```

1  createCSP := procedure() {
2      Variables := {"S", "E", "N", "D", "M", "O", "R", "Y", "C1", "C2", "C3"};
3      Values    := { 0 .. 9 };
4      Constraints := allDifferent({ "S", "E", "N", "D", "M", "O", "R", "Y" });
5      Constraints += { '(D + E) % 10 == Y', '(D + E) \ 10 == C1',
6                      '(N + R + C1) % 10 == E', '(N + R + C1) \ 10 == C2',
7                      '(E + O + C2) % 10 == N', '(E + O + C2) \ 10 == C3',
8                      '(S + M + C3) % 10 == O', '(S + M + C3) \ 10 == M'
9                      };
10     Constraints += { "S != 0", "M != 0" };
11     return [Variables, Values, Constraints];
12 };

```

Figure 3.15: Formulating “SEND + MORE = MONEY” as a CSP.

two numbers, i.e. we have

$$D + E = C1 \cdot 10 + Y.$$

This equation still contains four variables. We can reduce it to two equations that each involve only three variables as follows:

$$(D + E) \% 10 = Y \quad \text{and} \quad (D + E) \setminus 10 = C1.$$

Here, the symbol “\” denotes [integer division](#), e.g. we have $7 \setminus 3 = 2$. If we try to solve the cryptarithmic puzzle as coded in Figure 3.15 on page 58 using the constraint solver developed in the previous section, we will be disappointed. The reason is that most constraints involve three variables and therefore the constraint propagation developed in the previous section is of no help. However, we can solve the problem in a few seconds if we add the following constraints for the variables “C1”, “C2”, “C3”:

$$"C1 < 2", "C2 < 2", "C3 < 2".$$

Although these constraints are certainly true, the problem with this approach is that we would prefer if our constraint solver would be able to figure out these constraints by itself. After all, since D and E are both less than 10, their sum is obviously less than 20 and hence the carry $C1$ has to be less than 2. This line of reasoning is known as [consistency maintenance](#): Assume that the formula f is a constraint and the set of variables occurring in f has the form

$$\text{Var}(f) = \{x\} + R \quad \text{where } x \notin R,$$

i.e. the variable x occurs in the constraint f and, furthermore, R is the set of all variables occurring in f that are different from x . Furthermore, assume that we have a dictionary `AllowableValues` such that for every variable y , the dictionary entry `AllowableValues[y]` is the set of values that can be substituted for the variable y . Now a value v is consistent for x with respect to the constraint f iff the partial assignment $\{x, v\}$ can be extended to an assignment A satisfying the constraint f , i.e. for every variable $y \in R$ we have to find a value $w \in \text{AllowableValues}[y]$ such that the resulting assignment A satisfies the equations

$$\text{evaluate}(f, A) = \text{true}.$$

Here, the function `evaluate` takes a formula f and an assignment A and evaluates f using the assignment A . Now, [consistency maintenance](#) works as follows.

1. The dictionary `AllowableValues` is initialized as follows:

$$\text{AllowableValues}[x] := \text{Values} \quad \text{for all } x \in \text{Variables},$$

i.e. initially every variable x can take any value from the set of `Values`.

2. Next, the set `UncheckedVariables` is initialized to the set of all `Variables`:

`UncheckedVariables := Variables.`

3. As long as the set `UncheckedVariables` is not empty, we remove one variable x from this set:

$x := \text{from}(\text{UncheckedVariables})$

4. We iterate over all constraints f such that x occurs in f .

- (a) For every value $v \in \text{AllowableValues}[x]$ we check whether v is consistent with f .
- (b) If v is not consistent with f , then v is removed from `AllowableValues` $[x]$. Furthermore, if R is the set of all variables occurring in f that are different from x , then this set R is added to the set of `UncheckedVariables`.

5. Once `UncheckedVariables` is empty, the algorithm terminates. Otherwise, we jump back to step 3 and remove the next variable from the set `UncheckedVariables`.

The algorithm terminates as every iteration removes either a variable from the set `UncheckedVariables` or it removes a value from one of the sets `AllowableValues` $[y]$ for some variable y . Although the set `UncheckedVariables` can grow during the algorithm, the union

$$\bigcup_{x \in \text{Vars}} \text{AllowableValues}[x]$$

can never grow: Every time the set `UncheckedVariables` grows, for some variable x the set `AllowableValues` $[x]$ shrinks. As the sets `AllowableValues` $[x]$ are finite for all variables x , the set `UncheckedVariables` can only grow a finite number of times. Once the set `UncheckedVariables` does not grow any more, every iteration of the algorithm removes one variable from this set and hence the algorithm terminates eventually.

```

1  enforceConsistency := procedure(rw ValuesPerVar,
2                                Var2Formulas, Annotated, Connected)
3  {
4      UncheckedVars := domain(Var2Formulas);
5      while (UncheckedVars != {}) {
6          variable := from(UncheckedVars);
7          Constraints := Var2Formulas[variable];
8          Values := ValuesPerVar[variable];
9          RemovedVals := {};
10         for (f in Constraints) {
11             OtherVars := Annotated[f] - { variable };
12             for (value in Values) {
13                 if(!existsValue(variable,value,f,OtherVars,ValuesPerVar)) {
14                     RemovedVals += { value };
15                     UncheckedVars += Connected[variable];
16                 }
17             }
18         }
19         Remaining := Values - RemovedVals;
20         if (Remaining == {}) { backtrack; }
21         ValuesPerVar[variable] := Remaining;
22     }
23 };

```

Figure 3.16: Consistency maintenance in SETLX.

Figure 3.16 on page 59 shows how consistency maintenance can be implemented in SETLX. The procedure `enforceConsistency` takes four arguments.

- (a) **ValuesPerVar** is a dictionary associating the set of possible values with each variable.
- (b) **Var2Formulas** is a dictionary. For every variable v , **Var2Formulas**[v] is the set of those constraints f such that v occurs in f .
- (c) **Annotated** is the set of annotated constraints, i.e. this set contains pairs of the form $\langle f, V \rangle$ where f is a constraint and V is the set of all variables occurring in f . This last argument is needed only for efficiency: In order to avoid computing the set V of variables occurring in a constraint f every time the constraint f is encountered, we compute these sets at the beginning of our computation and store them in the dictionary **Annotated**.
- (d) **Connected** is a dictionary that takes a variable x and returns the set V of all variables that are related to x via a common constraint f , i.e. we have $y \in \text{Connected}[x]$ if there exists a constraint f such that both x and y occur in f and, furthermore, $x \neq y$.

The procedure **enforceConsistency** modifies the dictionary **ValuesPerVar** so that once the procedure has terminated, for every variable x the set **ValuesPerVar**[x] is consistent with the constraints for x . The implementation works as follows:

1. Initially, all variables need to be checked for consistency. Therefore, **UncheckedVars** is defined to be the set of all variables that occur in any of the constraints.
2. The **while**-loop iterates as long as there are still variables x left in **UncheckedVars** such that the consistency of **ValuesPerVar**[x] has not been established.
3. Next, a **variable** is selected and removed from **UncheckedVars**.
4. **Constraints** is the set of all constraints f such that this **variable** occurs in f .
5. **Values** is the set of those values that can be assigned to **variable**.
6. **RemovedVals** is the subset of those values that are found to be inconsistent with some constraint.
7. We iterate over all constraints $f \in \text{Constraints}$.
8. **OtherVars** is the set of variables occurring in f that are different from the chosen **variable**.
9. We iterate over all **value** $\in \text{Values}$ that can be substituted for **variable** and check whether **value** is consistent with f . To this end, we need to find values that can be assigned to the variables in the set **OtherVars** such that f evaluates as **true**. This is checked using the function **existsValue**.
10. If we do not find such values, then **value** is inconsistent for **variable** w.r.t. f and needs to be removed from the set **ValuesPerVar**[**variable**]. Furthermore, all variables that are connected to **variables** have to be added to the set **UncheckedVars**. The reason is that once a value is removed for **variable**, the value assigned to another variable y occurring in a constraint that mentions both **variable** and y might now become inconsistent.
11. If there are no consistent values for **variable** left, we have to backtrack.

Figure 3.17 on page 61 shows the implementation of the function **existsValue** that is used in the implementation of **enforceConsistency**. This procedure is called with five arguments.

- (a) **var** is variable.
- (b) **val** is a value that is to be assigned to **var**.
- (c) f is a constraint such that **var** occurs in f
- (d) **Vars** is the set of all those other variables occurring in f , i.e. the set of those variables that occur in f but that are different from **var**.

```

1  existsValue := procedure(var, val, f, Vars, ValuesPerVar) {
2      AllVars := { var } + Vars;
3      return exists( A in createAllAssignments(Vars, ValuesPerVar) |
4                  eval_constraint(A + { [var, val] }, f, AllVars) );
5  };
6  createAllAssignments := procedure(Vars, ValuesPerVar) {
7      if (Vars == {}) {
8          return { {} }; // set containing empty assignment
9      }
10     var      := from(Vars);
11     Values    := ValuesPerVar[var];
12     Assignments := createAllAssignments(Vars, ValuesPerVar);
13     return { { [var, val] } + A : val in Values, A in Assignments };
14 };

```

Figure 3.17: The implementation of `existsValue`.

(e) `ValuesPerVar` is a dictionary associating the set of possible values with each variable.

The procedure checks whether the partial assignment $\{\langle \text{var}, \text{val} \rangle\}$ can be extended so that the constraint `f` is satisfied. To this end it needs to create the set of all possible assignments. This set is generated using the function `createAllAssignments`. This function gets a set of variables `Vars` and a dictionary that assigns to every variable `var` in `Vars` the set of values that might be assigned to `var`.

1. If the set of variables `Vars` is empty, the empty set can serve as a dictionary that assigns a value to every variable in `Vars`.
2. Otherwise, we remove a variable `var` from `Vars` and get the set of `Values` that can be assigned to `var`.
3. Recursively, we create the set of all `Assignments` that associate values with the remaining variables.
4. Finally, the set of all possible assignments is the set of all combinations of assigning a value $\text{val} \in \text{Values}$ to `var` and assigning the remaining variables according to an assignment $A \in \text{Assignments}$.

On one hand, consistency checking creates a lot of overhead.² Therefore, it might actually slow down the solution of some constraint satisfaction problems that are easy to solve using just backtracking and propagation. On the other hand, many difficult constraint satisfaction problems can not be solved without consistency checking.

3.6 Local Search

There is another approach to solve constraint satisfaction problems. This approach is known as [local search](#). The basic idea is simple: Given a constraint satisfaction problem \mathcal{C} of the form

$$\mathcal{P} := (\text{Variables}, \text{Values}, \text{Constraints}),$$

local search works as follows:

1. Initialize the values of the variables in `Variables` randomly.
2. If all `Constraints` are satisfied, return the solution.
3. For every $x \in \text{Variables}$, count the number of unsatisfied constraints that involve the variable x .

² To be fair, the implementation shown in this section is far from optimal. In particular, by remembering which combinations of variables and values work for a given formula, the overhead can be reduced greatly. I have refrained from implementing this optimization because I did not want the code to get too complex.

4. Set **maxNum** to be the maximum of these numbers, i.e. **maxNum** is maximal number of unsatisfied constraints for any variable.
5. Compute the set **maxVars** of those variables that have **maxNum** unsatisfied constraints.
6. Randomly choose a variable x from the set **maxVars**.
7. Find a value $d \in \mathbf{Values}$ such that by assigning d to the variable x , the number of unsatisfied constraints for the variable x is minimized.
If there is more than one value d with this property, choose the value d randomly from those values that minimize the number of unsatisfied constraints.
8. Goto step 2 and repeat until either a solution is found or the sun rises in the west.

```

1  solve := procedure(n) {
2      Queens := [];
3      for (row in [1 .. n]) {
4          Queens[row] := rnd({1 .. n});
5      }
6      iteration := 0;
7      while (true) {
8          Conflicts := { [numConflicts(Queens, row), row] : row in [1 ..n] };
9          [maxNum, _] := last(Conflicts);
10         if (maxNum == 0) {
11             return Queens;
12         }
13         if (iteration % 10 != 0) { // avoid infinite loops
14             row := rnd({ row : [num, row] in Conflicts | num == maxNum });
15         } else {
16             row := rnd({ 1 .. n });
17         }
18         Conflicts := {};
19         for (col in [1 .. n]) {
20             Board := Queens;
21             Board[row] := col;
22             Conflicts += { [numConflicts(Board, row), col] };
23         }
24         [minNum, _] := first(Conflicts);
25         Queens[row] := rnd({ col : [num, col] in Conflicts | num == minNum });
26         iteration += 1;
27     }
28 };

```

Figure 3.18: Solving the n queens problem using local search.

Figure 3.18 on page 62 shows an implementation of these ideas in SETLX. Instead of solving an arbitrary constraint satisfaction problem, the program solves the n queens problem. We proceed to discuss this program line by line.

1. The procedure **solve** takes one parameter n , which is the size of the chess board. If the computation is successful, **solve**(n) returns a list of length n . Lets call this list **Queens**. For every row $r \in \{1, \dots, n\}$, the value **Queens**[r] specifies that the queen that resides in row r is positioned in column **Queens**[r].

2. The **for** loop initializes the positions of the queens to random values from the set $\{1, \dots, n\}$. Effectively, for every row on the chess board, this puts a queen in a random column.
3. The variable **iteration** counts the number of times that we need to reassign a queen in a given row.
4. All the remaining statements are surrounded by a **while** loop that is only terminated once a solution has been found.
5. The variable **Conflicts** is a set of pairs of the form $[c, r]$, where c is the number of times the queen in row r is attacked by other queens. Hence, c is the same as the number of unsatisfied conflicts for the variable specifying the column of the queen in row r .
6. **maxNum** is the maximum of the number of conflicts for any row.
7. If this number is 0, then all constraints are satisfied and the list **Queens** is a solution to the n queens problem.
8. Otherwise, we compute those rows that exhibit the maximal number of conflicts. From these rows we select one **row** arbitrarily.
9. The reason for enclosing the assignment to **row** in an **if** statement is explained later. On a first reading of this program, this **if** statement should be ignored.
10. Now that we have identified the **row** where the number of conflicts is biggest, we need to reassign **Queens[row]**. Of course, when reassigning this variable, we would like to have fewer conflicts after the reassignment. Hence, we test all columns to find the best column that can be assigned for the queen in the given **row**. This is done in a **for** loop that runs over all possible columns. The set **Conflicts** that is maintained in this loop is a set of pairs of the form $[k, c]$ where k is the number of times the queen in **row** would be attacked if it would be placed in column c .
11. We compute the minimum number of conflicts that is possible for the queen in **row** and assign it to **minNum**.
12. From those columns that minimize the number of violated constraints, we choose a column randomly and assign it for the specified **row**.

There is a technical issue, that must be addressed: It is possible there is just one row that exhibits the maximum number of conflicts. It is further possible that, given the placements of the other queens, there is just one optimal column for this row. In this case, the procedure **solve** would loop forever. To avoid this case, every 10 iterations we pick a random row to change.

```

1  numConflicts := procedure(Queens, row) {
2      n      := #Queens;
3      result := 0;
4      for (r in {1 .. n} | r != row) {
5          if ( Queens[r] == Queens[row]           ||
6              r - Queens[r] == row - Queens[row] ||
7              r + Queens[r] == row + Queens[row]
8          )
9              { result += 1; }
10     }
11     return result;
12 };

```

Figure 3.19: The procedure **numConflicts**.

The procedure `numConflicts` shown in Figure 3.19 on page 63 implements the function `numConflicts`. Given a board `Queens` that specifies the positions of the queens on the board and a `row`, this function computes the number of ways that the queen in `row` is attacked by other queens. If all queens are positioned in different rows, then there are only three ways left that a queen can be attacked by another queen.

1. The queen in row `r` could be positioned in the same column as the queen in `row`.
2. The queen in row `r` could be positioned in the same falling or rising diagonal as the queen in `row`. These diagonals are specified by the linear equations given in line 6 and 7 of Figure 3.19.

Using the program discussed in this section, the n queens problem can be solved for a $n = 1000$ in 30 minutes. As the memory requirements for local search are small, even much higher problem sizes can be tackled if sufficient time is available. Hence, it seems that local search is much better than the algorithms discussed previously. However, we have to note that local search is *incomplete*: If a constraint satisfaction problem \mathcal{P} has no solution, then local search loops forever. Therefore, in practice a dual approach is used to solve a constraint satisfaction problem. The constraint solver starts two threads: The first search does local search, the second thread tries to solve the problem via some refinement of backtracking. The first thread that terminates wins. The resulting algorithm is complete and, for a solvable problem, will have similar performance as local search. If the problem is unsolvable, this will *eventually* be discovered by backtracking. Note, however, that the constraint satisfaction problem is NP-complete. Hence, it is unlikely that there is an efficient algorithm that works always. However, many practically occurring constraint satisfaction problems can be solved today in a reasonably short time.

Chapter 4

Playing Games

One major success for the field of artificial intelligence happened in 1997 when the chess-playing computer **Deep Blue** was able to beat the World Chess Champion **Garry Kasparov** by $3\frac{1}{2} - 2\frac{1}{2}$. While **Deep Blue** was based on special hardware, according to the **computer chess rating list** of the 18th of March 2017, the chess program **Stockfish** runs on ordinary desktop computers and has an **Elo rating** of 3390. To compare, according to the **Fide** list of March 2017, the current World Chess Champion **Magnus Carlsen** has an Elo rating of just 2838. Hence, he wouldn't stand a chance to win a game against Stockfish. More recently, the computer program **AlphaGo** was able to beat **Lee Sedol**, who is currently¹ considered to be the second best **go** player in the world. Besides go and chess, there are many other games where today the performance of a computer exceeds the performance of human players. To name just one more example, at the beginning of 2017 the program **Libratus** was able to **beat** four professional poker players resoundingly.

In this chapter we want to investigate how a computer can play a game. To this end we define a **game** \mathcal{G} as a six-tuple

$$\mathcal{G} = \langle \text{States}, s_0, \text{Players}, \text{nextStates}, \text{finished}, \text{utility} \rangle$$

where the components are interpreted as follows:

1. **States** is the set of all possible **states** of the game.
2. $s_0 \in \text{startState}$ is the **start state**.
3. **Players** is the set of **players** of the game.
4. **nextStates** is a function that takes a state $s \in \text{States}$ and a player $p \in \text{Players}$ and returns the set of states that can be reached if p has to make a move in the state s . Hence, the signature of **nextStates** is given as follows:

$$\text{nextStates} : \text{States} \times \text{Players} \rightarrow 2^{\text{States}}.$$

5. **finished** is a function that takes a state s and decides whether the game is finished. Therefore

$$\text{finished} : \text{States} \rightarrow \mathbb{B}.$$

Here, \mathbb{B} is the set of Boolean values, i.e. we have $\mathbb{B} := \{\text{true}, \text{false}\}$.

Using the function **finished**, we define the set **TerminalStates** as the set of those states such that the game has finished, i.e. we define

$$\text{TerminalStates} := \{s \in \text{States} \mid \text{finished}(s)\}.$$

6. **utility** is a function that takes a state $s \in \text{TerminalStates}$ and a player $p \in \text{Players}$. It returns the **value** that the game has for player p . In all of our examples, this value will be an element from the set $\{-1, 0, +1\}$. The value -1 indicates that player s has lost the game, if the value is $+1$ the player p has

¹This assessment is of March 2017.

won the game and if this value is 0 then the game is drawn. Hence the signature of `utility` is

$$\text{utility} : \text{TerminalStates} \times \text{Players} \rightarrow \{-1, 0, +1\}.$$

Exercise 8: The definition given above does not capture all types of games. In particular, the definition only captures **deterministic games**: A game is called **deterministic** iff there is no randomness involved. Games like chess and go are certainly deterministic. However, other games like **dice chess** involve randomness. In dice chess a pair of dice is rolled at the beginning of each turn. The outcome of this roll determines which pieces may be moved. Extend the definition of a game so that also a games like **dice chess** can be described. \diamond

In this chapter we will only consider so called **two person zero sum games**. This means that the set `Players` has exactly two elements. If we call these players A and B, i.e. if we have

$$\text{Players} = \{A, B\},$$

then the game is called a **zero sum game** iff we have

$$\forall s \in \text{TerminalStates} : \text{utility}(s, A) + \text{utility}(s, B) = 0,$$

i.e. the losses of player A are compensated by the wins of player B and vice versa. Games like **go** and **chess** are two person zero sum games. We proceed to discuss an example of a game.

4.1 Tic-Tac-Toe

The game **tic-tac-toe** is played on a square board of size 3×3 . On every turn, one player puts an “X” on one of the free squares of the board, while the other player puts an O onto a free square when it is his turn. If the first player manages to place three Xs in a row, column, or diagonal, she has won the game. Similarly, if the other player manages to put three Os in a row, column, or diagonal, this player is the winner. Otherwise, the game is drawn. Figure 4.1 on page 67 shows a SETLX implementation of tic-tac-toe.

1. The function `players` returns the set `Players`. Traditionally, the players in tic-tac-toe are called “X” and “O”.
2. The function `startState` returns the start state, which is an empty board. States are represented as list of lists. The entries in these lists are the characters “X”, “O”, and “ ”. As the `StartState` is the empty board, it is represented as as follows:

```
[ [ " ", " ", " " ],
  [ " ", " ", " " ],
  [ " ", " ", " " ]
].
```

The function `startState` receives one optional argument `n`. By default `n` is equal to 3. This argument specifies the size of the board. Although traditionally tic-tac-toe is played on a 3×3 board, it can also be played on a bigger board.

3. The function `nextStates` takes a `State` and a `player` and computes the set of states that can be reached from `State` if `player` is to move next. To this end, it first computes the set of **empty** positions. Every position is represented as pair of the form `[row, col]` where `row` specifies the row and `col` specifies the column of the position. A position is **empty** iff

$$[\text{row}, \text{col}] = \text{ " " }.$$

The computation of the empty position has been sourced out to the function `find_empty`. The function `nextStates` then iterates over the set of empty positions. For every empty position `[row, col]` it creates a new state `NextState` that results from the current `State` by putting the mark of `player` in this position. The resulting states are collected in the set `Result` and returned.

4. The function `find_empty` takes a `State` and returns the set of empty positions.

```

1  players := procedure() { return { "X", "O" }; };
2  startState := procedure(n) {
3      L := [1 .. n];
4      return [ [ " " : col in L ] : row in L ];
5  };
6  nextStates := procedure(State, player) {
7      Empty := find_empty(State);
8      Result := {};
9      for ([row, col] in Empty) {
10         NextState := State;
11         NextState[row][col] := player;
12         Result += { NextState };
13     }
14     return Result;
15 };
16 find_empty := procedure(State) {
17     n := #State;
18     L := [1 .. n];
19     return { [row, col] : row in L, col in L | State[row][col] == " " };
20 };
21 utility := procedure(State, player) {
22     Lines := all_lines(State);
23     for (Line in Lines) {
24         if (#Line == 1 && Line != { " " }) {
25             if (Line == { player }) { return 1; } else { return -1; }
26         }
27     }
28     if (find_empty(State) == {}) { return 0; } // else return om
29 };
30 all_lines := procedure(State) {
31     n := #State;
32     L := [1 .. n];
33     Lines := { { State[row][ col] : col in L } : row in L };
34     Lines += { { State[row][ col] : row in L } : col in L };
35     Lines += { { State[idx][ idx] : idx in L } };
36     Lines += { { State[idx][-idx] : idx in L } };
37     return Lines;
38 };
39 finished := procedure(State) {
40     if (utility(State, "X") != om) { return true; } else { return false; }
41 };

```

Figure 4.1: A SETLX description of tic-tac-toe.

5. The function `utility` takes a `State` and a `player`. If the game is finished in the given `State`, it returns the value that this `State` has for the current `player`. If the game is not yet decided, Ω is returned instead.

In order to achieve its goal, the procedure first computes the set of all [line-sets](#). Given a `State`, a [line-set](#) is a set of those markers that either form a horizontal, vertical, or diagonal line in `State`, e.g. the set

$$\{\text{State}[1, 1], \text{State}[2, 2], \text{State}[3, 3]\}$$

is the set of entries in a diagonal line. The game is decided if all entries in this set are either “X” or “O”.

In this case, the set has exactly one element which is different from the blank. If this element is the same as **player**, then the game is won by **player**, otherwise it is lost.

Finally, if there are no empty squares left, then the game is a draw.

6. The auxiliary procedure **all_lines** takes a **State** and computes the set **Lines** of all **line-sets**.

- (a) For any $\text{row} \in \{1, 2, 3\}$ the set

$$\{ \text{State}[\text{row}][\text{col}] : \text{col} \text{ in } [1..3] \}$$

forms a horizontal line.

- (b) Likewise, for any $\text{col} \in \{1, 2, 3\}$ the set

$$\{ \text{State}[\text{row}][\text{col}] : \text{row} \text{ in } [1..3] \}$$

forms a vertical line.

- (c) Given that the top row is indexed with 1, and the bottom row is row number 3, the set

$$\{ \text{State}[\text{idx}][\text{idx}] : \text{idx} \text{ in } [1..3] \};$$

is the falling diagonal.

- (d) Finally, the set

$$\{ \text{State}[\text{idx}][n - (\text{idx} - 1)] : \text{idx} \text{ in } [1..3] \}$$

is the rising diagonal.

7. The procedure **finished** takes a **State** and checks whether the game is finished. To this end it computes the **utility** of the state for the player “X”. If this **utility** is different from Ω , the game is finished. Note that it does make no difference whether we take the utility of the state for the player “X” or for the player “O”: If the game is finished for “X”, then it is also finished for “O” and vice versa.

4.2 The Minimax Algorithm

Having defined the notion of a game, our next task is to come up with an algorithm that can play a game. The easiest algorithm that works is the **minimax algorithm**. This algorithm is based on the notion of the **value** of a state. To this end, we define a function

$$\text{value} : \text{States} \times \text{Players} \rightarrow \{-1, 0, +1\}$$

that takes a state $s \in \text{States}$ and a player $p \in \text{Players}$ and returns the value of s provided both the player p and his opponent play optimally. The easiest way to define this function is via recursion. The base case is simple:

$$\text{finished}(s) \rightarrow \text{value}(s, p) = \text{utility}(s, p).$$

If the game is not yet finished, assume that player o is the opponent of player p . Then we define

$$\neg \text{finished}(s) \rightarrow \text{value}(s, p) = \max(\{-\text{value}(n, o) \mid n \in \text{nextStates}(s, p)\}).$$

The reason is that if the game is not finished yet, the player p has to evaluate all possible moves. From these, the player p will choose the move that maximizes the value of the game. In order to do so, the player p computes the set **nextStates**(s, p) of all states that can be reached from the state s in any one move of the player p . Now if n is a state that results from player p making some move, then it is the turn of the other player o to make a move. Hence, in order to evaluate the state n , we have to call the function **value** recursively as **value**(n, o). As we are dealing with zero sum games here, we have **value**(n, p) = $-\text{value}(n, o)$. Figure 4.2 on page 69 shows an implementation of this strategy.

1. The implementation of the function **value** follows the reasoning outlined above. However, note that we have implemented **value** as a **cachedProcedure**, i.e. as a procedure that **memorizes** its results. Hence, when the function **value** is called a second time with the same pair of arguments, it does not recompute

```

1  value := cachedProcedure(State, player) {
2      if (finished(State)) {
3          return utility(State, player);
4      }
5      other := arb(players() - { player });
6      return max({ -value(ns, other) : ns in nextStates(State, player) });
7  };
8  best_move := procedure(State, player) {
9      NS      := nextStates(State, player);
10     other    := arb(players() - { player });
11     bestVal  := value(State, player);
12     return [bestVal, rnd({ns : ns in NS | -value(ns,other) == bestVal})];
13 };
14 play_game := procedure(n := 3) {
15     State := startState(n);
16     print(stateToString(State));
17     while (true) {
18         [val, State] := best_move(State, "O");
19         print("For me, the game has the value $val$. My move:");
20         print(stateToString(State));
21         if (final_msg(State)) { return; }
22         State := getMove(State);
23         print(stateToString(State));
24         if (final_msg(State)) { return; }
25     }
26 };

```

Figure 4.2: The Minimax algorithm.

the value but rather the value is looked up in a so called [cache](#) that stores all previous results computed by the function `value`. To understand why this is important, let us consider how many states would be explored in the case of tic-tac-toe if we would not use the idea of memorizing previous results, a technique which is known as [memoization](#). In this case, we have 9 moves for player “O” from the start state, then 8 moves for player “X”, then again 7 moves for player “X”. If we disregard the fact that some games are decided after fewer than 9 moves, the function `value` needs to consider

$$9 \cdot 8 \cdot 7 \cdot \dots \cdot 2 \cdot 1 = 9! = 362880$$

moves. However, if we count the number of possibilities of putting 5 “O”s and 4 “X”s on a 3×3 board, we see that there are only

$$\binom{9}{5} = \frac{9!}{5! \cdot 4!} = 126$$

possibilities, i.e. there are a factor of $5! \cdot 4! = 2880$ less states to evaluate!

2. The function `best_move` takes a `State` and a `player` and returns a pair $\langle v, s \rangle$ where s is a state that is optimal for the `player` and such that s can be reached in one step from `State`. Furthermore, v is the value of this state.
 - (a) To this end, it first computes the set `NS` of all states that can be reached from the given `State` in one step if `player` is to move next.
 - (b) Since there are only two players, the opponent `other` is found by subtracting `player` from the set of all players.

- (c) `bestValue` is the best value that `player` can achieve in the given `State`.
 - (d) The function returns randomly one of those states `ns` \in `NS` such that the value of `ns` is optimal, i.e. is equal to `bestValue`. We use randomization here since we want to have more interesting games. If we would always choose the first state that achieves the best values, then our program would always make the same move in a given state. Hence, playing the program would get boring much sooner.
3. The function `play_game` is used to play a game.
- (a) Initially, `State` is the `startState`.
 - (b) As long as the game is not finished, the procedure keeps running.
 - (c) We assume that the computer is player “O” and that the computer goes first. Hence, the function `best_move` is used to compute the state that results from the best move of player “O”.
 - (d) After that, it is checked whether the game is finished.
 - (e) If the game is not yet finished, the human is asked to make its move via the function `getMove` that takes a `State`, displays it, and asks the user to enter a move. The state resulting from this move is then returned and displayed.
 - (f) Next, we have to check whether the game is finished after our move has been executed.
 - (g) The `while`-loop keeps iterating until the game is finished. We do not have to put a test into the condition of this `while`-loop as we call the function `final_msg(State)` every time that a new `State` has been reached. This function returns `true` iff `finished(State)` is true. Additionally, it prints a message when the game has finished.

4.3 α - β -Pruning

The efficiency of the minimax algorithm can be improved if we provide two additional arguments to the function `value`. Traditionally, these arguments are called α and β . In order to be able to distinguish between the old function `value` and its improved version, we call the improved version `alphaBeta`. The idea is that the function `alphaBeta` and the function `value` are related by the following requirements:

1. As long as `value(s, p)` is between α and β , the function `alphaBeta` computes the same result as the function `value`, i.e. we have

$$\alpha \leq \text{value}(\text{State}, \text{player}) \leq \beta \rightarrow \text{alphaBeta}(s, p, \alpha, \beta) = \text{value}(s, p).$$
2. If `value(State, player) $\leq \alpha$` , we require that the value returned by `alphaBeta` is less than or equal to α , i.e. we have

$$\text{value}(s, p) \leq \alpha \rightarrow \text{alphaBeta}(s, p, \alpha, \beta) \leq \alpha.$$
3. Similarly, if `value(State, player) $\geq \beta$` , we require that the value returned by `valueAlphaBeta` is bigger than or equal to β , i.e. we have

$$\beta \leq \text{value}(s, p) \rightarrow \beta \leq \text{alphaBeta}(s, p, \alpha, \beta).$$

Therefore, `alphaBeta(State, player)` is only an approximation of `value(State, player)`. However, it turns out that this approximation is all that is needed. Figure 4.3 on page 71 shows an implementation of `alphaBeta` that achieves this specification. Once the function `alphaBeta` is implemented, the function `value` can then be realized as follows:

$$\text{value}(s, p) := \text{alphaBeta}(s, p, -1, +1).$$

The reason is that we already know that $-1 \leq \text{value}(s, p) \leq +1$ and hence the first case of the specification of `alphaBeta` guarantees that the equation

$$\text{value}(s, p) = \text{alphaBeta}(s, p, -1, +1)$$

```

1  alphaBeta := cachedProcedure(State, player, alpha := -1, beta := 1) {
2      if (finished(State)) {
3          return utility(State, player);
4      }
5      other := arb(players() - { player });
6      val   := alpha;
7      for (ns in nextStates(State, player)) {
8          val := max({ val, -alphaBeta(ns, other, -beta, -val) });
9          if (val >= beta) {
10             return val;
11         }
12     }
13     return val;
14 };

```

Figure 4.3: α - β -Pruning.

holds. Since `alphaBeta` is implemented as a recursive procedure, the fact that the implementation of `alphaBeta` shown in Figure 4.3 on page 71 satisfies the specification given above can be shown by computational induction.

We proceed to discuss the implementation of the function `alphaBeta`.

1. If `State` is a terminal state, the function returns the `utility` of the given `State` with respect to `player`.
2. Since there are only two players, the `other` player is the player that remains if `player` is removed from the set of all players.
3. The variable `val` is supposed to store the maximum of the values of all states that can be reached from the given `State` if `player` makes one move.

According to the specification of `alphaBeta`, we are not interested in values that are below α . Hence, it suffices to initialize `val` with α . This way, in the case that we have

$$\text{value}(\text{State}, \text{player}) < \alpha,$$

instead of returning the true value of the given `State`, the function `alphaBeta(State, player, α , β)` will instead return the value α , which is permitted by its specification.

4. Next, we iterate over all successor states `ns` \in `nextStates(State, player)`.
5. We have to recursively evaluate the states `ns` for the `other` player. Since the value of a state for the `other` player is the negative of the value for `player`, we have to exchange the roles of α and β and prefix them with a negative sign.

Note that the value of `val` computed in the previous iteration of the `for`-loop can serve as the value of the parameter α for the next recursive call, since once we have found a state that has value `val`, we know already that `value(s, p)` is at least `val`. Therefore, if we encounter states with a value less than `val` during our recursive evaluation of the states reachable from `State`, then those states can be ignored. This is achieved via the using `val` for the parameter α .

6. As the specification of `alphaBeta` ask us to compute the value of `State` only in those cases where it is less than or equal to β , once we find a successor state `s` that has a value that is at least as big a β we can stop the evaluation and return this value.

This shortcut can result in significant savings of computation time!

4.4 Depth Limited Search

In practice, most games are far too complex to be evaluated completely, i.e. the size of the set `States` is so big that even the fastest computer does not stand a chance to explore this set completely. For example, it is impossible to consider all possible games in chess. Instead, we have to limit the exploration in a way that is similar to the way professional players evaluate their game: Usually, a player considers all variations of the game for, say, the next three moves. After a given number of moves, the value of a position is estimated using an evaluation function. In order to implement this idea, we add a parameter `limit` to the procedure `value`. On every recursive invocation of the function `value`, the function `limit` is decreased. Once the limit reaches 0, instead of invoking the function `value` again recursively, instead we try to estimate the value of the given `State` using a heuristic function. This leads to the code shown in Figure 4.4 on page 72.

```

1  value := cachedProcedure(State, player, limit, alpha := -1, beta := 1) {
2      if (finished(State)) {
3          return utility(State, player);
4      }
5      if (limit == 0) { return heuristic(State, player); }
6      other := arb(players() - { player });
7      val   := alpha;
8      for (ns in nextStates(State, player)) {
9          val := max({ val, -value(ns, other, limit - 1, -beta, -val) });
10         if (val >= beta) {
11             return val;
12         }
13     }
14     return val;
15 };

```

Figure 4.4: Depth-limited α - β -pruning.

For a game like tic-tac-toe it is difficult to come up with a decent heuristic. A very crude approach would be to define:

```
heuristic := [State, player] |-> 0;
```

This heuristic would simply estimate the value of all states to be 0. As this heuristic is only called after it has been tested that the game has not yet been decided, this approach is not utterly unreasonable. For a more complex game like chess, the heuristic could instead be a weighted count of all pieces. Concretely, the algorithm for estimating the value of a state would work as follows:

1. Initially, the variable `sum` is set to 0:

```
sum := 0;
```

2. We would count the number of white rooks `Rookwhite` and black rooks `Rookblack`, subtract these numbers from each other and multiply the difference by 5. The resulting number would be added to `sum`:

```
sum += (Rookwhite - Rookblack) · 5;
```

3. We would count the number of white bishops `Bishopwhite` and black bishops `Bishopblack`, subtract these numbers from each other and multiply the difference by 3. The resulting number would be added to `sum`:

```
sum += (Bishopwhite - Bishopblack) · 3;
```

4. In a similar way we would count knights, queens, and pawns. Approximately, the weights of knights are 3, a queen is worth 9 and a pawn is worth 1.

The resulting `sum` can then be used as an approximation of the value of a state. More details about the weights of the pieces can be found in the Wikipedia article “[chess piece relative value](#)”.

Exercise 9: Read up on the game [Connect Four](#). You can play it online at

<http://www.4-gewinnt.de>

Your task is to implement this game. At the address

<https://github.com/karlstroetmann/Artificial-Intelligence/blob/master/SetlX/connect-four-frame.stlx>

is a frame that can be used to solve this exercise. In this frame, only the following procedures need to be implemented.

1. The procedure `nextStates(s, p)` is called with a state s and a player p . It is supposed to return the set of all states that can be reached from the given state s if player p is next to make a move. In order to implement this function efficiently, you should make use of the function `find_empty` that is described below.
2. The procedure `find_empty(s, c)` is called with a state s and a column c . It searches the column c of the board specified by s bottom up for the first empty square. Once an empty square is found, the corresponding row is returned. If all the squares in the column c are filled, then the number 7 is returned instead.
3. The procedure `utility(s, p)` is called with a state s and a player p . It computes the utility that this state has for player p if player p is the next to move. For efficiency reasons, the implementation of `utility` should make use of the auxiliary functions `all_lines` and `last_line_filled` that are described below.
4. The procedure `all_lines()` takes no arguments. It is supposed to compute all sets of coordinates that form a line of four consecutive squares. For example, one such set would be the set

$$\{\langle 1, 1 \rangle, \langle 1, 2 \rangle, \langle 1, 3 \rangle, \langle 1, 4 \rangle\}.$$

This set would be the vertical line that starts at $\langle 1, 1 \rangle$. The horizontal line that starts at the location $\langle 1, 1 \rangle$ is the set

$$\{\langle 1, 1 \rangle, \langle 2, 1 \rangle, \langle 3, 1 \rangle, \langle 4, 1 \rangle\},$$

while the rising diagonal starting at $\langle 1, 1 \rangle$ is represented by the set

$$\{\langle 1, 1 \rangle, \langle 2, 2 \rangle, \langle 3, 3 \rangle, \langle 4, 4 \rangle\}.$$

Note that the procedure `all_lines` should be implemented as a `cachedProcedure`. This way it is guaranteed that the necessary computation is only performed once.

5. The procedure `last_line_filled(s)` takes a state s . It checks whether the game is drawn. The game is drawn iff the board is completely filled. This can be checked most efficiently by checking the final row of the board. If this row is completely filled, all other rows must have been filled, too. \diamond

Chapter 5

Linear Regression

A great deal of the current success of artificial intelligence is due to recent advances in [machine learning](#). In order to get a first taste of what machine learning is about, we introduce [linear regression](#) in this chapter, since linear regression is one of the most basic algorithms in machine learning. It is also the foundation for more advanced forms of machine learning like [logistic regression](#) and [neural networks](#). Furthermore, linear regression is surprisingly powerful. Finally, many of the fundamental problems of machine learning can already be illustrated with linear regression. Therefore it is only natural that we begin our study of machine learning with the study of linear regression.

5.1 Simple Linear Regression

Assume we want to know how the [engine displacement](#) of a car engine relates to the fuel consumption of the car. Of course, we could try to create a theoretical model that relates the engine displacement to its fuel consumption. However, due to our lack of understanding of engine physics, this is not an option for us. Instead, we could try to take a look at different engines and compare their engine displacement with the corresponding fuel consumption. This way, we would collect a set of m observations of the form $\langle x_1, y_1 \rangle, \dots, \langle x_m, y_m \rangle$ where x_i is the engine displacement of the engine in the i -th car, while y_i is the fuel consumption of the i -th car. We call x the [independent variable](#), while y is the [dependent variable](#). We define the vectors \mathbf{x} and \mathbf{y} as follows:

$$\mathbf{x} := \langle x_1, \dots, x_m \rangle^\top \quad \text{and} \quad \mathbf{y} := \langle y_1, \dots, y_m \rangle^\top.$$

Here, the operator $^\top$ is interpreted as the [transpose](#) operator, i.e. \mathbf{x} and \mathbf{y} are considered to be column vectors.

In linear regression, we use a [linear hypothesis](#) and assume that the dependent variable y_i is related to the independent variable x_i via a linear equation of the form

$$y_i := \vartheta_1 \cdot x_i + \vartheta_0.$$

We do not expect this equation to hold exactly. The reason is that there are many other factors besides the engine displacement that influence the fuel consumption. For example, both the weight of a car and its [aerodynamics](#) certainly influence the fuel consumption. We want to calculate those values ϑ_0 and ϑ_1 such that the [mean squared error](#), which is defined as

$$\text{MSE}(\vartheta_0, \vartheta_1) := \frac{1}{m-1} \cdot \sum_{i=1}^m (\vartheta_1 \cdot x_i + \vartheta_0 - y_i)^2, \quad (5.1)$$

is minimized. It can be shown that the solution to this minimization problem is given as follows:

$$\vartheta_1 = r_{x,y} \cdot \frac{s_y}{s_x} \quad \text{and} \quad \vartheta_0 = \bar{\mathbf{y}} - \vartheta_1 \cdot \bar{\mathbf{x}}. \quad (5.2)$$

This solution makes use of the values $r_{x,y}$, s_x , and s_y . In order to define these values, we first define the [sample mean values](#) $\bar{\mathbf{x}}$ and $\bar{\mathbf{y}}$ of \mathbf{x} and \mathbf{y} respectively, i.e. we have

$$\bar{\mathbf{x}} = \frac{1}{m} \cdot \sum_{i=1}^m x_i \quad \text{and} \quad \bar{\mathbf{y}} = \frac{1}{m} \cdot \sum_{i=1}^m y_i.$$

Furthermore, s_x and s_y are the *sample standard deviations* of \mathbf{x} and \mathbf{y} , i.e. we have

$$s_x = \sqrt{\frac{1}{m-1} \cdot \sum_{i=1}^m (x_i - \bar{\mathbf{x}})^2} \quad \text{and} \quad s_y = \sqrt{\frac{1}{m-1} \cdot \sum_{i=1}^m (y_i - \bar{\mathbf{y}})^2}.$$

Finally, $r_{x,y}$ is the *sample correlation coefficient* that is defined as

$$r_{x,y} = \frac{1}{(m-1) \cdot s_x \cdot s_y} \cdot \sum_{i=1}^m (x_i - \bar{\mathbf{x}}) \cdot (y_i - \bar{\mathbf{y}}).$$

The number $r_{x,y}$ is also known as the *Pearson correlation coefficient* or *Pearson's r*. It is named after *Karl Pearson* (1857 – 1936). Note that the formula for the parameter ϑ_1 can be simplified to

$$\vartheta_1 = \frac{\sum_{i=1}^m (x_i - \bar{\mathbf{x}}) \cdot (y_i - \bar{\mathbf{y}})}{\sum_{i=1}^m (x_i - \bar{\mathbf{x}})^2} \quad (5.3)$$

This latter formula should be used to calculate ϑ_1 . However, the previous formula is also useful because it shows the the correlation coefficient is identical to the coefficient ϑ_1 , provided the variables \mathbf{x} and \mathbf{y} have been normalized so that their standard deviation is 1.

Exercise 10: Prove Equation 5.2 and Equation 5.3.

Hint: Take the partial derivatives of $\text{MSE}(\vartheta_0, \vartheta_1)$ with respect to ϑ_0 and ϑ_1 . In order to minimize the expression $\text{MSE}(\vartheta_0, \vartheta_1)$ with respect to ϑ_0 and ϑ_1 , these derivatives have to be set to 0. \diamond

5.1.1 Assessing the Quality of Linear Regression

Assume that we have been given a set of m observations of the form $\langle x_1, y_1 \rangle, \dots, \langle x_m, y_m \rangle$ and that we have calculated the parameters ϑ_0 and ϑ_1 according to Equation 5.2 and Equation 5.3. Provided that not all x_i have the same value, these formulæ will return two numbers for ϑ_0 and ϑ_1 that define a linear model for \mathbf{y} in terms of \mathbf{x} . However, at the moment we still lack a number that tell us how good this linear model really is. In order to judge the quality of the linear model given by

$$y = \vartheta_0 + \vartheta_1 \cdot x$$

we can compute the mean squared error according to Equation 5.1. However, the mean squared error is an absolute number that, by itself, is difficult to interpret. The reason is that the variable \mathbf{y} might be inherently noisy and we have to relate this noise to the mean squared error. Now the noise contained in \mathbf{y} can be measured by the *sample variance* of \mathbf{y} and is given by the formula

$$\text{Var}(\mathbf{y}) := \frac{1}{m-1} \cdot \sum_{i=1}^m (y_i - \bar{\mathbf{y}})^2. \quad (5.4)$$

If we compare this formula to the formula for the mean squared error

$$\text{MSE}(\vartheta_0, \vartheta_1) := \frac{1}{m-1} \cdot \sum_{i=1}^m (\vartheta_1 \cdot x_i + \vartheta_0 - y_i)^2,$$

we see that the sample variance of \mathbf{y} is an upper bound for the mean squared error since we have

$$\text{Var}(\mathbf{y}) = \text{MSE}(\bar{\mathbf{y}}, 0),$$

i.e. the sample variance is the value that we would get for the mean squared error if we set ϑ_0 to the average value of \mathbf{y} and ϑ_1 to zero. Since ϑ_0 and ϑ_1 are chosen to minimize the mean squared error, we have

$$\text{MSE}(\vartheta_0, \vartheta_1) \leq \text{MSE}(\bar{\mathbf{y}}, 0) = \text{Var}(\mathbf{y}).$$

The mean squared error is an absolute value and, therefore, difficult to interpret. The fraction

$$\frac{\text{MSE}(\vartheta_0, \vartheta_1)}{\text{Var}(y)}$$

is called the proportion of the *unexplained variance* because it is the variance that is still left if we use our linear model to predict the values of \mathbf{y} given the values of \mathbf{x} . The *explained variance* which is also known as the R^2 *statistic* is defined as

$$R^2 := \frac{\text{Var}(\mathbf{y}) - \text{MSE}}{\text{Var}(\mathbf{y})} = 1 - \frac{\text{MSE}}{\text{Var}(\mathbf{y})}. \quad (5.5)$$

Here MSE is short for $\text{MSE}(\vartheta_0, \vartheta_1)$ where we have substituted the values from equation 5.2 and 5.3. The R^2 statistic measures the quality of our model: If it is small, then our model does not explain the variation of the value of \mathbf{y} when the value of \mathbf{x} changes. On the other hand, if it is near to 100%, then our model does a good job in explaining the variation of \mathbf{y} when \mathbf{x} changes.

Since the formulæ for $\text{Var}(\mathbf{y})$ and $\text{MSE}(\vartheta_0, \vartheta_1)$ have the same denominator $m - 1$, this denominator can be cancelled when the R^2 statistic is computed. To this end we define the *total sum of squares* TSS as

$$\text{TSS} := \sum_{i=1}^m (y_i - \bar{\mathbf{y}})^2 = (m - 1) \cdot \text{Var}(\mathbf{y})$$

and the *residual sum of squares* RSS as

$$\text{RSS} := \sum_{i=1}^m (\vartheta_1 \cdot x_i + \vartheta_0 - y_i)^2 = (m - 1) \cdot \text{MSE}(\vartheta_0, \vartheta_1).$$

Then the formula for the R^2 statistic can be written as

$$R^2 = 1 - \frac{\text{RSS}}{\text{TSS}}.$$

This is the formula that we will use when we implement simple linear regression.

It should be noted that R^2 is the square of Pearson's r . The notation is a bit inconsistent since Pearson's r is written in lower case, while R^2 is written in upper case. Finally R^2 is also known as the *coefficient of determination*. It tells us to what extent the value of the variable y is determined by the value of x .

5.1.2 Putting the Theory to the Test

In order to get a better feeling for linear regression, we want to test it to investigate the factors that determine the fuel consumption of cars. Figure 5.1 on page 77 shows the head of the data file “cars.csv” which I have adapted from the file

<http://www-bcf.usc.edu/~gareth/ISL/Auto.csv>.

Figure 5.1 on page 77 shows the column headers and the first ten data entries contained in this file. Altogether, this file contains data of 392 different car models.

The file “cars.csv” is part of the data set accompanying the excellent book *Introduction to Statistical Learning* by Gareth James, Daniela Witten, Trevor Hastie, and Robert Tibshirani [JWHT14]. The file “cars.csv” contains the fuel consumption of a number of different cars that were in widespread use during the seventies and early eighties of the last century. The first column of this data set lists the *miles per gallon*, i.e. the number of miles a car can go with one gallon of gas. Note that this number is inverse to the fuel consumption: If a car A can go twice as many miles per gallon than another car B, then the fuel consumption of A is half of the fuel consumption of B. Furthermore, besides the miles per gallon, for every car the following other parameters are listed:

1. `cyl` is the number of cylinders,
2. `displacement` is the engine displacement in cubic inches,

	mpg	cyl	displacement	hp	weight	acc	year	name
1	18.0	8	307.0	130.0	3504.0	12.0	70	chevrolet chevelle malibu
2	15.0	8	350.0	165.0	3693.0	11.5	70	buick skylark 320
3	18.0	8	318.0	150.0	3436.0	11.0	70	plymouth satellite
4	16.0	8	304.0	150.0	3433.0	12.0	70	amc rebel sst
5	17.0	8	302.0	140.0	3449.0	10.5	70	ford torino
6	15.0	8	429.0	198.0	4341.0	10.0	70	ford galaxie 500
7	14.0	8	454.0	220.0	4354.0	9.0	70	chevrolet impala
8	14.0	8	440.0	215.0	4312.0	8.5	70	plymouth fury iii
9	14.0	8	455.0	225.0	4425.0	10.0	70	pontiac catalina
10	15.0	8	390.0	190.0	3850.0	8.5	70	amc ambassador dpl
11								

Figure 5.1: The head of the file `cars.csv`.

3. `hp` is the engine power given in units of **horsepower**,
4. `weight` is the weight in pounds,
5. `acc` is the acceleration given as the time in seconds needed to accelerate from 0 miles per hour to 60 miles per hour,
6. `year` is the year in which the model was introduced, and
7. `name` is the name of the model.

Our aim is to determine which of these parameters can be used best to explain the fuel consumption of a car. To this end, I have written the program shown in Figure 5.2 on page 78.

The procedure `simple_linear_regression` take three arguments:

- (a) `fileName` is the name of the file containing the data.

It is assumed that this file is a **CSV**-file containing the data that need to be analyzed.

- (b) `types` is a list of type names of the columns present in the **CSV** file. This is needed by SETLX in order to correctly interpret the data.
- (c) `name` is the column name that is to be used for linear regression, i.e. it specifies the column that contains the values for the independent variable x . The dependent variable y is assumed to be given in the first column of the **CSV**-file.

The implementation of the procedure `simple_linear_regression` works as follows:

1. The function `readTable` is used to read the file specified in `fileName`. It returns the object `csv` which is of class `Table`. This class and the implementation of the function `readTable` is shown in Figure 5.3 on page 80. The object `csv` contains both the column names as well as the data that is present in the given file. The list `types` is needed in order to convert the strings that are read into their proper data types.
2. The function `getData` extracts the `data` from the object `csv`. Technically, `data` is a list of lists. The list `data` has the same length as there are lines in the file specified by `fileName`. Every single line in this file is converted into a list of entries of the appropriate types. These lists are the elements of the list `data`.
3. `number` is the number of data lines in the file specified by `fileName`, i.e. `number` is what we have called m in the formulæ given previously.
4. `index` is the index of the column that has the given `name`. For example, in the case of the file shown in Figure 5.1 on page 77, the `index` of the `name` “displacement” is 3 as the engine displacement is given in the third column of the file “cars.csv”.

```

1  simple_linear_regression := procedure(fileName, types, name) {
2      csv      := readTable(fileName, types);
3      data     := csv.getData();
4      number  := #data;
5      index   := find_index(name, csv.getColumnNames());
6      Y       := [];
7      X       := [];
8      L       := [1 .. number];
9      for (i in L) {
10         Y[i] := 1 / data[i][1];
11         X[i] := data[i][index];
12     }
13     xMean   := +/ X / number;
14     yMean   := +/ Y / number;
15     theta1  := (+/ [(X[i]-xMean)*(Y[i]-yMean) : i in L])
16               / (+/ [(X[i]-xMean)**2 : i in L]);
17     theta0  := yMean - theta1 * xMean;
18     TSS     := +/ [(Y[i] - yMean) ** 2 : i in L];
19     RSS     := +/ [(Y[i] - theta0 - theta1 * X[i]) ** 2 : i in L];
20     R2      := 1 - RSS / TSS;
21     canvas  := plot_createCanvas("Fuel consumption vs. $name$");
22     plot_addBullets(canvas, [ [X[i], Y[i]] : i in L ], [0,0,255], 2.0);
23     return R2;
24 };
25 find_index := procedure(x, List) {
26     return arb({ idx : idx in [1 .. #List] | List[idx] == x });
27 };

```

Figure 5.2: Simple Linear Regression

5. The values of the dependent variable y are stored in the list Y and the values of the independent variable x are stored in the list X .
6. The list L is used as an abbreviation so that iterations over all lists stored in **data** are simplified.
7. The **for**-loop fills the lists X and Y . Since the file “**cars.csv**” contains the miles per gallon in the first column labelled **mpg**, we need to convert this data into fuel consumption. As the fuel consumption is the reciprocal of the miles per gallon, the list Y is filled with the reciprocal of the first column of the file “**cars.csv**”.
8. **xMean** is the mean value \bar{x} of the independent variable \mathbf{x} .
9. **yMean** is the mean value \bar{y} of the dependent variable \mathbf{y} .
10. The coefficient **theta1** is computed according to Equation 5.3, which is repeated here for convenience:

$$\vartheta_1 = \frac{\sum_{i=1}^m (x_i - \bar{x}) \cdot (y_i - \bar{y})}{\sum_{i=1}^m (x_i - \bar{x})^2}.$$

11. The coefficient **theta0** is computed according to Equation 5.2, which reads

$$\vartheta_0 = \bar{y} - \vartheta_1 \cdot \bar{x}.$$

dependent variable	explained variance
displacement	0.75
cyl	0.70
hp	0.73
weight	0.78
acc	0.21
year	0.31

Table 5.1: Explained variance for various dependent variables.

12. TSS is the **total sum of squares** and is computed using the formula

$$\text{TSS} = \sum_{i=1}^m (y_i - \bar{y})^2.$$

13. RSS is the **residual sum of squares** and is computed as

$$\text{RSS} := \sum_{i=1}^m (\vartheta_1 \cdot x_i + \vartheta_0 - y_i)^2.$$

14. R^2 is the R^2 statistic and measures the **proportion of the explained variance**. It is computed using the formula

$$R^2 = \frac{\text{TSS} - \text{RSS}}{\text{TSS}}.$$

15. The function `plot_createCanvas` creates a **canvas** that is used for plotting.

16. The function `plot_addBullets` plots data point onto the **canvas**. The second argument to this function is a list of pairs of the form

$$[\langle x_1, y_1 \rangle, \dots, \langle x_m, y_m \rangle]$$

The pair $\langle x_i, y_i \rangle$ is then plotted as a blue circle with radius 2.0 at the position $\langle x_i, y_i \rangle$.

Figure 5.3 on page 80 shows the implementation of the class `table` and the function `readTable`. This figure is only given for completeness.

Figure 5.4 on page 81 shows how to call the procedure `simple_linear_regression`. We need to define the types of the various columns that are present in the file “`cars.csv`”. Next, for all those column names present in this file, we use this variable to compute the explained variance. The resulting values are shown in Table 5.1. It seems that, given the data in the file “`cars.csv`”, the best indicator for the fuel consumption is the **weight** of a car. The **displacement**, the power **hp** of an engine, and the number of cylinders **cyl** are also good predictors. But notice that the **weight** is the real cause of fuel consumption: If a car has a big weight, it will also need a more powerful engine. Hence the variable **hp** is correlated with the variable **weight** and will therefore also provide a reasonable explanation of the fuel consumption, although the high engine power is not the most important cause of the fuel consumption.

```

1  class Table(columnNames, types, data) {
2      mColumnNames := columnNames;
3      mTypes       := types;
4      mData        := data;
5      static {
6          getColumnNames := [ ] |-> mColumnNames;
7          getTypes       := [ ] |-> mTypes;
8          getData        := [ ] |-> mData;
9          getRow         := [r] |-> mData[r];
10         getLength      := [ ] |-> #mData;
11
12         head := procedure(limit := 10) {
13             print(mColumnNames);
14             print(mTypes);
15             for (i in [1 .. limit]) {
16                 print(mData[i]);
17             }
18         };
19     }
20 }
21 readTable := procedure(fileName, types) {
22     allData := readFile(fileName); // list of lines
23     columnNames := split(allData[1], '\s*');
24     data := [];
25     for (i in [2 .. #allData]) {
26         row := split(allData[i], '\s*');
27         data[i-1] := [evalType(type, s) : [type, s] in types >< row];
28     }
29     return Table(columnNames, types, data);
30 };
31 evalType := procedure(type, s) {
32     switch {
33         case type == "double": return double(s);
34         case type == "int" : return int(s);
35         case type == "string": return s;
36         default: abort("unknown type $type$ in evalType");
37     }
38 };

```

Figure 5.3: The class `Table`.

5.1.3 Testing the Statistical Significance

In this section we answer the question of how to assess the **statistical significance** of our results. In the case of simple linear regression, the **null hypothesis** is given as follows:

H_0 : There is no relationship between the variables x and y .

In order to compute the probability that the null hypothesis is true, we have to compute the ***t*-statistic**, which is given by the following formula:

$$t := |\vartheta_1| \cdot \sqrt{\frac{(m-2)}{\text{RSS}} \cdot \sum_{i=1}^m (x_i - \bar{x})^2}. \quad (5.6)$$

```

1  test := procedure() {
2      types := ["double", "int", "double", "double", "double", "double", "int", "string"]
3      for (varName in ["displacement", "cyl", "hp", "weight", "acc", "year"]) {
4          R2 := simple_linear_regression("cars.csv", types, varName);
5          print("The explained variance for $varName$ vs fuel consumption is $R2$.");
6      }
7  };

```

Figure 5.4: Calling the procedure `simple_linear_regression`.

The t statistics is distributed according to **Student's t -distribution** with $m - 2$ degrees of freedom. In the upcoming version of SETLX, a preview of which is already available at

<https://github.com/jonassiefker/setlX>,

the cumulative Student's t -distribution can be computed via the function

`stat_studentCDF(t, ν)`,

where t is the value of the t -statistic and ν is the number of degrees of freedom. If we substitute $\nu := m - 2$, then the expression

`1 - stat_studentCDF(t, ν)`,

computes the probability that the null hypothesis is true. This probability is also known as the **p -value**. If the p -value is small, e.g. less than 0.01, then the null hypothesis is refuted and we can conclude that there is some relationship between x and y . In the example previously discussed, i.e. the example relating the fuel consumption of a car to various other parameters, the p -values are all very small, i.e. less than $2 \cdot 10^{-16}$. This is due to the fact that we have had enough data at our disposal: As a rule of thumb it can be said that once we have more than 30 pairs $\langle x_i, y_i \rangle$ and there is indeed a relationship between x and y , then simple linear regression will detect this relationship.

5.2 General Linear Regression

In practise, it is rarely the case that a given observed variable y only depends on a single variable x . To take the example of the fuel consumption of a car further, in general we would expect that the fuel consumption of a car depends not only on the mass of the car but is also related to the other parameters. To this end, we present the theory of **general linear regression**. In a **general regression problem** we are given a list of m pairs of the form $\langle \mathbf{x}^{(i)}, y^{(i)} \rangle$ where $\mathbf{x}^{(i)} \in \mathbb{R}^p$ and $y^{(i)} \in \mathbb{R}$ for all $i \in \{1, \dots, m\}$. The number p is called the number of **features**, while the pairs are called the **training examples**. Our goal is to compute a function

$$F : \mathbb{R}^p \rightarrow \mathbb{R}$$

such that $F(\mathbf{x}^{(i)})$ approximates $y^{(i)}$ as precisely as possible for all $i \in \{1, \dots, m\}$, i.e. we want to have

$$\forall i \in \{1, \dots, m\} : F(\mathbf{x}^{(i)}) \approx y^{(i)}.$$

In order to make the notation $F(\mathbf{x}^{(i)}) \approx y^{(i)}$ more precise, we define the **mean squared error**

$$\text{MSE} := \frac{1}{m-1} \cdot \sum_{i=1}^m \left(F(\mathbf{x}^{(i)}) - y^{(i)} \right)^2. \quad (5.7)$$

Then, given the list of training examples $[\langle \mathbf{x}^{(1)}, y^{(1)} \rangle, \dots, \langle \mathbf{x}^{(n)}, y^{(n)} \rangle]$, our goal is to minimize MSE. In order to proceed, we need to have a model for the function F . The simplest model is a linear model, i.e. we assume that F is given as

$$F(\mathbf{x}) = \sum_{j=1}^p w_j \cdot x_j + b = \mathbf{x}^\top \cdot \mathbf{w} + b \quad \text{where } \mathbf{w} \in \mathbb{R}^p \text{ and } b \in \mathbb{R}.$$

Here, the expression $\mathbf{x}^\top \cdot \mathbf{w}$ denotes the matrix product of the vector \mathbf{x}^\top , which is viewed as a 1-by- m matrix, and the vector \mathbf{w} . Alternatively, this expression could be interpreted as the dot product of the vector \mathbf{x} and the vector \mathbf{w} . At this point you might wonder why it is useful to introduce matrix notation here. The reason is that this notation shortens the formula and, furthermore, is more efficient to implement since most programming languages used in machine learning have special library support for matrix operations. Provided the computer is equipped with a graphics card, some programming languages are even able to delegate matrix operations to the graphics unit. This results in a considerable speed-up.

The definition of F given above is the model used in [linear regression](#). Here, \mathbf{w} is called the [weight vector](#) and b is called the [bias](#). It turns out that the notation can be simplified if we extend the p -dimensional feature vector \mathbf{x} to an $p+1$ -dimensional vector \mathbf{x}' such that

$$x'_j := x_j \quad \text{for all } j \in \{1, \dots, p\} \quad \text{and} \quad x'_{m+1} := 1.$$

To put it in words, the vector \mathbf{x}' results from the vector \mathbf{x} by appending the number 1:

$$\mathbf{x}' = \langle x_1, \dots, x_p, 1 \rangle^\top \quad \text{where } \langle x_1, \dots, x_p \rangle = \mathbf{x}^\top.$$

Furthermore, we define

$$\mathbf{w}' := \langle w_1, \dots, w_p, b \rangle^\top \quad \text{where } \langle w_1, \dots, w_p \rangle = \mathbf{w}^\top.$$

Then we have

$$F(\mathbf{x}) = \mathbf{w} \cdot \mathbf{x} + b = \mathbf{w}' \cdot \mathbf{x}'.$$

Hence, the bias has been incorporated into the weight vector at the cost of appending the number 1 at the end of input vector. As we want to use this simplification, from now on we assume that the input vectors $\mathbf{x}^{(i)}$ have all been extended so that their last component is 1. Using this assumption, we define the function F as

$$F(\mathbf{x}) := \mathbf{x}^\top \cdot \mathbf{w}.$$

Now equation (5.7) can be rewritten as follows:

$$\text{MSE}(\mathbf{w}) = \frac{1}{m-1} \cdot \sum_{i=1}^m \left((\mathbf{x}^{(i)})^\top \cdot \mathbf{w} - y^{(i)} \right)^2. \quad (5.8)$$

Our aim is to rewrite the sum appearing in this equation as a scalar product of a vector with itself. To this end, we first define the vector \mathbf{y} as follows:

$$\mathbf{y} := \langle y^{(1)}, \dots, y^{(m)} \rangle^\top.$$

Note that $\mathbf{y} \in \mathbb{R}^m$ since it has a component for all of the m training examples. Next, we define the [design matrix](#) X as follows:

$$X := \begin{pmatrix} (\mathbf{x}^{(1)})^\top \\ \vdots \\ (\mathbf{x}^{(m)})^\top \end{pmatrix}$$

Defined this way, the row vectors of the matrix X are the vectors $\mathbf{x}^{(i)}$ transposed. Now we have the following:

$$X \cdot \mathbf{w} - \mathbf{y} = \begin{pmatrix} (\mathbf{x}^{(1)})^\top \\ \vdots \\ (\mathbf{x}^{(m)})^\top \end{pmatrix} \cdot \mathbf{w} - \mathbf{y} = \begin{pmatrix} (\mathbf{x}^{(1)})^\top \cdot \mathbf{w} - y_1 \\ \vdots \\ (\mathbf{x}^{(m)})^\top \cdot \mathbf{w} - y_m \end{pmatrix}$$

Taking the square of the vector $X \cdot \mathbf{w} - \mathbf{y}$ we discover that we can rewrite equation (5.8) as follows:

$$\text{MSE}(\mathbf{w}) = \frac{1}{m-1} \cdot (X \cdot \mathbf{w} - \mathbf{y})^\top \cdot (X \cdot \mathbf{w} - \mathbf{y}). \quad (5.9)$$

5.2.1 Some Useful Gradients

In the last section, we have computed the mean squared error $\text{MSE}(\mathbf{w})$ using equation (5.9). Our goal is to minimize the $\text{MSE}(\mathbf{w})$ by choosing the weight vector \mathbf{w} appropriately. A necessary condition for $\text{MSE}(\mathbf{w})$ to be minimal is

$$\nabla \text{MSE}(\mathbf{w}) = \mathbf{0},$$

i.e. the gradient of $\text{MSE}(\mathbf{w})$ needs to be zero. In order to prepare for the computation of $\nabla \text{MSE}(\mathbf{w})$, we first compute the gradient of two simpler functions.

Computing the Gradient of $f(\mathbf{x}) = \mathbf{x}^\top \cdot C \cdot \mathbf{x}$

Suppose the function $f : \mathbb{R}^n \rightarrow \mathbb{R}$ is defined as

$$f(\mathbf{x}) := \mathbf{x}^\top \cdot C \cdot \mathbf{x} \quad \text{where } C \in \mathbb{R}^{n \times n}.$$

If we write the matrix C as $C = (c_{i,j})_{\substack{i=1,\dots,n \\ j=1,\dots,n}}$ and the vector \mathbf{x} as $\mathbf{x} = \langle x_1, \dots, x_n \rangle^\top$, then $f(\mathbf{x})$ can be computed as follows:

$$f(\mathbf{x}) = \sum_{i=1}^n x_i \cdot \sum_{j=1}^n c_{i,j} \cdot x_j = \sum_{i=1}^n \sum_{j=1}^n x_i \cdot c_{i,j} \cdot x_j.$$

We compute the partial derivative of f with respect to x_k and use the product rule together with the definition of the **Kronecker delta** $\delta_{i,j}$, which is defined as 1 if $i = j$ and as 0 otherwise:

$$\begin{aligned} \frac{\partial f}{\partial x_k} &= \sum_{i=1}^n \sum_{j=1}^n \left(\frac{\partial x_i}{\partial x_k} \cdot c_{i,j} \cdot x_j + x_i \cdot c_{i,j} \cdot \frac{\partial x_j}{\partial x_k} \right) \\ &= \sum_{i=1}^n \sum_{j=1}^n \left(\delta_{i,k} \cdot c_{i,j} \cdot x_j + x_i \cdot c_{i,j} \cdot \delta_{j,k} \right) \\ &= \sum_{j=1}^n c_{k,j} \cdot x_j + \sum_{i=1}^n x_i \cdot c_{i,k} \\ &= (C \cdot \mathbf{x})_k + (C^\top \cdot \mathbf{x})_k \end{aligned}$$

Hence we have shown that

$$\nabla f(\mathbf{x}) = (C + C^\top) \cdot \mathbf{x}.$$

If the matrix C is symmetric, i.e. if $C = C^\top$, this simplifies to

$$\nabla f(\mathbf{x}) = 2 \cdot C \cdot \mathbf{x}.$$

Next, if the function $g : \mathbb{R}^n \rightarrow \mathbb{R}$ is defined as

$$g(\mathbf{x}) := \mathbf{b}^\top \cdot A \cdot \mathbf{x}, \quad \text{where } \mathbf{b} \in \mathbb{R}^n \text{ and } A \in \mathbb{R}^{n \times n},$$

then a similar calculation shows that

$$\nabla g(\mathbf{x}) = A^\top \cdot \mathbf{b}.$$

Exercise 11: Prove this equation.

5.2.2 Deriving the Normal Equation

Next, we derive the so called **normal equation** for linear regression. To this end, we first expand the product in equation (5.9):

$$\begin{aligned}
\text{MSE}(\mathbf{w}) &= \frac{1}{m-1} \cdot (X \cdot \mathbf{w} - \mathbf{y})^\top \cdot (X \cdot \mathbf{w} - \mathbf{y}) \\
&= \frac{1}{m-1} \cdot (\mathbf{w}^\top \cdot X^\top - \mathbf{y}^\top) \cdot (X \cdot \mathbf{w} - \mathbf{y}) && \text{since } (A \cdot B)^\top = B^\top \cdot A^\top \\
&= \frac{1}{m-1} \cdot (\mathbf{w}^\top \cdot X^\top \cdot X \cdot \mathbf{w} - \mathbf{y}^\top \cdot X \cdot \mathbf{w} - \mathbf{w}^\top \cdot X^\top \cdot \mathbf{y} + \mathbf{y}^\top \cdot \mathbf{y}) \\
&= \frac{1}{m-1} \cdot (\mathbf{w}^\top \cdot X^\top \cdot X \cdot \mathbf{w} - 2 \cdot \mathbf{y}^\top \cdot X \cdot \mathbf{w} + \mathbf{y}^\top \cdot \mathbf{y}) && \text{since } \mathbf{w}^\top \cdot X^\top \cdot \mathbf{y} = \mathbf{y}^\top \cdot X \cdot \mathbf{w}
\end{aligned}$$

The fact that

$$\mathbf{w}^\top \cdot X^\top \cdot \mathbf{y} = \mathbf{y}^\top \cdot X \cdot \mathbf{w}$$

might not be immediately obvious. It follows from two facts:

1. For two matrices A and B such that the matrix product $A \cdot B$ is defined we have

$$(A \cdot B)^\top = B^\top \cdot A^\top.$$

2. The matrix product $\mathbf{w}^\top \cdot X^\top \cdot \mathbf{y}$ is a real number. The transpose r^\top of a real number r is the number itself, i.e. $r^\top = r$ for all $r \in \mathbb{R}$. Therefore, we have

$$\mathbf{w}^\top \cdot X^\top \cdot \mathbf{y} = (\mathbf{w}^\top \cdot X^\top \cdot \mathbf{y})^\top = \mathbf{y}^\top \cdot X \cdot \mathbf{w}.$$

Hence we have shown that

$$\text{MSE}(\mathbf{w}) = \frac{1}{m-1} \cdot (\mathbf{w}^\top \cdot (X^\top \cdot X) \cdot \mathbf{w} - 2 \cdot \mathbf{y}^\top \cdot X \cdot \mathbf{w} + \mathbf{y}^\top \cdot \mathbf{y}) \quad (5.10)$$

holds. The matrix $X^\top \cdot X$ used in the first term is symmetric because

$$(X^\top \cdot X)^\top = X^\top \cdot (X^\top)^\top = X^\top \cdot X.$$

Using the results from the previous section we can now compute the gradient of $\text{MSE}(\mathbf{w})$ with respect to \mathbf{w} . The result is

$$\nabla \text{MSE}(\mathbf{w}) = \frac{2}{m-1} \cdot (X^\top \cdot X \cdot \mathbf{w} - X^\top \cdot \mathbf{y}).$$

If the squared error $\text{MSE}(\mathbf{w})$ has a minimum for the weights \mathbf{w} , then we must have

$$\nabla \text{MSE}(\mathbf{w}) = \mathbf{0}.$$

This leads to the equation

$$\frac{2}{m-1} \cdot (X^\top \cdot X \cdot \mathbf{w} - X^\top \cdot \mathbf{y}) = \mathbf{0}.$$

This equation can be rewritten as

$$(X^\top \cdot X) \cdot \mathbf{w} = X^\top \cdot \mathbf{y}. \quad (5.11)$$

This equation is called the [normal equation](#).

Remark: Although the matrix $X^\top \cdot X$ will often be invertible, for numerical reasons it is not advisable to rewrite the normal equation as

$$\mathbf{w} = (X^\top \cdot X)^{-1} \cdot X^\top \cdot \mathbf{y}.$$

Instead, when solving the normal equation we will use the SETLX function `la_solve(A,b)`, which takes a matrix $A \in \mathbb{R}^{n \times n}$ and a vector $\mathbf{b} \in \mathbb{R}^n$ and solves the equation

$$A \cdot \mathbf{x} = \mathbf{b}.$$

◇

5.2.3 Testing the Statistical Significance

The F-statistic is defined according to the following formula:

$$F = \frac{TSS - RSS}{RSS} \cdot \frac{m - p - 1}{p} \quad (5.12)$$

The F-statistic is distributed according to the [Fisher-Snedecor-distribution](#) with $p - 1$ degrees of freedom in the nominator and $m - p$ degrees of freedom in the denominator.

5.2.4 Implementation

Figure 5.5 on page 85 shows an implementation of general linear regression. The procedure `run_linear_regression(f, t)` takes a file name f and a list of types t as its input.

```

1  run_linear_regression := procedure(fileName, types) {
2      csv      := readTable(fileName, types);
3      data     := csv.getData();
4      number  := #data;
5      xList   := [];
6      yList   := [];
7      for (i in [1 .. number]) {
8          yList[i] := 1 / data[i][1];
9          xList[i] := la_vector(data[i][2..-2] + [1.0]);
10     }
11     X := la_matrix(xList);
12     y := la_vector(yList);
13     w := la_solve(X! * X, X! * y);
14     d := X * w - y;
15     RSS := d * d;
16     yMean := +/ yList / number;
17     TSS := +/ [(yList[i] - yMean) ** 2 : i in [1 .. number]];
18     R2 := 1 - RSS / TSS;
19     return R2;
20 };

```

Figure 5.5: General linear regression.

1. It reads the file and stores the resulting `Table` object in the variable `csv`.
2. The data from this object is stored as a list of lists in the variable `data`.
3. `number` is the number of training examples.
4. `xList` stores the data of the independent variable x . Therefore, `xList` is a list of vectors. Note that we have extended the data vectors by putting a 1.0 at the end.
5. `yList` stores the data of the depend variable y .
6. The data in `xList` is then stored in the matrix `X`.
7. The data in `yList` is stored in the vector `y`.
8. The normal equation is formulated and solved using the function `la_solve`. Note that the postfix-operator “!” computes the transpose of a matrix when applied to a matrix.

9. The variable `d` is the difference between the predictions of the linear model and the observed values y .
10. `RSS` is the residual sum of squares.
11. `yMean` is the mean value of the variable y .
12. `TSS` is the total sum of squares.
13. `R2` is the proportion of the explained variance.

When we run the program shown in Figure 5.5 on page 85 with the data stored in `cars.csv`, which had been discussed previously, then the proportion of explained variance is 81%. Considering that our data does not take account of the aerodynamics of the cars, this is a reasonable result.

Exercise 12: The file “`advertising.csv`”, which is available at

<https://github.com/karlstroetmann/Artificial-Intelligence/blob/master/Set1X/advertising.csv>,

contains data about advertising. Concretely, let us think that this data lists the amount of money that has been spent in advertising a luxury refrigerator that is able to detect the amount of food stored and that is able to reorder supplies with `AmazonFresh™`. Every line of this csv file has five entries:

- (a) The first column is interpreted as the number of a store. All together, there are 200 stores that have sold this refrigerator. It is assumed that all of these stores are located in different counties.
- (b) The second column gives the amount of money spent in TV advertising in thousands of dollar for the store in that county.
- (c) The third column gives the amount of money spent in radio advertising in thousands of dollar.
- (d) The fourth column gives the amount of money spent in newspaper advertising in thousands of dollar.
- (e) The fifth column gives the number of refrigerators in thousands that have been sold in the corresponding store.

With the exception of the first column, all columns contain floating point numbers. Assume you are the marketing manager that is responsible for the advertising campaign. Your task is to evaluate the effectiveness of the different type of advertisements. To this end, answer the following questions.

1. Should you spend more money on newspaper adds and less on TV adds or is it the other way around? Or are radio advertisements the way to go? Assume that your advertising budget is fixed and that you want to sell as many refrigerators as possible.

Use linear regression to solve this question.

2. At this point, you will have discovered that only two of the three different types of advertisements contribute to the number sales. Check, whether there is an interaction between these two types of advertisements. For example, check whether TV adds have a positive or negative effect on the effectiveness of radio adds.

You can use a product term as a new features in your linear regression model to answer this question. For example, you can include the product of money spend in TV adds and money spend in radio adds as a new feature.

3. In order to further increase the accuracy of our model, we can use a `quadratic model`. If we have only two independent features x and y , a quadratic model has the following form:

$$f(x, y) = w_1 \cdot x^2 + w_2 \cdot x \cdot y + w_3 \cdot y^2 + w_4 \cdot x + w_5 \cdot y + w_6.$$

Compute the fraction of the explained variance that is achieved if a quadratic model is used to predict the sales.

4. In the previous part of the exercise you will have discovered that for a quadratic model, the fraction of the explained variance is quite high. Hence, the quadratic model seems to be accurate. Now that you have found an accurate model, by how much can you increase the sales by optimizing the strategy? Assume that for every county the company wants to spend the same amount of money for advertisement that is has been spent previously for that county.

Note: I have stolen the file “`advertising.csv`” from the data set accompanying the book “[Introduction to Statistical Learning](#)” [JWHT14]. As far as I understand it, the data contained in this file is fictional data, i.e. it has been made up. \diamond

Partial Solution to part 4: The total budget for advertising in a given county is fixed. Assume this total budget is called b . If x is the amount of money spent for TV advertisements and y is the amount of money spent for radio advertisements, then we must have

$$x + y = b \quad \text{and therefore} \quad y = b - x.$$

Hence the number of units sold is given by the formula

$$g(x) := f(x, b - x) = w_1 \cdot x^2 + w_2 \cdot x \cdot (b - x) + w_3 \cdot (b - x)^2 + w_4 \cdot x + w_5 \cdot (b - x) + w_6.$$

In order to compute the value x that maximizes g we have to compute the derivative of $g(x)$:

$$\begin{aligned} \frac{dg}{dx}(x) &= 2 \cdot w_1 \cdot x + w_2 \cdot b - 2 \cdot w_2 \cdot x + 2 \cdot w_3 \cdot (x - b) + w_4 - w_5 \\ &= 2 \cdot (w_1 - w_2 + w_3) \cdot x + w_2 \cdot b - 2 \cdot w_3 \cdot b + w_4 - w_5 \end{aligned}$$

This expression is 0 if and only if

$$x = \frac{-w_2 \cdot b + 2 \cdot w_3 \cdot b - w_4 + w_5}{2 \cdot (w_1 - w_2 + w_3)}.$$

It turns out that when the advertising strategy is optimized as outline above, the quadratic model predicts a sale of 5 511 887 units, while the number of units sold previously was 2 804 500 units. Hence, the optimized strategy is expected to sell about 97 % units more than the strategy applied previously. \square

Chapter 6

Classification

One of the earliest application of artificial intelligence is **classification**. A good example of classification is **spam detection**. A system for spam detection classifies an email as either spam or not spam. To do so, it first computes various **features** of the email and then uses these features to determine whether the email is likely to be spam. For example, a possible feature would be the number of occurrences of the word “**pharmacy**” in the text of the email.

6.1 Introduction

Formally, the classification problem in machine learning can be stated as follows. We are given a set of objects $S := \{o_1, \dots, o_n\}$ and a set of classes $C := \{c_1, \dots, c_k\}$. Furthermore, there exists a function

$$\text{classify} : S \rightarrow C$$

that assigns a class $\text{classify}(o)$ to every object $o \in S$. The set S is called the **sample space**. In the example of spam detection, the sample space S is the set of all emails that we might receive, i.e. S is the set of all strings of the ASCII alphabet, while

$$C = \{\text{spam}, \text{ham}\}.$$

If $\text{classify}(o) = \text{spam}$, we consider the email o to be spam, while we would consider it a genuine message otherwise. Our goal is to compute the function **classify**. In order to do this, we use an approach known as **supervised learning**: We take a subset $S_{\text{Train}} \subseteq S$ of emails where we already know whether the emails are spam or not. This set S_{Train} is called the **training set**. Next, we define a set of D **features** for every $o \in S$. These features have to be computable, i.e. we must have a function

$$\text{feature} : S \times \{1, \dots, D\} \rightarrow \mathbb{R}$$

such that $\text{feature}(o, j)$ computes the j -th feature and we have to be able to implement this function with reasonable efficiency. In general, the values of the features are real values. However, there are cases where these values are just Booleans. If

$$\text{feature}(o, j) \in \mathbb{B} \text{ for all } o \in S,$$

then the j -th feature is a **binary feature**. If we encode **false** as 0 and **true** as 1, then the set of Boolean values \mathbb{B} can be considered a subset of \mathbb{R} and hence Boolean features can be considered as real numbers. For example, in the case of spam detection, the first feature could be the occurrence of the string “**pharmacy**”. In this case, we would have

$$\text{feature}(o, 1) := (\text{pharmacy} \in o),$$

i.e. the first feature would be to check whether the email o contains the string “**pharmacy**”. If we want to be more precise, we can instead define the first feature as

$$\text{feature}(o, 1) := \text{count}(\text{"pharmacy"}, o),$$

i.e. we would count the number of occurrences of the string “**pharmacy**” in our email o . As `count("pharmacy", o)` is always a natural number, in this case the first feature would be a **discrete** feature. However, we can be even more precise than just counting the number of occurrences of “**pharmacy**”. After all, there is a difference if the string “**pharmacy**” occurs once in an email containing but a hundred characters or whether it occurs once in an email with a length of several thousand characters. To this end, we would then define the first feature as

$$\text{feature}(o, 1) := \frac{\text{count}(\text{"pharmacy"}, o)}{\#o},$$

where $\#o$ defines the number of characters in the string o . In this case, the first feature would be a **continuous** feature and as this is the most general case, unless stated otherwise, we deal with the continuous case.

Having defined the features, we next need a **model** of the function **classify** that tries to approximate the function **classify** via the features. This model is given by a function

$$\text{model} : \mathbb{R}^D \rightarrow C$$

such that

$$\text{model}(\text{feature}(o, 1), \dots, \text{feature}(o, D)) \approx \text{classify}(o).$$

Using the function **model**, we can then approximate the function **classify** using a function **guess** that is defined as

$$\text{guess}(o) := \text{model}(\text{feature}(o, 1), \dots, \text{feature}(o, D))$$

Most of the time, the function **guess** will only approximate the function **classify**, i.e. we will have

$$\text{guess}(o) = \text{classify}(o)$$

for most objects of $o \in S$ but not for all of them. The **accuracy** of our model is then defined as the fraction of those objects that are classified correctly, i.e.

$$\text{accuracy} := \frac{\#\{o \in S \mid \text{guess}(o) = \text{classify}(o)\}}{\#S}.$$

If the set S is infinite, this equation has to be interpreted as a limit, i.e. we can define S_n as the set of strings that have a length of at most n . Then, the accuracy can be defined as follows:

$$\text{accuracy} := \lim_{n \rightarrow \infty} \frac{\#\{o \in S_n \mid \text{guess}(o) = \text{classify}(o)\}}{\#S_n}.$$

The function **model** is usually determined by a set of **parameters** or **weights** \mathbf{w} . In this case, we have

$$\text{model}(\mathbf{x}) = \text{model}(\mathbf{x}; \mathbf{w})$$

where \mathbf{x} is the vector of features, while \mathbf{w} is the vector of weights. Later, when we introduce **logistic regression**, we will assume that the number of weights is the same as the number of features. When it comes to the choice of the model, it is important to understand that, at least in practical applications, all models are wrong. Nevertheless, some models are useful. There are two reasons for this:

1. We do not fully understand the function **classify** that we want to approximate by the function **model**.
2. The function **classify** is so complex, that even if we could compute it exactly, the resulting model would be much too complicated.

The situation is similar in physics: Let us assume that we intend to model the fall of an object. A model that is a hundred percent accurate would have to include not only air friction, but also the possible effects of tidal forces or, in case we have a metallic object, the effects of the magnetic field of the earth have to be taken into account. On top of that we need some corrections from relativistic physics and even quantum physics. A model of this kind would be so complicated that it would be useless.

Let us summarize our introductory discussion of machine learning in general and classification in particular. A set S of objects and a set C of classes are given. Our goal is to approximate a function

$$\text{classify} : S \rightarrow C$$

using certain **features** of our objects. The function `classify` is then approximated using a function `model` as follows:

$$\text{model}(\text{feature}(o, 1), \dots, \text{feature}(o, D); \mathbf{w}) \approx \text{classify}(o).$$

The model depends on a vector of parameters \mathbf{w} . In order to **learn** these parameters, we are given a **training set** S_{Train} that is a subset of S . As we are dealing with **supervised learning**, the function `classify` is known for all objects $o \in S_{\text{Train}}$. Our goal is to determine the parameters \mathbf{w} such that the number of mistakes we make on the training set is minimized.

6.1.1 Notation

We conclude this introductory section by fixing some notation. Let us assume that the objects $o \in S_{\text{Train}}$ are numbered from 1 to N , while the features are numbered from 1 to D . Then we define

1. $\mathbf{x}_i := [\text{feature}(o_i, 1), \dots, \text{feature}(o_i, D)]$ for all $i \in \{1, \dots, N\}$.
i.e. \mathbf{x}_i is a D -dimensional vector that collects the features of the i -th training object.
2. $x_{i,j} := \text{feature}(o_i, j)$ for all $i \in \{1, \dots, N\}$ and $j \in \{1, \dots, D\}$.
i.e. $x_{i,j}$ is the j -th feature of the i -th object.
3. $y_i := \text{classify}(o_i)$ for all $i \in \{1, \dots, N\}$
i.e. y_i is the class of the i -th object.

Mathematically, our goal is now to maximize the accuracy of our model as a function of the parameters \mathbf{w} .

6.1.2 Applications of Classification

Besides spam detection, there are many other classification problems that can be solved using machine learning. To give just one more example, imagine a general practitioner that receives a patient and examines his symptoms. In this case, the symptoms can be seen as the features of the patient. For example, these features could be

1. body temperature,
2. blood pressure,
3. heart rate,
4. body weight,
5. breathing difficulties,
6. age,

to name but a few of the possible features. Based on these symptoms, the general practitioner would then decide on an illness, i.e. the set of classes for the classification problem would be

$$\{\text{commonCold}, \text{pneumonia}, \text{asthma}, \text{flu}, \dots, \text{unknown}\}.$$

Hence, the task of disease diagnosis is a classification problem. This was one of the earliest problem that was tackled by artificial intelligence. As of today, **computer-aided diagnosis** has been used for more than 40 years in many hospitals.

6.2 Digression: The Method of Gradient Ascent

In machine learning, it is often the case that we have to find either the maximum or the minimum of a function

$$f : \mathbb{R}^n \rightarrow \mathbb{R}.$$

For example, when we discuss [logistic regression](#) in the next section, we will have to find the maximum of a function. First, let us introduce the arg max function. The idea is that

$$\hat{\mathbf{x}} = \arg \max_{\mathbf{x} \in \mathbb{R}^n} f(\mathbf{x})$$

is that value of $\mathbf{x} \in \mathbb{R}^n$ that maximizes $f(\mathbf{x})$. Formally, we have

$$\forall \mathbf{x} \in \mathbb{R}^n : f(\mathbf{x}) \leq f\left(\arg \max_{\mathbf{x} \in \mathbb{R}^n} f(\mathbf{x})\right).$$

Of course, the function argmax is only defined when the maximum is unique. If the function f is differentiable, we know that a necessary condition for a vector $\hat{\mathbf{x}} \in \mathbb{R}^n$ to satisfy

$$\hat{\mathbf{x}} = \arg \max_{\mathbf{x} \in \mathbb{R}^n} f(\mathbf{x}) \quad \text{is that we must have} \quad \nabla f(\hat{\mathbf{x}}) = \mathbf{0},$$

i.e. the [gradient](#) of f , which we will write as ∇f , vanishes at the maximum. Remember that the gradient of f is defined as

$$\nabla f := \begin{pmatrix} \frac{\partial f}{\partial x_1} \\ \vdots \\ \frac{\partial f}{\partial x_n} \end{pmatrix}.$$




Figure 6.1: The function $x \mapsto \sin(x) - \frac{1}{2} \cdot x^2$.

Unfortunately, in many cases the equation

$$\nabla f(\hat{\mathbf{x}}) = \mathbf{0}$$

can not be solved explicitly. This is already true in the one-dimensional case, i.e. if $n = 1$. For example, consider the function $f : \mathbb{R} \rightarrow \mathbb{R}$ that is defined as

$$f(x) := \sin(x) - \frac{1}{2} \cdot x^2.$$

This function is shown in Figure 6.1 on page 91. From the graph of the function it is obvious that this function has a maximum somewhere between 0.6 and 0.8. In order to compute this maximum, we can compute the derivative of f . This derivative is given as

$$f'(x) = \cos(x) - x$$

As it happens, the equation $\cos(x) - x = 0$ does not seem to have a solution in **closed form**. Hence, we can only approximate the solution numerically.

The method of **gradient ascent** is a numerical method that can be used to find the maximum of a function

$$f : \mathbb{R}^n \rightarrow \mathbb{R}.$$

The basic idea is to take a vector $\mathbf{x}_0 \in \mathbb{R}^n$ as start value and define a sequence of vectors $(\mathbf{x}_n)_{n \in \mathbb{N}}$ such that we have

$$f(\mathbf{x}_{n+1}) \geq f(\mathbf{x}_n) \quad \text{for all } n \in \mathbb{N}.$$

Hopefully, this sequence will converge against $\hat{\mathbf{x}} = \arg \max_{\mathbf{x} \in \mathbb{R}^n} f(\mathbf{x})$. If we do not really know where to start our search, we define $\mathbf{x}_0 := \mathbf{0}$. In order to compute \mathbf{x}_{n+1} given \mathbf{x}_n , the idea is to move from \mathbf{x}_n in that direction where we have the biggest change in the values of f . This direction happens to be the gradient of f at \mathbf{x}_n . Therefore, the definition of \mathbf{x}_{n+1} is given as follows:

$$\mathbf{x}_{n+1} := \mathbf{x}_n + \alpha \cdot \nabla f(\mathbf{x}_n) \quad \text{for all } n \in \mathbb{N}_0.$$

Here, α is called the **step size**. It determines by how much we move in the direction of the gradient. In practice, it is best to adapt the step size dynamically during the iteration. Figure 6.2 shows how this is done.

```

1  findMaximum := procedure(f, gradF, start, eps) {
2      x      := start;
3      fx     := f(x);
4      alpha  := 1.0;
5      while (true) {
6          [xOld, fOld] := [x, fx];
7          x += alpha * gradF(x);
8          fx := f(x);
9          if (fx < fOld) {
10             alpha *= 0.5;
11             [x, fx] := [xOld, fOld];
12             continue; // start over
13         } else {
14             alpha *= 1.2;
15         }
16         if (abs(fx - fOld) <= abs(fx) * eps) {
17             return [x, fx];
18         }
19     }
20 };

```

Figure 6.2: The gradient ascent algorithm.

The function `findMaximum` takes four arguments:

1. f is the function that is to be maximized. It is assumed that f takes a vector $\mathbf{x} \in \mathbb{R}^n$ as its input and

that it returns a real number.

2. `gradF` is the gradient of `f`. It takes a vector $\mathbf{x} \in \mathbb{R}^n$ as its input and returns the vector $\nabla f(\mathbf{x})$.
3. `start` is the a vector from \mathbb{R}^n that is used as the value of \mathbf{x}_0 . In practice, we will often use $\mathbf{0} \in \mathbb{R}^n$ as the start vector.
4. `eps` is the precision that we need for the maximum. We will have to say more on how `eps` is exactly related to the precision later. As we are using double precision floating point arithmetic, it won't make sense to use a value for `eps` that is smaller than 10^{-15} .

Next, let us discuss the implementation of gradient ascent.

1. `x` is initialized with the parameter `start`. Hence, `start` is really the same as \mathbf{x}_0 .
2. `fx` is the value that the function `f` takes for the argument `x`.
3. `alpha` is the step size α . We initialize `alpha` as 1.0. It will be adapted dynamically.
4. The `while` loop starting in line 5 executes the iteration.
5. In each iteration, we store the values of \mathbf{x}_n and $f(\mathbf{x}_n)$ in the variables `xOld` and `fOld`.
6. Next, we compute \mathbf{x}_{n+1} in line 7 and compute the corresponding value $f(\mathbf{x}_{n+1})$ in line 8.
7. If we are unlucky, $f(\mathbf{x}_{n+1})$ is smaller than $f(\mathbf{x}_n)$. This happens if the step size α is too large. Hence, in this case we decrease the value of α , discard both \mathbf{x}_{n+1} and $f(\mathbf{x}_{n+1})$ and start over again.
8. Otherwise, \mathbf{x}_{n+1} is a better approximation of the maximum than \mathbf{x}_n . In order to increase the speed of the convergence of our algorithm we will then increase the step size α by 20%.
9. The idea of our implementation is to stop the iteration when the function values of $f(\mathbf{x}_{n+1})$ and $f(\mathbf{x}_n)$ do not differ by more than ε percent, or, to be more precise, if

$$f(\mathbf{x}_{n+1}) < f(\mathbf{x}_n) \cdot (1 + \varepsilon).$$

As the sequence $(f(\mathbf{x}_n))_{n \in \mathbb{N}}$ will be monotonically increasing, i.e. we have

$$f(\mathbf{x}_{n+1}) \geq f(\mathbf{x}_n) \quad \text{for all } n \in \mathbb{N},$$

the condition given above is sufficient. Now, if the increase of $f(\mathbf{x}_{n+1})$ is less than $f(\mathbf{x}_n) \cdot (1 + \varepsilon)$ we assume that we have reached the maximum with the required precision. In this case we return both the value of `x` and the corresponding function value `f(x)`.

The implementation of gradient ascent given above is not the most sophisticated variant of this algorithm. It should also be noted that there are algorithms that are more powerful than gradient ascent. The first of these methods is the [conjugate gradient method](#). A refinement of this method is the [BFGS-algorithm](#) that has been invented by Broyden, Fletcher, Goldfarb, and Shanno. Unfortunately, we do not have the time to discuss this algorithm. However, our implementation of gradient ascent is sufficient for our applications and as this is not a course on numerical analysis but rather on artificial intelligence we will not delve deeper into this topic but, instead, we refer readers interested in more efficient algorithms to the literature [[Sny05](#)]. If you ever need to find the maximum of a function numerically, you should try to use a predefined library routine that implements a state of the art algorithm.

6.3 Logistic Regression

In [logistic regression](#) we use a linear model that is combined with the [sigmoid function](#). Before we can discuss the details of logistic regression we need to define this function and state some of its properties.

6.3.1 The Sigmoid Function

Definition 5 (Sigmoid Function) The **sigmoid function** $S : \mathbb{R} \rightarrow [0, 1]$ is defined as

$$S(t) = \frac{1}{1 + \exp(-t)}.$$

Figure 6.3 on page 94 shows the sigmoid function. ◇




Figure 6.3: The sigmoid function.

Let us note some immediate consequences of the definition of the sigmoid function. As we have

$$\lim_{x \rightarrow -\infty} \exp(-x) = \infty, \quad \lim_{x \rightarrow +\infty} \exp(-x) = 0, \quad \text{and} \quad \lim_{x \rightarrow \infty} \frac{1}{x} = 0,$$

the sigmoid function has the following properties:

$$\lim_{t \rightarrow -\infty} S(t) = 0 \quad \text{and} \quad \lim_{t \rightarrow +\infty} S(t) = 1.$$

Another important property is the symmetry of the sigmoid function. Figure 6.3 shows that if the sigmoid function is shifted down by $\frac{1}{2}$, the resulting function is **centrally symmetric**, i.e. we have

$$S(-t) - \frac{1}{2} = -\left(S(t) - \frac{1}{2}\right).$$

Adding $\frac{1}{2}$ on both sides of this equation shows that this is equivalent to the equation

$$S(-t) = 1 - S(t),$$

The proof of this fact runs as follows:

$$\begin{aligned}
1 - S(t) &= 1 - \frac{1}{1 + \exp(-t)} && \text{by definition of } S(t) \\
&= \frac{1 + \exp(-t) - 1}{1 + \exp(-t)} && \text{common denominator} \\
&= \frac{\exp(-t)}{1 + \exp(-t)} && \text{simplify} \\
&= \frac{1}{1 + \exp(+t)} && \text{expand fraction by } \exp(t) \\
&= S(-t). \quad \square && \text{by definition of } S(-t)
\end{aligned}$$

The exponential function can be expressed via the sigmoid function. Let us start with the definition of the sigmoid function.

$$S(t) = \frac{1}{1 + \exp(-t)}$$

Multiplying this equation with the denominator yields

$$S(t) \cdot (1 + \exp(-t)) = 1.$$

Dividing both sides by $S(t)$ gives:

$$\begin{aligned}
1 + \exp(-t) &= \frac{1}{S(t)} \\
\Leftrightarrow \exp(-t) &= \frac{1}{S(t)} - 1 \\
\Leftrightarrow \exp(-t) &= \frac{1 - S(t)}{S(t)}
\end{aligned}$$

We highlight this formula, as we need it later

$$\exp(-t) = \frac{1 - S(t)}{S(t)}.$$

If we take the reciprocal of both sides of this equation, we have

$$\exp(t) = \frac{S(t)}{1 - S(t)}.$$

Applying the natural logarithm on both sides of this equation yields

$$t = \ln \left(\frac{S(t)}{1 - S(t)} \right).$$

This shows that the inverse of the sigmoid function is given as

$$S^{-1}(y) = \ln \left(\frac{y}{1 - y} \right).$$

This function is known as the **logit function**. Next, let us compute the derivative of $S(t)$, i.e. $S'(t) = \frac{dS}{dt}$. We have

$$\begin{aligned}
S'(t) &= -\frac{-\exp(-t)}{(1 + \exp(-t))^2} \\
&= \exp(-t) \cdot S(t)^2 \\
&= \frac{1 - S(t)}{S(t)} \cdot S(t)^2 \\
&= (1 - S(t)) \cdot S(t)
\end{aligned}$$

We have shown

$$S'(t) = (1 - S(t)) \cdot S(t).$$

We will later need the derivative of the logarithm of the logistic function. We define

$$L(t) := \ln(S(t)).$$

Then we have

$$\begin{aligned}
L'(t) &= \frac{S'(t)}{S(t)} && \text{by the chain rule} \\
&= \frac{(1 - S(t)) \cdot S(t)}{S(t)} \\
&= 1 - S(t) \\
&= S(-t)
\end{aligned}$$

As this is our most important result, we highlight it:

$$L'(t) = S(-t) \quad \text{where} \quad L(t) := \ln(S(t)).$$

6.3.2 The Model of Logistic Regression

We use the following model to compute the probability that an object with features \mathbf{x} will be of the given class:

$$P(y = +1 \mid \mathbf{x}; \mathbf{w}) = S(\mathbf{x} \cdot \mathbf{w}).$$

Here $\mathbf{x} \cdot \mathbf{w}$ denotes the dot product of the vectors \mathbf{x} and \mathbf{y} . It is assumed that \mathbf{x} contains a constant feature which always takes the value of 1. Seeing this model the first time you might think that this model is not very general and that it can only be applied in very special circumstances. However, the fact is that the features can be functions of arbitrary complexity and hence this model is much more general than it appears on first sight.

We assume that y can only take the values $+1$ or -1 , e.g. in the example of spam detection $y = 1$ if the email is spam and $y = -1$ otherwise. Since complementary probabilities add up to 1, we have

$$P(y = -1 \mid \mathbf{x}; \mathbf{w}) = 1 - P(y = +1 \mid \mathbf{x}; \mathbf{w}) = 1 - S(\mathbf{x} \cdot \mathbf{w}) = S(-\mathbf{x} \cdot \mathbf{w}).$$

Hence, we can combine the equations for $P(y = -1 \mid \mathbf{x}; \mathbf{w})$ and $P(y = +1 \mid \mathbf{x}; \mathbf{w})$ into a single equation

$$P(y \mid \mathbf{x}; \mathbf{w}) = S(y \cdot (\mathbf{x} \cdot \mathbf{w})).$$

Given N objects o_1, \dots, o_n with feature vectors $\mathbf{x}_1, \dots, \mathbf{x}_n$ we want to determine the weight vector \mathbf{w} such that the likelihood $\ell(\mathbf{X}, \mathbf{y})$ of all of our observations is maximized. This approach is called the **maximum likelihood estimation** of the weights. As we assume the probabilities of different observations are independent, the individual probabilities have to be multiplied to compute the overall likelihood $\ell(\mathbf{X}, \mathbf{y}; \mathbf{w})$ of a given training set:

$$\ell(\mathbf{X}, \mathbf{y}; \mathbf{w}) = \prod_{i=1}^N P(y_i | \mathbf{x}_i; \mathbf{w}).$$

Here, we have combined the different attribute vectors \mathbf{x}_i into the matrix \mathbf{X} . Since it is easier to work with sums than with products, instead of maximizing the function $\ell(\mathbf{X}, \mathbf{y}; \mathbf{w})$ we maximize the function

$$\ell\ell(\mathbf{X}, \mathbf{y}; \mathbf{w}) := \ln(\ell(\mathbf{X}, \mathbf{y}; \mathbf{w})).$$

As the natural logarithm is a monotone function, the functions $\ell(\mathbf{X}, \mathbf{y}; \mathbf{w})$ and $\ell\ell(\mathbf{X}, \mathbf{y}; \mathbf{w})$ take their maximum at the same value of \mathbf{w} . As we have

$$\ln(a \cdot b) = \ln(a) + \ln(b),$$

the natural logarithm of the likelihood is

$$\ell\ell(\mathbf{X}, \mathbf{y}; \mathbf{w}) = \sum_{i=1}^N \ln \left(S(y_i \cdot (\mathbf{x}_i \cdot \mathbf{w})) \right) = \sum_{i=1}^N L(y_i \cdot (\mathbf{x}_i \cdot \mathbf{w})).$$

Our goal is to maximize the likelihood. Since this is the same as maximizing the log-likelihood, we need to determine those values of the coefficients \mathbf{w} that satisfy

$$\frac{\partial}{\partial w_j} \ell\ell(\mathbf{X}, \mathbf{y}; \mathbf{w}) = 0.$$

In order to compute the partial derivative of $\ell\ell(\mathbf{X}, \mathbf{y}; \mathbf{w})$ with respect to the coefficients \mathbf{w} we need to compute the partial derivative of the dot product $\mathbf{x}_i \cdot \mathbf{w}$ with respect to the weights w_j . We define

$$h(\mathbf{w}) := \mathbf{x}_i \cdot \mathbf{w} = \sum_{k=1}^D x_{i,k} \cdot w_k.$$

Then we have

$$\frac{\partial}{\partial w_j} h(\mathbf{w}) = x_{i,j}.$$

Now we are ready to compute the partial derivative of $\ell\ell(\mathbf{X}, \mathbf{y}; \mathbf{w})$ with respect to \mathbf{w} :

$$\begin{aligned} & \frac{\partial}{\partial w_j} \ell\ell(\mathbf{X}, \mathbf{y}; \mathbf{w}) \\ &= \frac{\partial}{\partial w_j} \sum_{i=1}^N L(y_i \cdot (\mathbf{x}_i \cdot \mathbf{w})) \\ &= \sum_{i=1}^N y_i \cdot x_{i,j} \cdot S(-y_i \cdot (\mathbf{x}_i \cdot \mathbf{w})), \quad \text{since} \quad \frac{dL(x)}{dx} = S(-x). \end{aligned}$$

Hence, the partial derivative of the log-likelihood function is given as follows:

$$\frac{\partial}{\partial w_j} \ell\ell(\mathbf{X}, \mathbf{y}; \mathbf{w}) = \sum_{i=1}^N y_i \cdot x_{i,j} \cdot S(-y_i \cdot \mathbf{x}_i \cdot \mathbf{w})$$

Next, we have to find the value of \mathbf{w} such that

$$\sum_{i=1}^N y_i \cdot x_{i,j} \cdot S(-y_i \cdot \mathbf{x}_i \cdot \mathbf{w}) = 0 \quad \text{for all } j \in \{1, \dots, D\}.$$

These are D equations for the D variables w_1, \dots, w_D . Due to the occurrence of the sigmoid function, these equations are nonlinear. We can not solve these equations explicitly. Nevertheless, our computation of the

gradient of the log-likelihood was not for nought: We will use gradient ascent in order to find the value of \mathbf{w} that maximizes the log-likelihood. This method has been outlined in the previous section.

6.3.3 Implementing Logistic Regression

In order to implement logistic regression we need a data structure for tabular data. Figure 6.4 on page 98 shows the class `table` that can be used to administer this kind of data. Figure 6.4 shows an example of tabular data that is stored in a `csv` file. In this case, the data stores the hours a student has learned for a particular exam and the fact whether the student has passed or failed. The first column stores pass or fail, where a pass is coded using the number 1, while a fail is coded as 0. The second column stores the number of hours that the student has learned in order to pass the exam.

```

1  class table(columnNames, types, data) {
2      mColumnNames := columnNames;
3      mTypes       := types;
4      mData        := data;
5
6      static {
7          getColumnNames := [ ] |-> mColumnNames;
8          getTypes       := [ ] |-> mTypes;
9          getData        := [ ] |-> mData;
10         getRow         := [r] |-> mData[r];
11         getLength      := [ ] |-> #mData;
12
13         head := procedure(limit := 10) {
14             print(mColumnNames);
15             print(mTypes);
16             for (i in [1 .. limit]) {
17                 print(mData[i]);
18             }
19         };
20     }
21 }
22 readTable := procedure(fileName, types) {
23     all := readFile(fileName);
24     columnNames := split(all[1], ',\s*');
25     data := [];
26     for (i in [2 .. #all]) {
27         row := split(all[i], ',\s*');
28         data[i-1] := [eval("$type$(ss$)") : [type, s] in types >< row];
29     }
30     return table(columnNames, types, data);
31 };

```

Figure 6.4: A class to represent tabular data.

There is no need for us to discuss every detail of the implementation of the class `table`. The important thing to note is that the data is stored as a list of lists in the member variable `mData`. Each of the inner lists corresponds to one row of the `csv` file. This member variable can be accessed using the function `getData`. The function `readTable` has the responsibility to read a `csv` file and to convert it into an object of class `table`. In order to do this, it has to be called with two arguments. The first argument is the file name, the second argument is a list of the types of each column in the `csv` file. For example, to read the file “`exam.csv`” we would call `readTable` as follows:

```
readTable("exam.csv", ["int", "double"]).
```

	Pass	Hours
1	0	0.50
2	0	0.75
3	0	1.00
4	0	1.25
5	0	1.50
6	0	1.75
7	1	1.75
8	0	2.00
9	1	2.25
10	0	2.50
11	1	2.75
12	0	3.00
13	1	3.25
14	0	3.50
15	1	4.00
16	1	4.25
17	1	4.50
18	1	4.75
19	1	5.00
20	1	5.50
21	1	5.50

Figure 6.5: Results of an exam.

The program shown in Figure 6.6 on page 100 implements logistic regression. As there are a number of subtle points that might easily be overlooked otherwise, we proceed to discuss this program line by line.

1. First, we have to load both the class `table` and our implementation of gradient ascent that has already been discussed in Section 6.2.
2. Line 4 implements the sigmoid function

$$S(x) = \frac{1}{1 + \exp(-x)}.$$

3. Line 5 starts the implementation of the natural logarithm of the sigmoid function, i.e. we implement

$$L(x) = \ln(S(X)) = \ln\left(\frac{1}{1 + \exp(-x)}\right) = -\ln(1 + \exp(-x)).$$

The implementation is more complicated than you might expect. The reason has to do with overflow. Consider values of x that are smaller than, say, -1000 . The problem is that the expression `exp(1000)` evaluates to `Infinity`, which represents the mathematical value ∞ . But then `1 + exp(1000)` is also `Infinity` and finally `log(1 + exp(1000))` is `Infinity`. However, in reality we have

$$\ln(1 + \exp(1000)) \approx 1000.$$

The argument works as follows:

```

1  load("table.stlx");
2  load("gradient-ascent.stlx");
3
4  sigmoid := procedure(x) { return 1.0 / (1.0 + exp(-x)); };
5  logSigmoid := procedure(x) {
6      if (x >= -100) {
7          return -log(1.0 + exp(-x));
8      } else {
9          return x;
10     }
11 };
12 ll := procedure(X, y, w) {
13     result := 0;
14     for (i in [1 .. #X]) {
15         result += logSigmoid(y[i] * (X[i] * w));
16     }
17     return result;
18 };
19 gradLL := procedure(X, y, w) {
20     result := [];
21     for (j in [1 .. #X[1]]) {
22         result[j] := 0;
23         for (i in [1 .. #X]) {
24             result[j] += y[i] * X[i][j] * sigmoid((-y[i]) * (X[i] * w));
25         }
26     }
27     return la_vector(result);
28 };
29 logisticRegressionFile := procedure(fileName, types) {
30     csv := readTable(fileName, types);
31     data := csv.getData();
32     number := #data;
33     dmnsn := #data[1];
34     yList := [];
35     xList := [];
36     for (i in [1 .. number]) {
37         yList[i] := data[i][1];
38         xList[i] := la_vector([1.0] + data[i][2..]);
39     }
40     X := la_matrix(xList);
41     y := la_vector([2 * y - 1 : y in yList]);
42     start := la_vector([0.0 : i in [1 .. dmnsn]]);
43     eps := 10 ** -15;
44     f := w |> ll(X, y, w);
45     gradF := w |> gradLL(X, y, w);
46     return findMaximum(f, gradF, start, eps)[1];
47 };

```

Figure 6.6: An implementation of logistic regression.

$$\begin{aligned}
\ln(1 + \exp(x)) &= \ln(\exp(x) \cdot (1 + \exp(-x))) \\
&= \ln(\exp(x)) + \ln(1 + \exp(-x)) \\
&= x + \ln(1 + \exp(-x)) \\
&\approx x + \ln(1) + \exp(-x) && \text{Taylor expansion of } \ln(1 + x) \\
&= x + 0 + \exp(-x) \\
&\approx x && \text{since } \exp(-x) \approx 0 \text{ for large } x
\end{aligned}$$

This is the reason that `logSigmoid` returns `x` if `x` is less than `-100`.

4. The function `ll(x, y, w)` computes the log-likelihood $\ell(\mathbf{X}, \mathbf{y}; \mathbf{w}) = \sum_{i=1}^N L(y_i \cdot (\mathbf{x}_i \cdot \mathbf{w}))$. Here L denotes the natural logarithm of the sigmoid of the argument. It is assumed that \mathbf{X} is a matrix. Every observation corresponds to a row in this matrix, i.e. the vector \mathbf{x}_i is the feature vector containing the features of the i -th observation. \mathbf{y} is a vector describing the outcomes, i.e. the elements of this vector are either `+1` or `-1`. Finally, \mathbf{w} is the vector of coefficients.

5. The function `gradLL(x, y, w)` computes the gradient of the log-likelihood according to the formula

$$\frac{\partial}{\partial w_j} \ell(\mathbf{X}, \mathbf{y}; \mathbf{w}) = \sum_{i=1}^N y_i \cdot x_{i,j} \cdot S(-y_i \cdot \mathbf{x}_i \cdot \mathbf{w}).$$

The different components of this gradient are combined into a vector. The arguments are the same as the arguments to the log-likelihood.

6. Finally, the function `logisticRegressionFile` takes two arguments. The first argument is the name of the csv file containing the data, while the second argument is a list specifying the types of the columns. The elements of this list have to be either `"int"` or `"double"`. The task of this function is to read the csv file, convert the data in the matrix \mathbf{X} and the vector \mathbf{y} , and then use the method of gradient ascent to find the coefficients \mathbf{w} that maximize the likelihood.

If we run the function `logisticRegressionFile` using the data shown in Figure 6.5 via the command

```
logisticRegressionFile("exam.csv", ["int", "double"]);
```

the resulting coefficients are:

```
<<-4.077649741107752 1.5046211108850898>>
```

This shows that the probability $P(h)$ that a student who has studied for h hours will pass the exam is given approximately as follows:

$$P(h) \approx \frac{1}{1 + \exp(4.1 - 1.5 \cdot h)}$$

Figure 6.7 shows a plot of the probability $P(x)$. This figure has been taken from the [Wikipedia article on logistic regression](#). It has been created by [Michaelg2015](#).




Figure 6.7: Probability of passing an exam versus hours of studying.

Chapter 7

Neural Networks

In this chapter, we discuss **neural networks**. Many of the most visible breakthroughs in artificial intelligence have been achieved through the use of neural networks:

1. The current system used by Google to **automatically translate** web pages is called “**Google Neural Machine Translation**” and, as the name suggests, is based on neural networks.
2. **AlphaGo** uses neural networks together with tree search [SHM⁺16]. It has **recently** beaten **Lee Sedol** in the game of go. At that time, Lee Sedol was ranked third among the **top go players**.
3. **Autonomous driving** makes heavy use of neural networks.

The list given above is far from being complete. In this chapter, we will only discuss **feedforward** neural networks. Although recently **recurrent neural networks** have gotten a lot of attention, these type of neural networks are more difficult to train and are therefore beyond the scope of this introduction. The rest of this chapter is strongly influenced by the online book

<http://neuralnetworksanddeeplearning.com/index.html>

that has been written by Michael Nielsen [Nie15]. This book is easy to read, carefully written, and free to access. I recommend this book to anybody who wants to dive deeper into the fascinating topic of neural networks.

7.1 Feedforward Neural Networks

A neural network is built from **neurons**. Neural networks are inspired by biological **neurons**. However, in order to understand artificial neural networks it is not necessary to know how biological neurons work and it is definitely not necessary to understand how networks of biological neurons, i.e. brains, work¹. Instead, we develop a mathematical abstraction of neurons that will serve as the foundation of the theory developed in this chapter. At the abstraction level that we are looking at neural networks, a single neuron with n inputs is defined as a pair $\langle \mathbf{w}, b \rangle$ where the vector $\mathbf{w} \in \mathbb{R}^m$ is called the **weight vector** and the number $b \in \mathbb{R}$ is called the **bias**. Conceptually, a neuron is a function p that maps an input vector $\mathbf{x} \in \mathbb{R}^m$ into the interval $[0, 1]$. This function is defined as follows:

$$p(\mathbf{x}; \mathbf{w}, b) := a(\mathbf{x} \cdot \mathbf{w} + b),$$

where a is called the **activation function**. In our applications, we will always use the sigmoid function as our activation function, i.e. we have

$$a(t) := S(t) = \frac{1}{1 + \exp(-t)}.$$

The function p modelling the neuron can be written more explicitly using index notation. If

¹ Actually, when it comes to brains, although there are many speculations, surprisingly little is known for a fact.

$$\mathbf{w} = \langle w_1, \dots, w_m \rangle^\top$$

is the weight vector and

$$\mathbf{x} = \langle x_1, \dots, x_m \rangle^\top$$

is the input vector, then we have

$$p(\mathbf{x}; \mathbf{w}, b) = S \left(\left(\sum_{i=1}^m x_i \cdot w_i \right) + b \right).$$

If we compare $p(\mathbf{x}; \mathbf{w}, b)$ to a similar function appearing in the last chapter, you will notice that so far a neuron works just like logistic regression. The only difference is that the bias b is now explicit in our notation. In logistic regression, we had assumed that the first component x_1 of our feature vector \mathbf{x} was always equal to 1. This assumption enabled us to incorporate the bias b into the weight vector \mathbf{w} .

A **feedforward neural network** is a layered network of neurons. Formally, the **topology** of a neural network is given by a number $L \in \mathbb{N}$ and a list $[m(1), \dots, m(L)]$ of L natural numbers. The number L is called the **number of layers** and for $i \in \{2, \dots, L\}$ the number $m(i)$ is the number of neurons in the i -th layer. The first layer is called the **input layer**. The input layer does not contain neurons but instead just contains **input nodes**. The last layer (i.e. the layer with index L) is called the **output layer** and the remaining layers are called **hidden layers**. If there is more than one hidden layer, the neural network is called a **deep neural network**.

As the first layer is the input layer, the **input dimension** is defined as $m(1)$. Similarly, the **output dimension** is defined as $m(L)$. Every node in the i -th layer is connected to every node in the $(i+1)$ -th layer via a **weight**. The weight $w_{j,k}^{(i)}$ is the weight of the connection from the k -th neuron in layer $i-1$ to the j -th neuron in layer i . The weights in layer i are combined into the **weight matrix** $W^{(i)}$ of the layer i : This matrix is defined as

$$W^{(i)} := (w_{j,k}^{(i)}).$$

Note that $W^{(i)}$ is an $m(i) \times m(i-1)$ matrix, i.e. we have

$$W^{(i)} \in \mathbb{R}^{m(i) \times m(i-1)}.$$

The j -th neuron in layer i has the **bias** $b_j^{(i)}$. These biases of layer i are combined into the **bias vector**

$$\mathbf{b}^{(i)} := \langle b_1^{(i)}, \dots, b_{m(i)}^{(i)} \rangle^\top.$$

Then, the **activation** of the j -th neuron in layer i is denoted as $a_j^{(i)}$ and is defined recursively as follows:

1. For the input layer we have

$$a_j^{(1)} := x_j. \tag{FF1}$$

To put it differently, the input vector \mathbf{x} is the activation of the input nodes.

2. For all other layers we have

$$a_j^{(i)}(\mathbf{x}) := S \left(\left(\sum_{k=1}^{m(i-1)} w_{j,k}^{(i)} \cdot a_k^{(i-1)}(\mathbf{x}) \right) + b_j^{(i)} \right) \quad \text{for all } i \in \{2, \dots, L\}. \tag{FF2}$$

The **activation vector** of layer i is defined as

$$\mathbf{a}^{(i)} := \langle a_1^{(i)}, \dots, a_{m(i)}^{(i)} \rangle^\top.$$

The output of our neural network for an input \mathbf{x} is given by the neurons in the output layer, i.e. the output vector $\mathbf{o}(\mathbf{x}) \in \mathbb{R}^{m(L)}$ is defined as

$$\mathbf{o}(\mathbf{x}) := \langle a_1^{(L)}(\mathbf{x}), \dots, a_{m(L)}^{(L)}(\mathbf{x}) \rangle^\top = \mathbf{a}^{(L)}(\mathbf{x}).$$

Note that the equations (FF1) and (FF2) describe how information propagates through the neural network:

1. Initially, the input vector \mathbf{x} is given and stored in the input layer of the neural network:

$$\mathbf{a}^{(1)}(\mathbf{x}) := \mathbf{x}.$$

2. The first layer of neurons, which is the second layer of nodes, is activated and computes the activation vector $\mathbf{a}^{(2)}$ according to the formula

$$\mathbf{a}^{(2)}(\mathbf{x}) := S(W^{(2)} \cdot \mathbf{a}^{(1)}(\mathbf{x}) + \mathbf{b}^{(2)}) = S(W^{(2)} \cdot \mathbf{x} + \mathbf{b}^{(2)}).$$

3. The second layer of neurons, which is the third layer of nodes, is activated and computes the activation vector $\mathbf{a}^{(3)}(\mathbf{x})$ according to the formula

$$\mathbf{a}^{(3)}(\mathbf{x}) := S(W^{(3)} \cdot \mathbf{a}^{(2)}(\mathbf{x}) + \mathbf{b}^{(3)}) = S(W^{(3)} \cdot S(W^{(2)} \cdot \mathbf{x} + \mathbf{b}^{(2)}) + \mathbf{b}^{(3)})$$

4. This proceeds until the output layer is reached and the output

$$\mathbf{o}(\mathbf{x}) := \mathbf{a}^{(L)}(\mathbf{x})$$

has been computed. Note that every neuron of the neural network performs logistic regression.

Next, we assume that we have n training examples

$$\langle \mathbf{x}^{(i)}, \mathbf{y}^{(i)} \rangle \quad \text{for } i = 1, \dots, n$$

such that

$$\mathbf{x}^{(i)} \in \mathbb{R}^{m(1)} \text{ and } \mathbf{y}^{(i)} \in \mathbb{R}^{m(L)}$$

Our goal is to choose the weight matrices $W^{(l)}$ and the bias vectors $\mathbf{b}^{(l)}$ in a way such that

$$\mathbf{o}(\mathbf{x}^{(i)}) = \mathbf{y}^{(i)} \quad \text{for all } i \in \{1, \dots, n\}.$$

Unfortunately, in general we will not be able to achieve equality for all $i \in \{1, \dots, n\}$. Therefore, our goal is to minimize the error instead. To be more precise, the quadratic error cost function is defined as

$$C(W^{(2)}, \dots, W^{(L)}, \mathbf{b}^{(2)}, \dots, \mathbf{b}^{(L)}; \mathbf{x}^{(1)}, \mathbf{y}^{(1)}, \dots, \mathbf{x}^{(n)}, \mathbf{y}^{(n)}) := \frac{1}{2 \cdot n} \cdot \sum_{i=1}^n (\mathbf{o}(\mathbf{x}^{(i)}) - \mathbf{y}^{(i)})^2.$$

Note that the cost function is additive in the training examples $\langle \mathbf{x}^{(i)}, \mathbf{y}^{(i)} \rangle$. In order to simplify the notation we define

$$C_{\mathbf{x}, \mathbf{y}}(W^{(2)}, \dots, W^{(L)}, \mathbf{b}^{(2)}, \dots, \mathbf{b}^{(L)}) := \frac{1}{2} \cdot (\mathbf{a}^{(L)}(\mathbf{x}) - \mathbf{y})^2,$$

i.e. $C_{\mathbf{x}, \mathbf{y}}$ is the part of the cost function that is associated with a single training example $\langle \mathbf{x}, \mathbf{y} \rangle$. Then, we have

$$C(W^{(2)}, \dots, W^{(L)}, \mathbf{b}^{(2)}, \dots, \mathbf{b}^{(L)}; \mathbf{x}^{(1)}, \mathbf{y}^{(1)}, \dots, \mathbf{x}^{(n)}, \mathbf{y}^{(n)}) := \frac{1}{n} \cdot \sum_{i=1}^n C_{\mathbf{x}^{(i)}, \mathbf{y}^{(i)}}(W^{(2)}, \dots, W^{(L)}, \mathbf{b}^{(2)}, \dots, \mathbf{b}^{(L)}).$$

As the notation

$$C_{\mathbf{x}, \mathbf{y}}(W^{(2)}, \dots, W^{(L)}, \mathbf{b}^{(2)}, \dots, \mathbf{b}^{(L)})$$

is far too heavy, we will abbreviate this term as $C_{\mathbf{x}, \mathbf{y}}$ in the following discussion of the backpropagation algorithm. Similarly, we abbreviate the quadratic error cost function as C . Our goal is to choose the weight matrices $W^{(l)}$ and the bias vectors $\mathbf{b}^{(l)}$ such that the quadratic error cost function C is minimized. We will use a variation of gradient descent to find this minimum².

² In logistic regression we have tried to *maximize* the log-likelihood. Here, instead we *minimize* the quadratic error cost function. Hence, instead of gradient *ascent* we use gradient *descent*.

7.2 Backpropagation

There are two reasons for the recent success of neural networks.

1. The computing power that is available today has vastly increased in the last 20 years. For example, today the **AMD Vega 10** graphic card offers about 12.5 teraflops in single precision performance. It consumes about 300 watt. Contrast this with **ASCI White**, which was the most powerful supercomputer in 2000: In 2000 when it topped the rankings of the supercomputers, it offered a performance of 7.2 teraflops and needed 6 megawatt to operate. The cost to build ASCI White where about 110,000,000\$. On the contrary, the AMD Vega 10 is expected to cost about 1,000\$.
2. The breakthrough in the theory of neural networks was the rediscovering of the **backpropagation algorithm** by David Rumelhart, Geoffrey Hinton, and Ronald Williams [RHW86] in 1986.

Essentially, the backpropagation algorithm is an efficient way to compute the partial derivatives of the cost function C with respect to the weights $w_{j,k}^{(l)}$ and the biases $b_j^{(l)}$. Before we can proceed to compute these partial derivatives, we need to define some auxiliary variables.

7.2.1 Definition of some Auxiliary Variables

We start by defining the auxiliary variables $z_j^{(l)}$. The expressions $z_j^{(l)}$ are defined as the inputs of the activation function S of the j -th neuron in layer l :

$$z_j^{(l)} := \left(\sum_{k=1}^{m(l-1)} w_{j,k}^{(l)} \cdot a_k^{(l-1)} \right) + b_j^{(l)} \quad \text{for all } j \in \{1, \dots, m(l)\} \text{ and } l \in \{2, \dots, L\}.$$

Of course, the term $a_k^{(l-1)}$ really is a function of the input vector \mathbf{x} . However, it is better to suppress this dependence in the notation since otherwise the formulæ get too cluttered. Essentially, $z_j^{(l)}$ is the input to the sigmoid function when the activation $a_j^{(l)}$ is computed, i.e. we have

$$a_j^{(l)} = S(z_j^{(l)}).$$

We will see that the partial derivatives of the cost function $C_{\mathbf{x},\mathbf{y}}$ with respect to both the weights $w_{j,k}^{(l)}$ and the biases $b_j^{(l)}$ can be computed easily if we first compute the partial derivatives of $C_{\mathbf{x},\mathbf{y}}$ with respect to $z_j^{(l)}$. Therefore we define

$$\varepsilon_j^{(l)} := \frac{\partial C_{\mathbf{x},\mathbf{y}}}{\partial z_j^{(l)}} \quad \text{for all } j \in \{1, \dots, m(l)\} \text{ and } l \in \{2, \dots, L\},$$

that is we regard $C_{\mathbf{x},\mathbf{y}}$ as a function of the $z_j^{(l)}$ and take the partial derivatives according to these variables. Note that $\varepsilon_j^{(l)}$ does depend on both \mathbf{x} and \mathbf{y} . Since the notation would get very cumbersome if we would write $\varepsilon(\mathbf{x},\mathbf{y})_j^{(l)}$, we regard \mathbf{x} and \mathbf{y} as fixed for now. Next, the quantities $\varepsilon_j^{(l)}$ are combined into a vector:

$$\boldsymbol{\varepsilon}^{(l)} := \begin{pmatrix} \varepsilon_1^{(l)} \\ \vdots \\ \varepsilon_{m(l)}^{(l)} \end{pmatrix}.$$

For reasons that will be explained later, this quantity $\boldsymbol{\varepsilon}^{(l)}$ is called the **error in layer l** .

7.2.2 The Hadamard Product

Later, we will have need of the **Hadamard product** of two vectors. Assume that $\mathbf{x}, \mathbf{y} \in \mathbb{R}^n$. The **Hadamard product** of \mathbf{x} and \mathbf{y} is defined by multiplying the vectors elementwise:

$$\begin{pmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{pmatrix} \odot \begin{pmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{pmatrix} := \begin{pmatrix} x_1 \cdot y_1 \\ x_2 \cdot y_2 \\ \vdots \\ x_n \cdot y_n \end{pmatrix},$$

i.e. the i -th component of the Hadamard product $\mathbf{x} \odot \mathbf{y}$ is the product of the i -th component of \mathbf{x} with the i -th component of \mathbf{y} .

7.2.3 Backpropagation: The Equations

Now we are ready to state the **backpropagation equations**. The first of these four equations reads as follows:

$$\varepsilon_j^{(L)} = (a_j^{(L)} - y_j) \cdot S'(z_j^{(L)}) \quad \text{for all } j \in \{1, \dots, m(L)\}, \quad (\text{BP1})$$

where $S'(x)$ denotes the derivative of the sigmoid function. We have shown in Chapter 6 that

$$S'(x) = (1 - S(x)) \cdot S(x)$$

holds. The equation (BP1) can also be written in vectorized form using the Hadamard product:

$$\boldsymbol{\varepsilon}^{(L)} = (\mathbf{a}^{(L)} - \mathbf{y}) \odot S'(\mathbf{z}^{(L)}) \quad (\text{BP1v})$$

Here, we have **vectorized** the application of the function S' to the vector $\mathbf{z}^{(L)}$, i.e. the expression $S'(\mathbf{z}^{(L)})$ is defined as follows:

$$S' \left(\begin{pmatrix} z_1^{(L)} \\ \vdots \\ z_{m(L)}^{(L)} \end{pmatrix} \right) := \begin{pmatrix} S'(z_1^{(L)}) \\ \vdots \\ S'(z_{m(L)}^{(L)}) \end{pmatrix}.$$

The next equation computes $\varepsilon_j^{(l)}$ for $l < L$.

$$\varepsilon_j^{(l)} = \sum_{i=1}^{m(l+1)} w_{i,j}^{(l+1)} \cdot \varepsilon_i^{(l+1)} \cdot S'(z_j^{(l)}) \quad \text{for all } j \in \{1, \dots, m(l)\} \text{ and } l \in \{2, \dots, L-1\}. \quad (\text{BP2})$$

This equation is more succinct in vectorized notation:

$$\boldsymbol{\varepsilon}^{(l)} = \left((\mathbf{W}^{(l+1)})^\top \cdot \boldsymbol{\varepsilon}^{(l+1)} \right) \odot S'(\mathbf{z}^{(l)}) \quad \text{for all } l \in \{2, \dots, L-1\}. \quad (\text{BP2v})$$

Note that this equation computes $\boldsymbol{\varepsilon}^{(l)}$ in terms of $\boldsymbol{\varepsilon}^{(l+1)}$: The error $\boldsymbol{\varepsilon}^{(l+1)}$ at layer $l+1$ is **propagated backwards** through the neural network to produce the error $\boldsymbol{\varepsilon}^{(l)}$ at layer l . This is the reason for calling the algorithm **backpropagation**.

Next, we have to compute the partial derivative of $C_{\mathbf{x}, \mathbf{y}}$ with respect to the bias of the j -th neuron in layer l , which is denoted as $b_j^{(l)}$. We have

$$\frac{\partial C_{\mathbf{x}, \mathbf{y}}}{\partial b_j^{(l)}} = \varepsilon_j^{(l)} \quad \text{for all } j \in \{1, \dots, m(l)\} \text{ and } l \in \{2, \dots, L\} \quad (\text{BP3})$$

In vectorized notation, this equation takes the following form:

$$\nabla_{\mathbf{b}^{(l)}} C_{\mathbf{x}, \mathbf{y}} = \boldsymbol{\varepsilon}^{(l)} \quad \text{for all } l \in \{2, \dots, L\} \quad (\text{BP3v})$$

Here, $\nabla_{\mathbf{b}^{(l)}} C_{\mathbf{x}, \mathbf{y}}$ denotes the gradient of $C_{\mathbf{x}, \mathbf{y}}$ with respect to the bias $\mathbf{b}^{(l)}$. Finally, we can compute the partial derivative of $C_{\mathbf{x}, \mathbf{y}}$ with respect to the weights:

$$\frac{\partial C_{\mathbf{x}, \mathbf{y}}}{\partial w_{j,k}^{(l)}} = a_k^{(l-1)} \cdot \varepsilon_j^{(l)} \quad \text{for all } j \in \{1, \dots, m(l)\}, k \in \{1, \dots, m(l-1)\}, \text{ and } l \in \{2, \dots, L\} \quad (\text{BP4})$$

In vectorized notation, this equation can be written as:

$$\nabla_{W^{(l)}} C_{\mathbf{x}, \mathbf{y}} = \boldsymbol{\varepsilon}^{(l)} \cdot (\mathbf{a}^{(l-1)})^\top \quad \text{for all } l \in \{2, \dots, L\} \quad (\text{BP4v})$$

Here, the expression $\boldsymbol{\varepsilon}^{(l)} \cdot (\mathbf{a}^{(l-1)})^\top$ denotes the matrix product of the column vector $\boldsymbol{\varepsilon}^{(l)}$ that is regarded as an $m(l) \times 1$ matrix and the row vector $(\mathbf{a}^{(l-1)})^\top$ that is regarded as an $1 \times m(l-1)$ matrix.

The equations (BP3) and (BP4) show why it was useful to introduce the numbers $\varepsilon_j^{(l)}$: These numbers enable us to compute the partial derivatives of the cost function with respect to both the biases and the weights. Furthermore, the equations (BP1) and (BP2) show how these numbers can be computed. An implementation of backpropagation should use the vectorized versions of these equations since this is more efficient for two reasons:

1. Interpreted languages like SETLX, *Python*, or *Octave* take much more time to execute a loop than to execute a simple matrix-vector multiplication. The reason is that in a loop, in addition to executing the statement a given number of times, the statement has to be interpreted every time it is executed.
2. Languages that are optimized for machine learning often take care to delegate the execution of matrix operations to the graphical coprocessor which is optimized for these kinds of operations.

7.2.4 Proof of the Backpropagation Equations

Next, we prove the backpropagation equations. Although the proof is a bit tedious, it should be accessible: The **chain rule** of **multivariate calculus** is all that is needed to understand why the backpropagation equations are true. As a reminder, the chain rule in multivariate calculus works as follows: Assume that the functions $f = f(\mathbf{y})$ and $g = g(\mathbf{x})$ where $\mathbf{y} \in \mathbb{R}^k$, $\mathbf{x} \in \mathbb{R}^n$, $g(\mathbf{x}) \in \mathbb{R}^k$, and $f(\mathbf{y}) \in \mathbb{R}$ are differentiable³. So we have

$$f : \mathbb{R}^k \rightarrow \mathbb{R} \quad \text{and} \quad g : \mathbb{R}^n \rightarrow \mathbb{R}^k.$$

If the function $h : \mathbb{R}^n \rightarrow \mathbb{R}$ is defined as

$$h(\mathbf{x}) := f(g(\mathbf{x})) \quad \text{for all } \mathbf{x} \in \mathbb{R}^n,$$

then the partial derivative of h with respect to x_j satisfies

$$\frac{\partial h}{\partial x_j} = \sum_{i=1}^k \frac{\partial f}{\partial y_i} \cdot \frac{\partial g_i}{\partial x_j}.$$

Remember that we have defined the numbers $\varepsilon_j^{(l)}$ as

$$\varepsilon_j^{(l)} = \frac{\partial C_{\mathbf{x}, \mathbf{y}}}{\partial z_j^{(l)}},$$

while the numbers $z_j^{(l)}$ have been defined as

$$z_j^{(l)} := \left(\sum_{k=1}^{m(l-1)} w_{j,k}^{(l)} \cdot a_k^{(l-1)}(\mathbf{x}) \right) + b_j^{(l)}.$$

Since the quadratic error cost function $C_{\mathbf{x}, \mathbf{y}}$ for the training example $\langle \mathbf{x}, \mathbf{y} \rangle$ has been defined in terms of the activation $\mathbf{a}^{(L)}$ as

$$C_{\mathbf{x}, \mathbf{y}} = \frac{1}{2} \cdot (\mathbf{a}^{(L)}(\mathbf{x}) - \mathbf{y})^2$$

³ If this text had been written in German, I would have said that f and g are “total differenzierbar”.

and we have $\mathbf{a}^{(L)}(\mathbf{x}) = S(\mathbf{z}^{(L)})$, the chain rule tells us that $\varepsilon_j^{(L)}$ can be computed as follows:

$$\begin{aligned}
\varepsilon_j^{(L)} &= \frac{\partial C_{\mathbf{x}, \mathbf{y}}}{\partial z_j^{(L)}} \\
&= \frac{\partial}{\partial z_j^{(L)}} \frac{1}{2} \cdot (\mathbf{a}^{(L)}(\mathbf{x}) - \mathbf{y})^2 \\
&= \frac{1}{2} \cdot \frac{\partial}{\partial z_j^{(L)}} \sum_{i=1}^{m(L)} (a_i^{(L)}(\mathbf{x}) - y_i)^2 \\
&= \frac{1}{2} \cdot \frac{\partial}{\partial z_j^{(L)}} \sum_{i=1}^{m(L)} (S(z_i^{(L)}) - y_i)^2 \\
&= \frac{1}{2} \cdot \sum_{i=1}^{m(L)} 2 \cdot (S(z_i^{(L)}) - y_i) \cdot \frac{\partial}{\partial z_j^{(L)}} S(z_i^{(L)}) \\
&= \sum_{i=1}^{m(L)} (S(z_i^{(L)}) - y_i) \cdot S'(z_i^{(L)}) \cdot \frac{\partial z_i^{(L)}}{\partial z_j^{(L)}} \\
&= \sum_{i=1}^{m(L)} (S(z_i^{(L)}) - y_i) \cdot S'(z_i^{(L)}) \cdot \delta_{i,j} \\
&= (S(z_j^{(L)}) - y_j) \cdot S'(z_j^{(L)}) \\
&= (a_j^{(L)} - y_j) \cdot S'(z_j^{(L)})
\end{aligned}$$

Thus we have proved equation [BP1](#). Next, let us compute $\varepsilon_j^{(l)}$ for $l < L$. We have

$$\begin{aligned}
\varepsilon_j^{(l)} &= \frac{\partial C_{\mathbf{x}, \mathbf{y}}}{\partial z_j^{(l)}} \\
&= \sum_{i=1}^{m(l+1)} \frac{\partial C_{\mathbf{x}, \mathbf{y}}}{\partial z_i^{(l+1)}} \cdot \frac{\partial z_i^{(l+1)}}{\partial z_j^{(l)}} \quad \text{using the chain rule} \\
&= \sum_{i=1}^{m(l+1)} \varepsilon_i^{(l+1)} \cdot \frac{\partial z_i^{(l+1)}}{\partial z_j^{(l)}} \quad \text{using the definition of } \varepsilon_i^{(l+1)}
\end{aligned}$$

In order to proceed, we have to remember the definition of $z_i^{(l+1)}$. We have

$$z_i^{(l+1)} = \left(\sum_{k=1}^{m(l)} w_{i,k}^{(l+1)} \cdot S(z_k^{(l)}) \right) + b_i^{(l+1)}$$

Therefore, the partial derivatives $\frac{\partial z_i^{(l+1)}}{\partial z_j^{(l)}}$ can be computed as follows:

$$\begin{aligned}
\frac{\partial z_i^{(l+1)}}{\partial z_j^{(l)}} &= \sum_{k=1}^{m(l)} w_{i,k}^{(l+1)} \cdot S'(z_k^{(l)}) \cdot \frac{\partial z_k^{(l)}}{\partial z_j^{(l)}} \\
&= \sum_{k=1}^{m(l)} w_{i,k}^{(l+1)} \cdot S'(z_k^{(l)}) \cdot \delta_{k,j} \\
&= w_{i,j}^{(l+1)} \cdot S'(z_j^{(l)})
\end{aligned}$$

If we substitute this expression back into the result we got for $\varepsilon_j^{(l)}$ we have shown the following:

$$\begin{aligned}\varepsilon_j^{(l)} &= \sum_{i=1}^{m^{(l+1)}} \varepsilon_i^{(l+1)} \cdot \frac{\partial z_i^{(l+1)}}{\partial z_j^{(l)}} \\ &= \sum_{i=1}^{m^{(l+1)}} \varepsilon_i^{(l+1)} \cdot w_{i,j}^{(l+1)} \cdot S'(z_j^{(l)}) \\ &= \sum_{i=1}^{m^{(l+1)}} w_{i,j}^{(l+1)} \cdot \varepsilon_i^{(l+1)} \cdot S'(z_j^{(l)})\end{aligned}$$

Therefore, we have now proven equation (BP2). We proceed to prove equation (BP4).

According to the chain rule we have

$$\frac{\partial C_{\mathbf{x},\mathbf{y}}}{\partial w_{j,k}^{(l)}} = \frac{\partial C_{\mathbf{x},\mathbf{y}}}{\partial z_j^{(l)}} \cdot \frac{\partial z_j^{(l)}}{\partial w_{j,k}^{(l)}}$$

Now by definition of $\varepsilon_j^{(l)}$, the first factor on the right hand side of this equation is equal to $\varepsilon_j^{(l)}$:

$$\varepsilon_j^{(l)} = \frac{\partial C_{\mathbf{x},\mathbf{y}}}{\partial z_j^{(l)}}.$$

In order to proceed, we need to evaluate the partial derivative $\frac{\partial z_j^{(l)}}{\partial w_{j,k}^{(l)}}$. The term $z_j^{(l)}$ has been defined as follows:

$$z_i^{(l)} = \left(\sum_{k=1}^{m^{(l-1)}} w_{i,k}^{(l)} \cdot S(z_k^{(l-1)}) \right) + b_i^{(l)}$$

Hence we have

$$\frac{\partial z_j^{(l)}}{\partial w_{j,k}^{(l)}} = S(z_k^{(l-1)}) = a_j^{(l-1)}$$

Combining these equations we arrive at

$$\frac{\partial C_{\mathbf{x},\mathbf{y}}}{\partial w_{j,k}^{(l)}} = \varepsilon_j^{(l)} \cdot a_j^{(l-1)}$$

Therefore, equation (BP4) has been verified.

Exercise 13: Prove equation (BP3). ◇

7.3 Stochastic Gradient Descent

The equations describing backpropagation describe the gradient of the cost function for a single training example $\langle \mathbf{x}, \mathbf{y} \rangle$. However, when we train a neural network, we need to take all training examples into account. If we have n training examples

$$\langle \mathbf{x}^{(1)}, \mathbf{y}^{(1)} \rangle, \langle \mathbf{x}^{(2)}, \mathbf{y}^{(2)} \rangle, \dots, \langle \mathbf{x}^{(n)}, \mathbf{y}^{(n)} \rangle,$$

then the quadratic error cost function has been previously defined as the sum

$$C(W^{(2)}, \dots, W^{(L)}, \mathbf{b}^{(2)}, \dots, \mathbf{b}^{(L)}; \mathbf{x}^{(1)}, \mathbf{y}^{(1)}, \dots, \mathbf{x}^{(n)}, \mathbf{y}^{(n)}) := \frac{1}{2 \cdot n} \cdot \sum_{i=1}^n \left(\mathbf{o}(\mathbf{x}^{(i)}) - \mathbf{y}^{(i)} \right)^2.$$

In practical applications of neural networks, the number of training examples is usually big. For example, when we later develop a neural network to classify handwritten digits, we will have 60,000 training examples. More ambitious projects that use neural networks to classify objects in images use millions of training examples.

When we compute the gradient of the quadratic error function with respect to a weight matrix $W^{(l)}$ or a bias $b^{(l)}$ we have to compute the sums

$$\frac{1}{2 \cdot n} \cdot \sum_{i=1}^n \frac{\partial C_{\mathbf{x}^{(i)}, \mathbf{y}^{(i)}}}{\partial w_{j,k}^{(l)}} \quad \text{and} \quad \frac{1}{2 \cdot n} \cdot \sum_{i=1}^n \frac{\partial C_{\mathbf{x}^{(i)}, \mathbf{y}^{(i)}}}{\partial b_j^{(l)}}$$

over all training examples in order to perform a single step of gradient descent. If n is large, this is computationally costly. Note that these sums can be regarded as computing average values. In **stochastic gradient descent**, instead we approximate these sums by randomly choosing a small subset of the training examples. In order to formulate this approximation in a convenient notation, let us assume that instead of using all n training examples, we just use the first m training examples. Then we approximate the sums show above as follows:

$$\frac{1}{2 \cdot n} \cdot \sum_{i=1}^n \frac{\partial C_{\mathbf{x}^{(i)}, \mathbf{y}^{(i)}}}{\partial w_{j,k}^{(l)}} \approx \frac{1}{2 \cdot m} \cdot \sum_{i=1}^m \frac{\partial C_{\mathbf{x}^{(i)}, \mathbf{y}^{(i)}}}{\partial w_{j,k}^{(l)}} \quad \text{and} \quad \frac{1}{2 \cdot n} \cdot \sum_{i=1}^n \frac{\partial C_{\mathbf{x}^{(i)}, \mathbf{y}^{(i)}}}{\partial b_j^{(l)}} \approx \frac{1}{2 \cdot m} \cdot \sum_{i=1}^m \frac{\partial C_{\mathbf{x}^{(i)}, \mathbf{y}^{(i)}}}{\partial b_j^{(l)}},$$

i.e. we approximate these sums by the average value of their first m training examples. Of course, in general we will not choose the first m training examples but rather we will choose m **random** training examples. The randomness of this choice is the reason this algorithm is called **stochastic** gradient descent. It turns out that if we take care that eventually all training examples are used during gradient descent, then the approximations given above can speed up the learning of neural networks substantially.

7.4 Implementation

Next, we will take a look at a neural network that is able to recognize digits.

```

1  zeros_vec := n      |-> la_vector([0] * n);
2  zeros_mat := [m, n] |-> la_matrix([0] * n * m);
3  rndVector := length |-> la_vector([random() - 0.5: i in [1..length]]);
4  my_rnd    := []     |-> (random() - 0.5) / 14; // 14 = 28 / 2
5  rndMatrix := [rs, cs] |-> la_matrix([my_rnd(): c in [1..cs]]: r in [1..rs]);
6  sigmoid   := [x]     |-> la_vector([1 / (1 + exp(-a)) : a in x]);
7  // derivative of the sigmoid of a vector
8  sigmoid_prime := procedure(x) {
9      s := sigmoid(x);
10     return la_vector([a * (1 - a): a in s]);
11 };
12 hadamard := [x, y] |-> la_vector([x[i] * y[i]: i in [1 .. #x]]);
13 // compute the index of the biggest value in x
14 argmax := procedure(x) {
15     [maxValue, maxIndex] := [x[1], 1];
16     for (i in [2 .. #x] | x[i] > maxValue) {
17         [maxValue, maxIndex] := [x[i], i];
18     }
19     return maxIndex;
20 };

```

Figure 7.1: Auxiliary functions.

```

1  class network(inputSize, hiddenSize, outputSize) {
2      mInputSize := inputSize; // 784
3      mHiddenSize := hiddenSize; // 30 .. 100
4      mOutputSize := outputSize; // 10
5      mBiasesH := rndVector(mHiddenSize);
6      mBiasesO := rndVector(mOutputSize);
7      mWeightsH := rndMatrix(mHiddenSize, mInputSize);
8      mWeightsO := rndMatrix(mOutputSize, mHiddenSize);
9      ...
10 }

```

Figure 7.2: The constructor of the class `network`.

```

1  sgd := procedure(training_data, epochs, mbs, eta, test_data) {
2      n_test := #test_data;
3      n := #training_data;
4      for (j in [1 .. epochs]) {
5          training_data := shuffle(training_data);
6          mini_batches := [training_data[k .. k+mbs-1]: k in [1, mbs..n]];
7          for (mini_batch in mini_batches) {
8              update_mini_batch(mini_batch, eta);
9          }
10         print("Epoch $j$: $evaluate(test_data)$ / $n_test$");
11     }
12 };
13 update_mini_batch := procedure(mini_batch, eta) {
14     nabla_BH := zeros_vec(mHiddenSize);
15     nabla_BO := zeros_vec(mOutputSize);
16     nabla_WH := zeros_mat(mHiddenSize, mInputSize);
17     nabla_WO := zeros_mat(mOutputSize, mHiddenSize);
18     for([x,y] in mini_batch) {
19         [dltNbl_BH, dltNbl_BO, dltNbl_WH, dltNbl_WO] := backprop(x, y);
20         nabla_BH += dltNbl_BH;
21         nabla_BO += dltNbl_BO;
22         nabla_WH += dltNbl_WH;
23         nabla_WO += dltNbl_WO;
24     }
25     alpha := eta / #mini_batch;
26     this.mBiasesH -= alpha * nabla_BH;
27     this.mBiasesO -= alpha * nabla_BO;
28     this.mWeightsH -= alpha * nabla_WH;
29     this.mWeightsO -= alpha * nabla_WO;
30 };

```

Figure 7.3: Stochastic gradient descent.

```

1  backprop := procedure(x, y) {
2      ZH := mWeightsH * x + mBiasesH;
3      AH := sigmoid(ZH);
4      ZO := mWeightsO * AH + mBiasesO;
5      AO := sigmoid(ZO);
6      delta := hadamard(AO - y, sigmoid_prime(ZO));
7      nabla_BO := delta;
8      nabla_WO := la_matrix(delta) * la_matrix(AH)!;
9      delta := hadamard(mWeightsO! * delta, sigmoid_prime(ZH));
10     nabla_BH := delta;
11     nabla_WH := la_matrix(delta) * la_matrix(x)!;
12     return [nabla_BH, nabla_BO, nabla_WH, nabla_WO];
13 };

```

Figure 7.4: Implementation of backpropagation.

```

1  feedforward := procedure(x) {
2      AH := sigmoid(mWeightsH * x + mBiasesH);
3      AO := sigmoid(mWeightsO * AH + mBiasesO);
4      return AO;
5  };
6  evaluate := procedure(test_data) {
7      test_results := [[argmax(feedforward(x)) - 1, y]: [x, y] in test_data];
8      return #[1 : [a, b] in test_results | a == b];
9  };

```

Figure 7.5: Evaluation functions.

```

1  load("nn-loader.stlx");
2  load("nn.stlx");
3
4  main := procedure() {
5      resetRandom();
6      training_length := 60000;
7      test_length := 10000;
8      inputSize := 784;
9      hiddenSize := 30;
10     outputSize := 10;
11     [training_data, test_data] := load_data_wrapper(60000, 10000);
12     print("Create Network");
13     net := network(inputSize, hiddenSize, outputSize);
14     print("Start SGD");
15     s1 := now();
16     net.sgd(training_data, 30, 10, 3.0, test_data);
17     s2 := now() - s1;
18     print("Time needed:\t" + s2);
19 };
20 main();

```

Figure 7.6: How to start the training.

```

1  load_data := procedure(training_length, test_length) {
2      test_data      := parse_csv("mnist_test.csv", test_length);
3      training_data := parse_csv("mnist_train.csv", training_length);
4      return [training_data, test_data];
5  };
6  parse_csv := procedure(file, len) {
7      images := [0 : i in [1..len]];
8      labels := [0 : i in [1..len]];
9      print("Reading file:\t$file$");
10     csv := readFile(file);
11     for (i in [1..len]) {
12         labels[i] := int(csv[i][1]);
13         images[i] := parse_line(csv[i][3..]);
14         if (i % 10000 == 0) {
15             print("Image $i$ of $len$ imported");
16         }
17     }
18     print("End reading:\t$file$");
19     return [images, labels];
20 };
21 parse_line := procedure(s) {
22     L := split(s, ",");
23     return [double(x)/255.0 : x in L];
24 };
25 vectorized_result := procedure(j) {
26     e      := la_vector([0] * 10);
27     e[j+1] := 1; // +1, because the Index starts with 1 in Set1X (not 0)
28     return e;
29 };
30 load_data_wrapper := procedure(training_length, test_length) {
31     [tr_d, te_d] := load_data(training_length, test_length);
32     test_data := format_pixels_label_pairs(te_d[1], te_d[2], test_length);
33     training_data := format_with_unit_vectors(tr_d[1], tr_d[2], training_length);
34     return [training_data, test_data];
35 };
36 format_pixels_label_pairs := procedure(pixels, labels, length) {
37     return [ [la_vector(pixels[i]), labels[i]] : i in [1..length] ];
38 };
39 format_with_unit_vectors := procedure(pixels, labels, length) {
40     return [ [la_vector(pixels[i]), vectorized_result(labels[i])] : i in [1..length] ];
41 };

```

Figure 7.7: Loading the data.

Bibliography

- [HNR68] Peter Hart, Nils Nilsson, and Bertram Raphael. A formal basis for the heuristic determination of minimum cost paths. *IEEE Transactions on Systems Science and Cybernetics SSC4*, 4(2):100–107, 1968.
- [HNR72] Peter Hart, Nils Nilsson, and Bertram Raphael. Correction to “A formal basis for the heuristic determination of minimum cost paths”. *SIGART Newsletter*, 37:28–29, 1972.
- [JWHT14] Gareth James, Daniela Witten, Trevor Hastie, and Robert Tibshirani. *An Introduction to Statistical Learning: with Applications in R*. Springer, 2014.
- [Kor85] Richard Korf. Depth-first iterative-deepening: An optimal admissible tree search. *Artificial Intelligence*, 27:97–109, 1985.
- [Nie15] Michael A. Nielsen. *Neural Networks and Deep Learning*. Determination Press, 2015.
- [RHW86] David E. Rumelhart, Geoffrey E. Hinton, and Ronald J. Williams. Learning internal representations by error propagation. In David E. Rumelhart and James L. McClelland, editors, *Parallel Distributed Processing: Explorations in the Microstructure of Cognition, Vol. 1*, pages 318–362. MIT Press, 1986.
- [RN09] Stuart Russell and Peter Norvig. *Artificial Intelligence: A Modern Approach*. Pearson Education Limited, 3rd edition, 2009.
- [SHM⁺16] David Silver, Aja Huang, Chris J. Maddison, Arthur Guez, Laurent Sifre, George van den Driessche, Julian Schrittwieser, Ioannis Antonoglou, Veda Panneershelvam, Marc Lanctot, Sander Dieleman, Dominik Grewe, John Nham, Nal Kalchbrenner, Ilya Sutskever, Timothy Lillicrap, Madeleine Leach, Koray Kavukcuoglu, Thore Graepel, and Demis Hassabis. Mastering the game of go with deep neural networks and tree search. *Nature*, 529:484, 2016.
- [Sny05] Jan A. Snyman. *Practical Mathematical Optimization: An Introduction to Basic Optimization Theory and Classical and New Gradient-Based Algorithms*. Springer Publishing, 2005.